

Future Internet Services based on LISP Technology

Yue LI

Department of INFRES
Telecom Paristech

This dissertation is submitted for the degree of
Doctor of Philosophy

Acknowledgements

And I would like to acknowledge ...

Abstract

Blabla ...

Table of contents

List of figures	xi
List of tables	xv
1 Introduction	1
1.1 Motivations	1
1.2 Contributions	3
1.3 Organization of Manuscript	4
2 LISP Overview	7
2.1 LISP Architecture	7
2.1.1 LISP Data Plane	8
2.1.2 LISP Control Plane	9
2.1.3 Communication between two LISP-sites	12
2.2 Interworking With Legacy Internet	14
2.3 Mapping Cache Update Mechanisms	15
2.3.1 Solicit-Map-Request (SMR)	15
2.3.2 Map-Versioning	16
2.4 LISP Mobility	17
2.4.1 LISP Mobile Node	18
2.4.2 LISP MN in LISP-Site	19
2.4.3 LISP MN in non-LISP-Site	20
2.4.4 MN in LISP-Site	21
2.5 LISP use-cases	22
2.5.1 LISP traffic engineering	22
2.5.2 LISP transition between IPv4 and IPv6	22
2.5.3 Mobility at client-side	23
2.5.4 Mobility in Data Center	24

2.5.5	LISP SDN	24
3	Related Work	27
3.1	LISP platforms	27
3.1.1	LISP Beta Network	27
3.1.2	LISP-Lab platform	29
3.2	LISP monitoring tools	29
3.2.1	LISPmon	29
3.2.2	LIG	30
3.2.3	RIG	30
3.2.4	Missing work	31
3.3	LISP implementations	31
3.3.1	Existing implementations	32
3.3.2	Missing work	34
3.4	LISP mapping system evaluation	34
3.4.1	Current studies	34
3.4.2	Missing work	35
3.5	LISP network evolution	35
3.5.1	Current studies	35
3.5.2	Missing work	36
3.6	LISP interworking with legacy Internet	36
3.6.1	Impact on traffic between LISP-site and the legacy Internet	36
3.6.2	Current studies	36
3.6.3	Missing work	37
3.7	LISP mobility	37
3.7.1	Current studies	37
3.7.2	Missing work	38
3.8	Troubleshooting problem	39
3.9	Other LISP Related studies	40
3.9.1	LISP Cache performance	40
3.9.2	LISP security	41
3.9.3	The maximum transmission unit (MTU) issue	42
4	Evaluation of LISP mapping system	43
4.1	Introduction	43
4.2	Methodology	44
4.2.1	Measurements Overview	44

4.2.2	Metrics	45
4.3	Mapping System Stability Evaluation	46
4.4	Mapping System Consistency Evaluation	50
4.4.1	Consistency Evaluation by MR	51
4.4.2	Consistency Evaluation by VP	54
4.5	Conclusion	56
5	LISP-Views: LISP mapping system monitor	59
5.1	Introduction	59
5.2	Proposed Monitoring Architecture	60
5.2.1	Motivation	60
5.2.2	Description of LISP-Views	61
5.3	LISP-Views Validation	63
5.3.1	Methodology	63
5.3.2	LISP-Views vs. LISPmon	64
5.4	Dissecting LISP with LISP-Views	66
5.5	Conclusion	71
6	Assessing LISP interworking performance through RIPE Atlas	73
6.1	Introduction	74
6.2	Related Work	74
6.2.1	RIPE Atlas	74
6.2.2	Alexa	75
6.3	Measurement Methodology	75
6.3.1	For Dataset2015	76
6.3.2	For Dataset2016	77
6.4	IPv4 Ping Results from Dataset2015	80
6.5	IPv4 Ping Results from Dataset2016	83
6.6	IPv6 Ping Results	88
6.7	Traceroute-related Results	92
6.8	Conclusion	95
7	ns-3 Implementation of LISP and LISP-MN	97
7.1	Introduction	97
7.2	Related work	98
7.3	LISP/LISP-MN Implementation based on ns3	99
7.3.1	Implementation of LISP Data Plane	99

7.3.2	Implementation of LISP Control Plane	100
7.3.3	Integration of TUN net interface card	101
7.3.4	Integration of DHCP	101
7.4	Evaluation	102
7.4.1	Simulation Setup	102
7.4.2	Results	102
7.5	Conclusion	104
8	Conclusion and Future Work	105
8.1	Major Contributions	105
8.1.1	Evaluation of LISP Mapping System	105
8.1.2	LISP-Views: LISP Mapping System Monitor	105
8.1.3	Assessing LISP interworking performance through RIPE Atlas . . .	105
8.1.4	ns-3 Implementation of LISP and LISP-MN	105
8.2	Discussion & Future work	105
	References	107
	Appendix A Source Code	117
A.1	LISP Stability and Consistency analyzer	117
A.2	LISP-Views	117
A.3	LISP measurements on Atlas	117
A.4	LISP implementation on ns-3	118

List of figures

1.1	Growth of the BGP Table - 1994 to Present	2
1.2	Structure of this dissertation	5
2.1	Hierarchical taxonomy of mapping systems (from [65])	9
2.2	Illustration of the LISP+ALT topology	10
2.3	Illustration of the LISP-DDT topology	11
2.4	LISP architecture overview	13
2.5	LISP architecture with the packets forwarding sequence	14
2.6	Packet sequence of LISP SMR mechanism	16
2.7	Packet sequence of LISP Map-Versioning mechanism	17
2.8	Illustration of LISP-MN in LISP-Site and non-LISP-Site	18
2.9	Packet flow of mobility when LISP-MN in LISP-Site	20
2.10	Packet flow of mobility when LISP-MN in non-LISP-Site	21
2.11	Illustration of LISP transmission between IPv4 and IPv6	23
2.12	LISP in OpenDaylight (source: [14])	25
3.1	International LISP Beta Network High Level Topology – October 2017 [12]	28
3.2	Illustration of LISP troubleshooting	39
4.1	Overall percentage of stability observed at five different VPs.	47
4.2	CDF of instability frequency. X-axis is instability frequency, indicating the ratio of instability occurrence number for one $\langle EID, MR, VP \rangle$ tuple over total experiment number.	48
4.3	CDF of stability duration (in number of experimental rounds). X-axis is stability duration, indicating during how many experiment rounds that the mapping for one $\langle EID, MR, VP \rangle$ tuple is stable.	49
4.4	CDF of the dissimilarity among different types of instabilities.	50
4.5	Overall consistency by MR at 5 different Vantage Points.	52

4.6	Cumulative Distribution Function of inconsistency frequency by MR. X-axis is inconsistency frequency by VP, indicating the ratio of inconsistency occurrence for each $\langle EID, *, VP \rangle$ tuple over total experiment number. . . .	53
4.7	Overall consistency by VP from the different Map-Resolvers.	54
4.8	Cumulative Distribution Function of inconsistency frequency by VP. X-axis is inconsistency frequency by VP, indicating the ratio of inconsistency occurrence for each $\langle EID, MR, * \rangle$ tuple over total experiment number.	55
5.1	Monitoring Architecture	62
5.2	Comparison between LISPmon and LISP-Views. The upper sub-figure shows the number of LISP Map-Reply from LISPmon generally at 7:00 and LISP-Views exactly at 8:00 over days. The bottom sub-figure indicates the LISP Map-Reply and overall Mapping from LISP-Views over time with an interval of 2 hours.	64
5.3	Comparison between the LISP Map-Reply of LISPmon and LISP-Views over days	66
5.4	Reliability of each MR	67
5.5	Median RTT per MR (In the most time, the Negative Map-Reply and overall are overlapped.)	68
5.6	ECDF RTT comparison between LISP and Negative Map-Reply	69
5.7	CDF of the number of RLOCs	70
5.8	EIDs by Quantity of Associated RLOCs [16]	71
6.1	Locations of probes and anchor	77
6.2	cdf of average RTT between different probes	80
6.3	Smallest mean RTT (left) and smallest median RTT (right) grouped by continent.	81
6.4	Relative mean (left) and median (right) RTT clustered by different continents.	82
6.5	CDF of median RTT between different probes (IPv4)	83
6.6	Smallest median RTT grouped by continent (IPv4).	84
6.7	IPv4 Relative median RTT clustered by different continents for LISP-Lab (left) and LIP6 (right)	86
6.8	ECDF of median RTT between different probes (IPv6)	88
6.9	IPv6 Relative median RTT clustered by different continents for LISP-Lab (left) and LIP6 (right)	89
6.10	Distribution of IPv4 Relative Hops Number for LISP-Lab (left) and LIP6 (right)	92

6.11	IPv4 Relative hops number clustered by different continents for LISP-Lab (left) and LIP6 (right)	93
6.12	Distribution of IPv6 Relative Hops Number for LISP-Lab (left) and LIP6 (right)	94
6.13	IPv6 Relative hops number clustered by different continents for LISP-Lab (left) and LIP6 (right)	95
7.1	UML diagram of LISP/LISP-MN implementation. The solid arrow refers to a composition relation, while the blank one refers to a inheritance relation.	99
7.2	LISP mobility simulation scenario for double encapsulation	102
7.3	Schema for LISP-MN mobility	103

List of tables

2.1	Catalogs of mobility	18
4.1	Map-Reply Differences	45
4.2	Percentage of Instabilities by Category	48
5.1	Percentage of Mapping Source	69
6.1	Different configurations of probes in 2015	76
6.2	Location of LISP Beta Network PxTRs in 2016	77
6.3	Configuration of two LISP probes	78
6.4	Different configurations of probes in 2016	79
6.5	Correlation coefficient to FranceIX	81
6.6	Correlation coefficient to FranceIX (IPv4)	85
6.7	Reliability of each probe (IPv4)	87
6.8	Correlation coefficient to FranceIX (IPv6)	91
6.9	Reliability of each probe (IPv6)	91

Chapter 1

Introduction

1.1 Motivations

Since the ARPANET (Advanced Research Projects Agency Network) emerged in 1966 [84], the Internet has obtained significant achievements over the last fifty years and developed very fast. Until the beginning of 2018, the Internet users in the world achieves 3.82 billions, i.e., around 40% of the world population has an Internet connection today [9]. Moreover, we come into an era of the mobile Internet. According to the latest research [30], the number of global mobile users has been higher than that of desktop since 2014. Although having a huge success, today's Internet is suffering the overloading of IP address semantics, which results in many shortcomings in the application landscape [53]. As in the current TCP/IP protocol stack, the IP address serves both *who* to identify the host and *where* to indicate the attachment point in the network topology, it leads to a serious routing scalability problem. As shown in Fig. 1.1, the BGP table size, which is the number of IP prefixes in the BGP (Border Gateway Protocol) routing table of DFZ (Default Free Zone) routers, has grown at a super-linear rate in recent years. After full use of IPv6, the routing scalability issue will be aggravated, since it has much larger addressing space than IPv4. Maintaining such routing table is very cumbersome. Particular, once a site changes its routing information, all the corresponding routers need to update, which causes a huge BGP churn. In addition, frequent routing information updates result in the increasing of convergence time and packet loss rates.

To cope with the current Internet scalability problems, the Loc/ID Split Networking is proposed, which separates the single IP addressing space into two roles: *who* and *where*. Many solutions based on this mechanism are proposed, which can be classified into two catalogs: *Host-based* and *Network-based*. The former one needs modify the host stack. The representative architectures are: HIP (Host Identity Protocol) [93] [91] [90], Six/One [114],




Fig. 1.1 Growth of the BGP Table - 1994 to Present

SHIM6 (Site Multihoming by IPv6 Intermediation) [61] [94], ILNP (Identifier-Locator Network Protocol) [37] [34], and LNA (Layered Naming Architecture) [35]. The later one needs to update the routers, and the representative proposals are: GSE (Global, Site, and End-system address elements) [95], LISP (Locator/ID Separation Protocol) [47], and TIDR (Tunneled Inter-domain Routing) [31]. Among them, LISP is outstanding by its relatively comprehensive network systems in both control and data planes. This characterization promotes the standardization, production and industrialization of LISP. It was initially proposed by Cisco in 2006 and has an exclusive IETF WG since 2009, which guarantees LISP can be actively under developed.

The key philosophy of LISP is to separate the IP addressing space into two sub-spaces: the *Routing LOCators* (RLOCs) and the *Endpoint IDentifiers* (EIDs). End-to-end communication is based on EIDs, which in an intra-domain context means just to performing conventional routing and forwarding. On the contrary, for inter-domain end-to-end communication, RLOCs are used for routing and forwarding. To this end, LISP does a map-and-encap operation, first mapping EIDs to RLOCs (through a DNS-like on-demand mapping system [55], where Map-Resolver (MR) is used to receive the query) and then encapsulating the original packet so to use RLOCs on the outer header. Internetworking between LISP and non-LISP networks is performed through two types of proxies [76]: *Proxy Ingress Tunnel Routers* (PITRs) and *Proxy Egress Tunnel Routers* (PETRs). They are collectively denoted by *PxTR*. PITRs advertise in the BGP routing infrastructure the EID-prefix space on behalf of the

LISP sites, so that non-LISP sites can reach them. In this way PITRs attract the traffic from non-LISP sites, which simply send packets with destination EID as destination IP address. Subsequently the PITR performs LISP encapsulation and forward the packets to the destination. PETRs are used as exit points from LISP sites when the destination is not an EID in a LISP site. Hence, PETRs decapsulate the LISP-encapsulated packets received from LISP-sites and forwards them in the legacy Internet. With help of this detaching mechanism, LISP is able to improve the Internet routing scalability, enable enhanced inter-domain traffic engineering, lightweight multi-tenant carrier-grade VPN, and VMs mobility in Data-Centers [109].

As LISP is a promising technology that brings many benefits to the future Internet, it attracts the attention from both academic and industrial world. Many companies, universities and research institutes are now involved in the development of LISP. Hence, the related standardization and researches are very active. In terms of deployment, two large-scale experimental LISP platforms: LISP Beta Network and LISP-Lab platform are deployed on the world respectively in 2008 and 2015. Up to date, all the efforts are mainly about the theoretic performance analysis, description of LISP architectures, implementations, extensions, applications and so on. However, it lacks of the comprehensive measurement works for LISP technology itself. These performance measurements work will give the research community a clear view about the LISP performance, reveal the points that can be improved, indicate possible research directions, and improve the performance of two platforms. Thus, the large-scale measurements in realistic networks are of paramount importance for the LISP development.

1.2 Contributions

With the goal of evaluating LISP from the different views, in this dissertation we aim at answering the following research questions:

1. Mapping system is essential in LISP, because it decides where the packets will be routed. Thus, does it always provide the identical mapping information for the same destination? Does every MR offer the same response at the same time? Does the mapping information replies keep consistent if we query them from the different Vantage Point (VP)s?

We continuously measured LISP Beta Network for seventeen days. Measurements show that the LISP mapping system is stable over most of time. The replies from the different MRs and VPs are also mostly consistent. Nevertheless, instability and inconsistency are observed although they are rare events. We define a new taxonomy

to classify these instability and inconsistency so to facilitate the deeper analysis. This work is published in WNM [81].

2. As the results of the last question show that there exists inconsistency between the MRs and VPs, is it possible to implement a more comprehensive monitor than today's LISPmon so to simultaneously supervise all the MRs as experimenter define instead of just supervising one MR?

We propose a dynamic LISP monitoring architecture, namely LISP-Views, so to deepen the understandings on LISP and to ease day-to-day operations and troubleshooting. After one full month monitoring to the whole LISP mapping system of two platforms, it demonstrates that LISP-Views provides more information by discovering more mapping information from all MRs and more complete mapping information of each MR. Furthermore, with our proposed monitoring platform, more mapping system performance metrics such as reliability, latency, and configuration issues can be assessed, which helps for further LISP improvements. This work is published in ITC [77].

3. By using proxy approach to communicate with the legacy Internet, a stretch in the path is introduced. Thus, how much is it in the real world performance when the platforms integrate with the legacy Internet? Under which conditions such stretch has an important impact on the performance? Conversely, under which situation such overhead is so small that it can be ignored?

To evaluate LISP interworking performance with legacy Internet, we conduct two times experiments. The results of twice experiments are coherent. Both of them show that LISP is stable although the stretch is introduced due to the introducing of proxy. We find that the proxy indeed introduces the negative effects for the destinations located in Europe and America, but can be ignored for the intercontinental long-distance transmission to Asia destinations. Besides, we observe that the position of proxy is very important. It either near to the sources or the destinations can decrease the latency a lot. Further, the performance of LISP-Lab PxTR is more reliable than the one of LISP Beta Network. This work is published in INFOCOM WKSHPS [79], SIGCOMM [78], and Computer Networks [80].

1.3 Organization of Manuscript

The remainder of this dissertation is as follows and shown in Fig. 1.2:

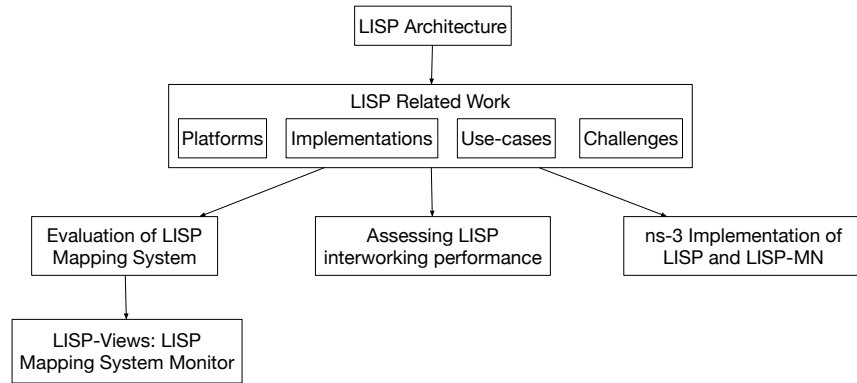


Fig. 1.2 Structure of this dissertation

- *Chapter. 2* describes the basic LISP knowledge, presenting how the packets are exchanged between two LISP-sites and with the legacy Internet. It also shows the different scenario of LISP mobility.
- *Chapter. 3* introduces the LISP current status, such as: the available LISP platforms, the first LISP monitor, the LISP manual management tool, the open source LISP implementations and LISP use-cases. The challenges brought by LISP, LISP related studies and missing works are also illustrated.
- *Chapter. 4* is about the evaluation of LISP Mapping System. In this chapter, we analyze the first question in Sec. 1.2 in details.
- *Chapter. 5* describes the proposition of LISP-Views: LISP Mapping System Monitor, i.e., the second contribution in Sec. 1.2.
- *Chapter. 6* assesses LISP interworking performance through RIPE Atlas, where introduces how the experiments are conducted and how to get the results of the third contribution in Sec. 1.2.
- *Chapter. 7* introduces ns-3 Implementation of LISP and LISP-MN.
- *Chapter. 8* concludes the dissertation and opens up new perspectives for further work.

Chapter 2

LISP Overview

In this Chapter, the LISP related concepts and its mechanism we refer to in the reminder of the dissertation are introduced. The basic architecture of LISP is introduced by presenting its Data Plane and Control Plane. We detail how the traffic between two LISP-sites is exchanged by using an example. A new network element, which is used to allow communication between LISP-site and the legacy Internet, is also presented, as well as the schema of packets exchange between them. Then, two mechanisms used to update the LISP mapping cache are illustrated. Finally, this chapter ends with the explanation of LISP mobile node, i.e., the LISP-enabled terminal, and the introduction about how the LISP procedure is performed when the LISP-enabled node is behind a LISP-speaking router.

2.1 LISP Architecture

The *Locator/Identifier Separation Protocol (LISP)* separates the traditional IP address role into two logical sub-spaces: (i) *Endpoint Identifier (EID)* and (ii) *Routing LOCator (RLOC)*. Both of them are the conventional IPv4 (32-bit) [103] or IPv6 (128-bit) [45] addresses. The former is the identifier of terminal hosts, which is locally routable in the LISP-site, i.e., within the subnet. Normally they are the addresses of source and destination host. Within the same LISP-site, the hosts can directly communicate with each other via EIDs as in a legacy IP network. The host gets a destination EID in the same way it obtains the destination IP address today, for instance through a Domain Name System (DNS) [87] lookup or Session Initiation Protocol (SIP) [105]. RLOCs denote the attachment points of EIDs in the Internet topology, i.e., the source/destination address of the edge routers, behind whom the EIDs can be found. They are topologically assigned to network attachment points and amenable to aggregation. RLOCs addresses are used to forward the packets in the core Internet.

When packets are transferred between two different LISP-sites, an encapsulation on the edge router is needed: the source and remote destination EIDs are in the inner packet header, while the source and destination RLOCs are in the outer header. The EID-prefixes are not necessary to be announced on the Internet core, so the BGP routing table size can be reduced. Similarly, in stub networks, it does not need to know the routing information of core Internet any more. The aforementioned encapsulation and forwarding operations are the responsibilities of LISP Data Plane. More details are introduced in Sec. 2.1.1. To bind the EIDs and RLOCs, a new network entity named Mapping Distribution System (MDS) is used, consisting of Map Resolver (MR) and Map Server (MS) [56]. The MDS is not only able to provide the mapping information between the EIDs and RLOCs for the LISP-sites (we simply call them LISP Map-Replies), but also indicate for which addresses are not LISP sites and have no RLOCs, meaning that are part of the legacy Internet (these are called Negative Map-Replies). The Control Plane is further described in Sec. 2.1.2.

2.1.1 LISP Data Plane

The LISP Data plane mainly takes care of the encapsulation and decapsulation operations done at the source and destination networks. LISP basically provides a level of indirection through a tunneling mechanism over the core Internet (the so called *Default-Free Zone* – DFZ), in order to provide end-to-end communication. More specifically, any host willing to communicate with another host residing at remote LISP-site, uses its own EID as source address and the EID of other host as destination address to generate regular IP packets. These packets are first transferred as usual to the border router, called *Ingress Tunnel Router (ITR)*. ITR performs a mapping lookup in its Cache or in case of miss query to the MDS for the mapping information. Then the ITR encapsulates the regular IP packets into a LISP packet by adding a tunnel header [47], where the RLOC of the ITR is used as source address and the destination RLOC as destination address. After encapsulation, the LISP packets are forwarded over the core Internet. When arriving at the destination border router, called *Egress Tunnel Router (ETR)*, the LISP packets are decapsulated (i.e., the tunnel header is removed) and transferred to the destination host. A device which acts as both an ITR and an ETR is denoted as *xTR*.

Mappings are stored in two data structures on the xTRs: the *LISP Database* and the *LISP Cache*. The LISP Database is populated by configuration and stores all known EID-to-RLOC mappings, for which the EID-Prefixes are behind the xTR. On an ITR, this helps select a source RLOC for an EID used in encapsulation. While on an ETR, it allows to verify whether itself is the proper ETR connecting to the destination EID, so that such ETR is able to forward the decapsulated packets to the final destination.




Fig. 2.1 Hierarchical taxonomy of mapping systems (from [65])

The LISP Cache temporally stores mappings for the EID-prefixes of the remote communicating end-points. On an ITR, it is used to select the destination RLOC of the outer header of the LISP-encapsulated packet. While on an ETR, it is used to perform a basic anti-spoof verification. Different from the LISP Database, the LISP Cache is populated on demand. The procedure is triggered by the first packet of a new flow, which can not find a suitable mapping for the destination EID among the mappings stored in the LISP Cache. The mappings in the LISP Cache are purged if not used upon expiration of a timeout ([66, 72, 116, 43]).

2.1.2 LISP Control Plane

The LISP Control plane is mainly tasked to discover the associations of EIDs with RLOCs, i.e., find the *mappings*. A new network entity named *mapping system* is introduced by LISP to perform such EID-to-RLOC distribution. In particular, a mapping contains an EID prefix and an associated list of RLOC tuples: $\langle RLOC, Priority, Weight \rangle$. The RLOC refers to the IP address of each xTR interface. Each RLOC is assigned a Priority. The RLOC with the highest priority is preferred. A value of 255 means that this RLOC must not be used for forwarding. The weight is taken into account only to perform load balancing in the case of several RLOCs having the same priority. Weight is encoded as a relative weight of total packets that match the mapping entry.

The current BGP-based Internet routing architecture relies on a push model, in which routing information is pushed to the whole Internet. The Locator/Identifier Split-based In-




Fig. 2.2 Illustration of the LISP+ALT topology

ternet routing instead relies on a pull model, the *Mapping Distribution System (MDS)* to provide a mapping upon an explicit query,¹ where routing information is pulled only when actually needed. This paradigm shift requires making routing information available in an on-demand manner. Fig. 2.1 shows the current possible taxonomy of mapping systems [65]. So far, several mapping systems have been proposed for LISP, such as: LISP-TREE [69], LISP-NERD (Not-so-novel EID to RLOC Database) [75] and LISP-CONS (Content distribution Overlay Network Service for LISP) [39]. However, only two have been deployed: *LISP Alternative Logical Topology (LISP+ALT)* [57] and *LISP Delegated Database Tree (LISP-DDT)* [58].

LISP+ALT

LISP+ALT was the initial mapping system for LISP, where the ETRs store mappings they are authoritative for [106]. The basic idea is to use BGP to construct an overlay, named Alternative Logical Topology (ALT), to establish reachability between ETRs via tunnels (e.g., GRE tunnels [50]). As shown in Fig. 2.2, the routers in the ALT are called ALT-Routers. Each ALT-Router maintains a BGP session with its neighbor and announces the

¹The terms *Mapping Distribution System (MDS)* and *mapping system* are used interchangeably in this dissertation.




Fig. 2.3 Illustration of the LISP-DDT topology

EID-prefixes it is authoritative for. These EIDs are then routable in the ALT. Note that the ALT-Routers only exchange aggregated EID-prefixes that can be reached through them between neighbors, but they do not contain the mapping information. To get a mapping, an ITR constructs a Map-Request for the queried EID by using its own RLOC as source address and the queried EID as destination address, then sends it to an ALT-Router. The Map-Request is forwarded over the ALT and eventually reaches an originator ETR for the EID-prefix that matches the destination EID. This ETR resolves the Map-Request and directly sends back a Map-Reply to the inquirer ITR, this time not using the ALT. LISP-ALT belongs to the Map-Bases with Partial Knowledge using Hierarchically Structured Overlay (MBPK-HSO) in Fig. 2.1.

LISP-DDT

LISP-DDT is a hierarchical distributed database that embodies the delegation of authority to provide mappings from LISP EIDs to RLOCs. Its topology is shown in Fig. 2.3. It consists of *Map-Resolvers (MRs)* and *Map-Servers (MSes)* to provide a simplified "front end" for one or more EID to RLOC mapping databases. The MR is a network infrastructure component that accepts LISP Encapsulated Map-Requests from an ITR and resolves the mapping using a DNS-like distributed database. The MS learns of EID-Prefix mapping entries (a list of

EID-Prefixes plus a set of RLOC tuples) from ETRs. The hierarchy is maintained as a tree where a root server is responsible for the entire EID space. This space is divided into several portions and each one is managed by one of the root's children. Mapping information is only stored at the tree leaves, which are made of MSes. Intermediate nodes only maintain pointers to their children.

If an ITR cannot find a mapping in LISP Cache for a new flow, it sends a query, called *Map-Request*, to a MR [56]. The MR determines whether the queried destination IP address is an EID. If not, a *Negative Map-Reply* is directly returned; otherwise, the root server is queried first and replies with a pointer to its child node responsible for the queried EID. The process is recursively repeated to find the child root of the sub-tree where a mapping can be retrieved for the queried EID. When the leaf is reached, the mapping is then retrieved by sending a Map-Request to the ETR authoritative for the matching EID-prefix. The ETR in turn directly replies a *Map-Reply* message containing the whole requested mapping information to the ITR initiating the query.

When an ETR first contacts an MS after restarting or changing its database, it sends the Map-Register messages to the MS to publish its EID-prefixes as well as the mapping information every 20 seconds. For maintaining the association, the registration interval should be expanded to at least 1 minute. If the MS has not received a valid Map-Register within the past 3 minutes, it removes the registration of this ETR. If the "want-map-notify" (M-bit) flag is set in Map-Register, the received MS needs to send back a Map-Notify message to inform the ETR that the Map-Register has been received and processed. If the "proxy Map-Reply" flag (P-bit) is set, the MS answers the Map-Requests to the queried ITR by providing the mapping information on behalf of ETR. MRs offer an interface of MDS to the xTRs, so that the whole complex MDS mechanism can be hidden.

2.1.3 Communication between two LISP-sites

We use Fig. 2.4 as an example to describe how the LISP packets are processed in details in the case of the packets exchanging between two LISP-sites.

When the host in the AS_s communicates with the host in the AS_d in Fig. 2.4, it needs to use EID_s as source address and EID_d as destination address. After the conventional IP packets routing to xTR_1 (the function of ITR and ETR are combined together in Fig. 2.4) labeled as "1", ITR_1 first checks in its mapping cache [66] to find out the association between EID_d and its RLOC. If there is no record, it sends a Map-Request to Map-Resolver (MR), labeled as "2". MR searches which Map-Server (MS) stores the mapping information of EID_d and forwards the Map-Request to that MS. If MS is authoritative (i.e., acts as a proxy), it directly returns back a Map-Reply containing the mapping information (this procedure is not labeled

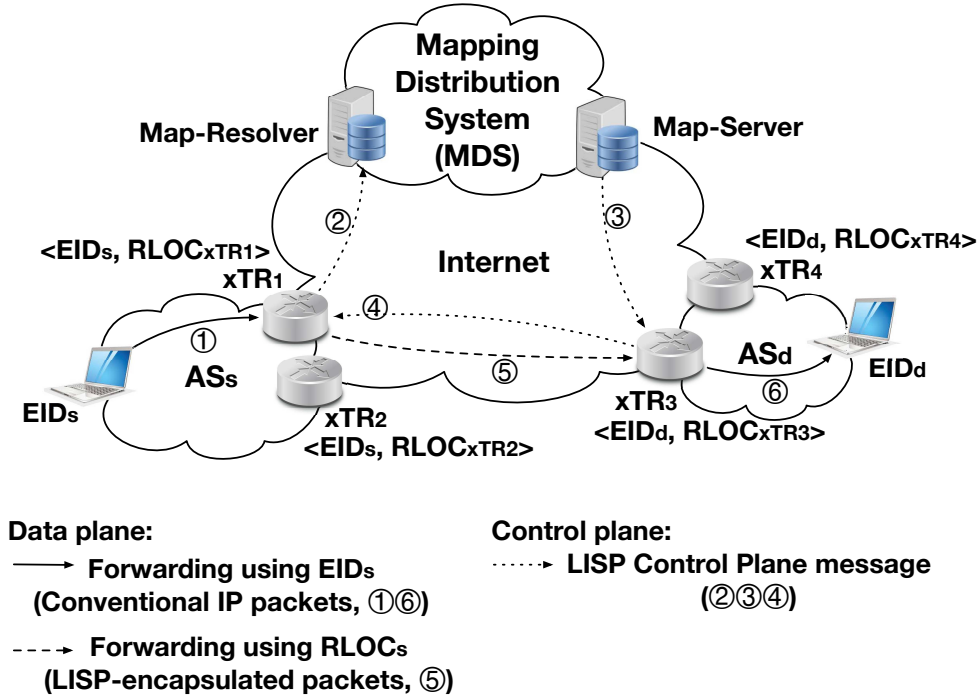


Fig. 2.4 LISP architecture overview

in the figure). If not, after forwarding the Map-Request within MDS with the procedure described in Sec. 2.1.2 to MS and then to the destination side (label "3"), ETR_1 receives a Map-Reply from one of ETRs of EID_d (label "4"). If the Map-Reply shows that the priority of ETR_3 is higher than ETR_4 , then the ITR_1 encapsulates the packets by adding $RLOC_{ETR_3}$ as the destination address and $RLOC_{ITR_1}$ as the source address in the outer header and send the encapsulated packets through Internet as labeled "5" in Fig. 2.4. Once ETR_3 receives these packets, it decapsulates and forwards them by using the original IP packets (label "6"). If the destination site is non LISP-site, the MR directly returns back the Map-Reply without mapping information to ITR_1 .

Based on LISP architecture, there are three types of outcomes for a query:

- *LISP Map-Reply*, means that the queried IP address belongs to a LISP site (i.e., EID) and the Map-Reply contains the mapping information for this site, including: the association between the EID-prefix that queried EID belongs to, and the list of RLOC tuples $\langle RLOC, Priority, Weight \rangle$ of the destination site.
- *Negative Map-Reply*, means that the prefix covering the queried IP address belongs to a non-LISP site (i.e., conventional site), and the Map-Reply contains no mapping information. We consider these two kinds of replies as a successful query.

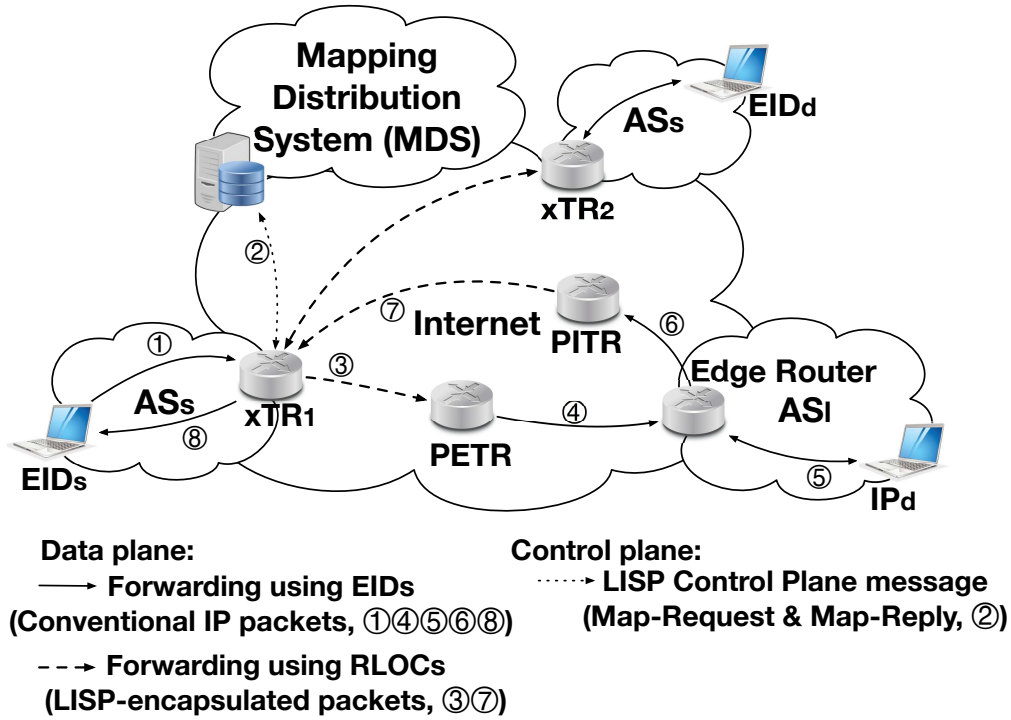


Fig. 2.5 LISP architecture with the packets forwarding sequence

- *No Map-Reply*, means that the xTR does not receive any reply during a certain time. According to the standard described in [47], the time-out is set to 3 seconds. But this value can be modified in the implementation to adapt the experiment needs. In this case, we consider the query as failed.

2.2 Interworking With Legacy Internet

Sec. 2.1.3 describes how the LISP packets are exchanged between two LISP-sites. This section discusses how the packets are processed when a terminal behind LISP-site communicates with a terminal on the legacy Internet.

Although EIDs are syntactically identical to IPv4 or IPv6 addresses, normally routes to them are not announced in the global routing system, so an interoperability mechanism is needed for non-LISP-speaking sites to exchange traffic with LISP-speaking sites. Thus, two new network elements are introduced: LISP Proxy Ingress Tunnel Router (PITR) and LISP Proxy Egress Tunnel Router (PETR). The PITR acts as an intermediate LISP ITR for the hosts on legacy Internet. It is in charge of announcing one or more highly aggregated EID-Prefixes on behalf of LISP-sites into the public Internet and receives the traffic from

the public Internet to encapsulate them as LISP packets. The PETR acts as an ETR for traffic destined to non-LISP sites. It receives the traffic from the ITRs, decapsulates the LISP packets into conventional IP packets and sends them to the legacy Internet.

Fig. 2.5 shows an example for the packets exchanging between LISP-site and non LISP-site (i.e., legacy Internet). The host with EID_s in the AS_s exchanges the packets with the host with the IP_d in the AS_l . After the same procedure introduced in Sec. 2.1.3 labeled as "1" and "2" in Fig. 2.4 that the original IP packets are sent from host to xTR_1 , and the later one sends the Map-Request to MDS to query the mapping information. Since the destination is in non LISP-site, the MR directly returns back the Map-Reply without mapping information (i.e., Negative Map-Reply) to xTR_1 , and the packet is not encapsulated but natively forwarded without using LISP. If leveraging LISP, the xTR_1 encapsulates the packets using the statically configured RLOC of PETR as the outer destination address and forwards the LISP packets there. When PETR receives the LISP-encapsulated packets, it conducts a decapsulation as ETR and then forwards the traditional IP packets to the legacy destination host by using the normal BGP routing, labeled as procedure "4" and "5" in Fig. 2.5.

When the packets come back, it is counted in a traditional way from the host with address IP_d in AS_l to PITR, denoted as procedure "5" and "6" in Fig. 2.5. PITR shares many characteristics with ITR so that it also queries the mapping information of destination to MDS and encapsulates the non-LISP Internet traffic into LISP-encapsulated packets and route them to their destination RLOCs (label "7"). Target ETR_1 then decapsulates the packets and send them to the destination host EID_s . The PETR and PITR can be noted together as PxTR if they are located on a same physical router.

2.3 Mapping Cache Update Mechanisms

Given that LISP mapping retrieve is based on pull model, in case that mapping changes occur in ETRs, the only way that makes remote ITRs be aware of such change and update the corresponding mapping information is to request again the mappings to ETRs. Thus, two mechanisms are proposed to get the latest mapping: Solicit-Map-Request and Map-Versioning.

2.3.1 Solicit-Map-Request (SMR)

Soliciting a Map-Request is used by ETRs to tell remote ITRs to update the mappings they have cached at the aspect to Control Plane. When the mappings in the Database of ETR

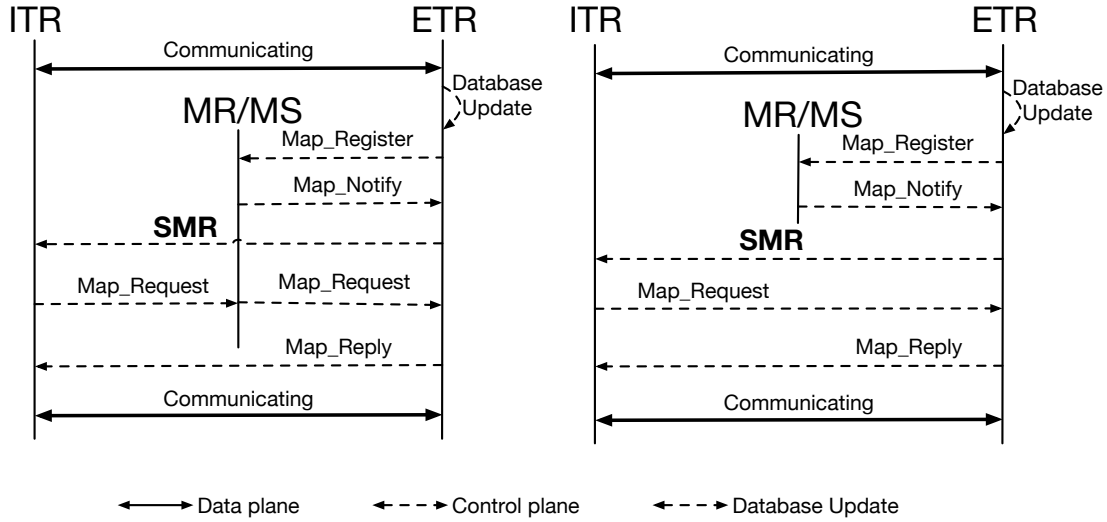


Fig. 2.6 Packet sequence of LISP SMR mechanism

change, the ETR sends the Map-Requests with the SMR bit set (called Solicit-Map-Request (SMR) message) for each RLOC in its Cache. A remote ITR that receives the SMR message will take a check in its Cache whether the source RLOC of the SMR message, i.e., the RLOC of ETR with mappings change is in its Cache. If yes, the remote ITR sends a Map-Request to the MR (as shown in the left figure of Fig. 2.6) or directly to the sender of SMR (as shown in the right figure of Fig. 2.6). If the RLOC of sender is not in the Cache of ITR any longer, the later drops the SMR and does not send the Map-Request. The later situation occurs when the ETR sends the SMR to all RLOCs stored in its Cache, but the remote ITR actually has not sent the packets to it for a long time. Between sending the SMR message and receiving the new Map-Reply, the ITR continues to use the mapping information previously cached, and that may cause the packet loss.

2.3.2 Map-Versioning

Map-Versioning is used to inform the xTR that the mappings of remote site change through Data Plane and update the mapping information with the help of SMR message. Map-Versioning associates a Map-Version number to each LISP mapping information and transports such a version number in the LISP-specific header [67]. When a mapping changes, a new version number, normally is incrementally higher than the previous one, is assigned to the updated mapping. When an ITR encapsulates a packet, it selects the version number assigned to the mapping stored in the Database as *source RLOC version number*, and selects

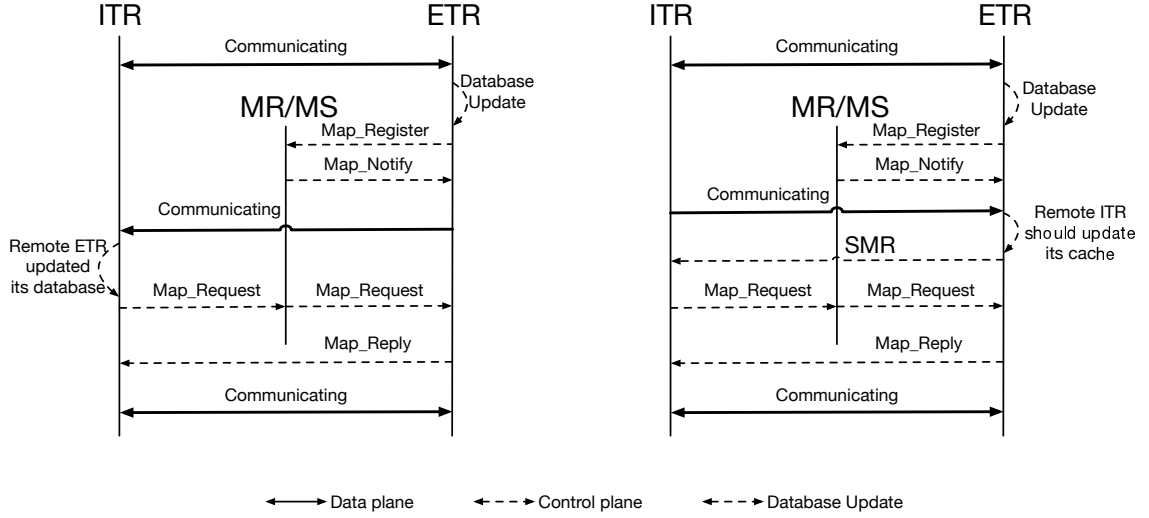


Fig. 2.7 Packet sequence of LISP Map-Versioning mechanism

the version number assigned to the mapping contained in the Cache as *destination RLOC version number*. When an ETR receives such packets and decapsulates them, it makes a comparison between the received version number and the local one. If the source RLOC is higher than the one that ETR stores in its Cache, i.e., the mappings of remote ITR change, so the ETR sends a Map-Request to the MR or directly to the remote ITR to retrieve the latest mapping information (the left figure of Fig. 2.7). If the destination RLOC is lower than the one contained in the Database of ETR, i.e., the local mappings change but the remote site does not know, then the ETR sends a SMR to the remote ITR to let it update the cached mapping (the right figure of Fig. 2.7).

2.4 LISP Mobility

LISP can be implemented on both border routers and terminals. As presented in Sec. 2.1.1, a router supporting LISP called xTR and a mobile host supporting LISP called LISP Mobile Node (LISP-MN), which will be introduced in Sec. 2.4.1. According to the adoption of LISP and host/router supports LISP, mobility can be divided into 4 catalogs: LISP-MN in LISP-Site, LISP-MN in non-LISP-Site, MN in LISP-Site, and MN in non-LISP-Site (shown in Tab. 2.1). However, the case "MN in non-LISP-Site" is the current conventional mobility, which has no relationship with LISP. There are some proposals to solve the seamless mobility issues for the traditional handover, such as Mobile IP (MIP) [98], Mobility Support in IPv6 (MIP6) [97] [86], Multipath TCP (MPTCP) [54] and so on. Thus, we will present

Table 2.1 Catalogs of mobility

Name of catalogs	MN	Site
LISP-MN in LISP-Site	LISP	LISP
LISP-MN in non-LISP-Site	LISP	non-LISP
MN in LISP-Site	non-LISP	LISP
normal mobility	non-LISP	non-LISP

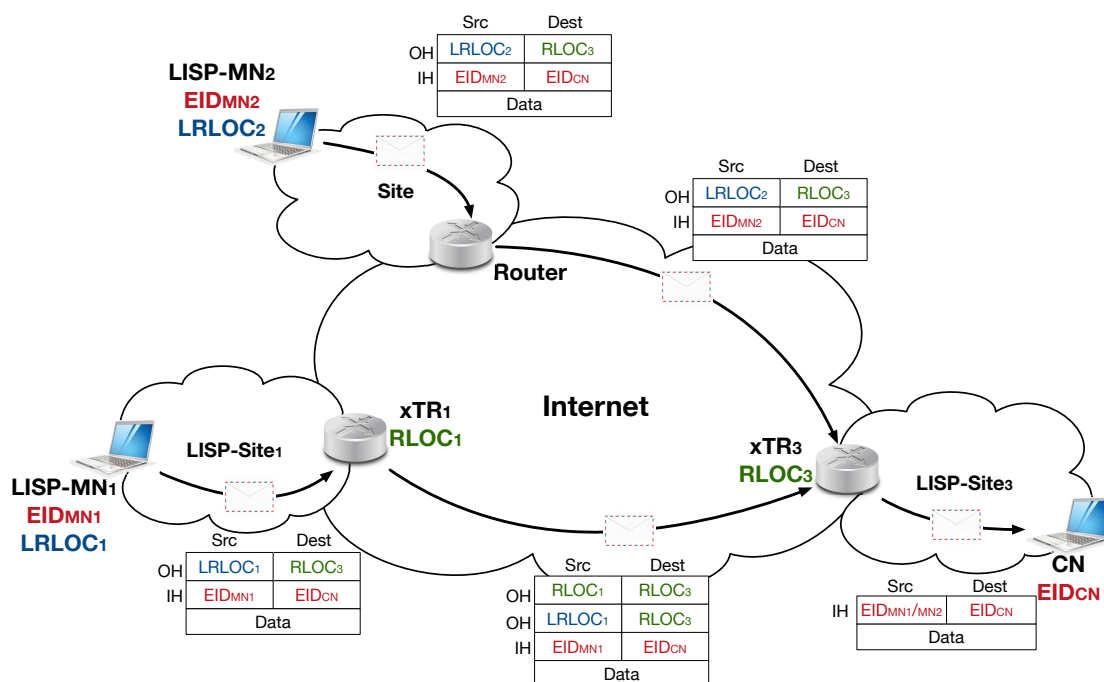


Fig. 2.8 Illustration of LISP-MN in LISP-Site and non-LISP-Site

LISP-MN first and then only describe how the packets exchange when the handover occurs for the first three cases in the following sections.

2.4.1 LISP Mobile Node

This section introduces a mobile LISP-speaking node, called LISP Mobile Node (LISP-MN). It can reside behind both LISP-Site and non-LISP-Site. It is able to change its attachment point during the communication without causing the interrupt of the conversation. The LISP-MN implements a subset of the standard xTR functionality [49] supporting roaming and discovering in an efficient and scalable manner. It can send Map-Request to MR to get the mapping information of the remote host.

When a LISP-MN resides in a LISP-Site, it is assigned an EID taken from the site's EID-prefix as its RLOC, called *Local RLOC (LRLOC)*. Thus, LISP-MN stores not only its permanent unique EID but also the LRLOC and registers this mapping information to the MS. The conventional IP packets produced by the LISP-MN are encapsulated by itself using the permanent EID as the inner source address and LRLOC as the outer source address. The encapsulated LISP packets are forwarded to xTR and encapsulated again by the method of basic encapsulation of xTR.

When a LISP-MN resides in a non-LISP-Site, it is assigned an IP address taken from the site's prefix as its *LRLOC*. LISP-MN also need to store both its permanent unique EID and the LRLOC and register this mapping information to the MS. The conventional IP packets produced by the LISP-MN are encapsulated by itself using the permanent EID as the inner source address and LRLOC as the outer source address. The encapsulated LISP packets are forwarded to the border router and natively forwarded by the method of BGP to the Internet core.

The packet flow sequences of LISP-MN in LISP-Site and non-LISP-Site are shown in Fig. 2.8. Two LISP-MNs, i.e., LISP-MN₁ and LISP-MN₂ send the packets to a remote CN. The only difference is that LISP-MN₁ resides in a LISP-Site, while LISP-MN₂ resides in a non-LISP-Site. The LISP-MN₁ produces the traditional IP packets with EID_{MN1} as source address and EID_{CN} as destination address. Then it encapsulates by adding $LRLOC_1$ as outer source address and $RLOC_3$ as outer destination address by itself after querying the mapping information to MR. The LISP packets with one time encapsulation are forwarded to xTR_1 . The latter gets the mapping information from MR and encapsulates once again the packets by adding the $RLOC_1$ as the outer source address and $RLOC_3$ as the outer destination address, then sends the double LISP encapsulated packets on core Internet. The xTR_3 needs to decapsulate the packets twice after receiving and verifying the packets, and finally sends to CN.

The situation for LISP-MN₂ is simpler. It also produces the traditional IP packets and encapsulates the packets by itself, then sends to its border router. The *Router* does nothing related to LISP but only forwards the already LISP-encapsulated packets to the Internet core in the method of BGP. Thus, the packets are encapsulated only once, whose inner source and destination address is respectively EID_{MN2} and EID_{CN} , outer source and destination address is respectively $LRLOC_2$ and $RLOC_3$.

2.4.2 LISP MN in LISP-Site

This section describes how a LISP-MN seamlessly moves from one LISP-Site to another and the packet flow is shown in Fig. 2.9. At first, packets exchanged between LISP-MN resid-



Fig. 2.9 Packet flow of mobility when LISP-MN in LISP-Site

ing in LISP-Site₁ and CN are double encapsulated LISP packets as mentioned in Sec. 2.4.1. After conducting the handover, LISP-MN still uses its permanent EID_{MN} . However, since it changes the attachment point, xTR_2 distributes a new EID from its EID-prefix to it, and it uses this as its new LRLOC, labeled $LRLOC_2$ in the figure. The LISP-MN should register itself with the new mapping information to the MDS and revoke CN to update its LISP-Cache by SMR or Map-Versioning (introduced in Sec. 2.3 of Chapter. 2). When receiving the packets from LISP-MN, the xTR_2 conducts the second encapsulation and sends the packets to xTR_3 . The benefit of such mobility is that the LISP-MN can roam through the different subnets without interrupting the communication with the remote CN, because LISP-MN never changes its EID. However, to conduct double encapsulation performs higher delay of handover.

2.4.3 LISP MN in non-LISP-Site

The Fig. 2.10 illustrates the packet flow of the LISP-MN roaming between two non-LISP-Sites. The Site₁ distributes an IP address from its prefix to LISP-MN, and the later uses this as its LRLOC, labeled as $LRLOC_1$ in the figure. The LISP-MN encapsulates the packets by itself and sends to Router₁. Router₁ just natively forwards the packets to xTR_3 by

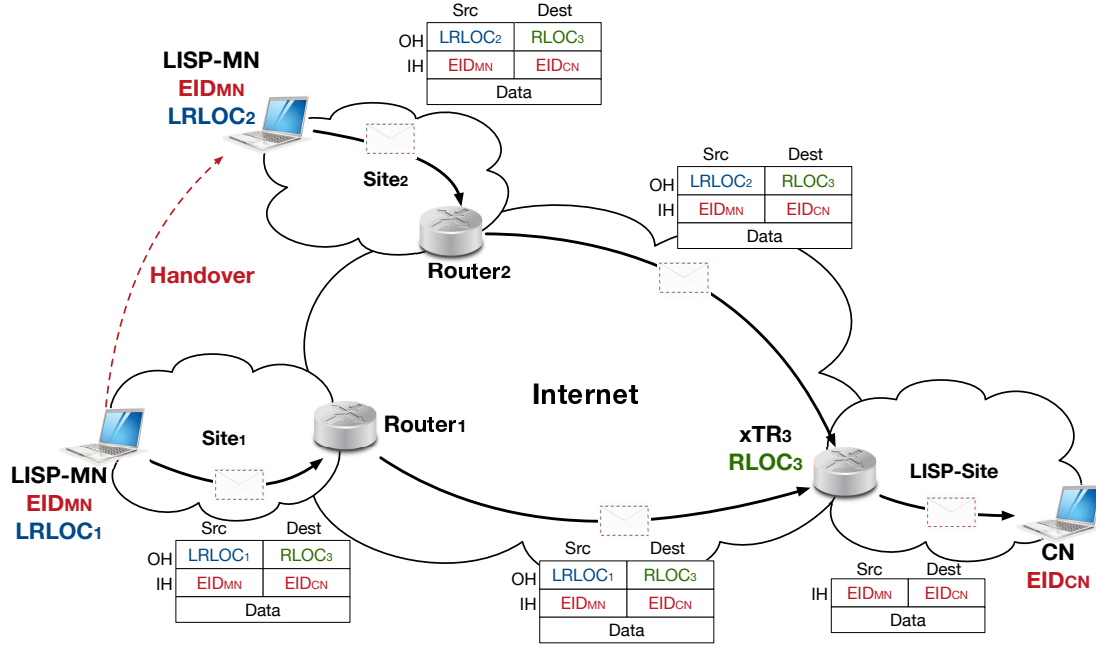


Fig. 2.10 Packet flow of mobility when LISP-MN in non-LISP-Site

using BGP. When LISP-MN moves to another site labeled *Site₂*, it receives an IP address from a different prefix and uses it as the new LRLOC. It registers itself to the MDS with the EID_{MN} -to- $LRLOC_2$, and informs the remote *CN* to update the mapping information by sending SMR or using Map-Versioning. In this scenario, LISP-MN can also permit to seamlessly roam through the different subnets, but it can also reduce the high latency of double encapsulation during the handover.

2.4.4 MN in LISP-Site

When the MN is normal host without any LISP function, while the border router supports LISP, i.e., xTR, the packets are exchanged with the remote *CN* as the basic communication between two LISP-Sites introduced in Sec. 2.1.3 of Chapter. 2. The EID of MN is distributed from the EID-prefix of the LISP-Site, if MN moves to a new LISP-Site but keeps the original EID, it can not exchange the packets with its new xTR, since its EID is not belongs to the EID-prefix of the new LISP-Site. Thus, in this case, the mobility of MN can only conduct within a subnet instead of roaming inter-domains.

2.5 LISP use-cases

Even though LISP is proposed to solve the Internet scalability issues, it also provides the advantages to the other uses, such as: traffic engineering, transition between IPv4 and IPv6, device's mobility at client-side, virtual machine's mobility in data center, LISP SDN, and so on [40].

2.5.1 LISP traffic engineering

For every mapping information, an EID-prefix associates to a list of RLOC tuples $\langle \text{RLOC}, \text{Priority}, \text{Weight} \rangle$, where each RLOC is annotated with a priority and a weight. In the case of multiple RLOCs existing, the ITR selects the one with the highest priority and sends all the LISP encapsulated packet to this RLOC. If several RLOCs all have the highest priority, the traffic is balanced proportionally according to their weight among such RLOCs to route to the multiple ETRs. Traffic engineering in LISP thus helps the mapping owner decide the primary and backup path for its incoming and outgoing packet use, and also allows it to control the traffic load among its links. For example, the received Map-Reply shows that an EID-prefix is associated to $\langle \text{RLOC_1}, 10, 50 \rangle$, $\langle \text{RLOC_2}, 10, 20 \rangle$, $\langle \text{RLOC_3}, 10, 20 \rangle$, $\langle \text{RLOC_4}, 10, 10 \rangle$, then the RLOC_1 gets 50% of the traffic, RLOC_2 and RLOC_3 are respectively assigned to 20% of the traffic, and the last 10% of the traffic are balanced by RLOC_4.

Traffic engineering in LISP can achieve one step further than nowadays BGP. Since every Map-Request contains the source EID of the packet, for which causes the ITR mapping miss and triggers the Map-Request. It is possible that a mapping owner (ETR) gives the different Map-Replies according to the sender of Map-Requests. Thus, the traffic engineering policy may be different to different ITRs although they query the same ETR for the same EID-prefix. This functionality is not available today with BGP, because a domain cannot control exactly the incoming routes from the other domains if they are not direct neighbor.

2.5.2 LISP transition between IPv4 and IPv6

The LISP encapsulation mechanism is designed to support any combination of address families for locators and identifiers, such as: IPv4, IPv6, MAC address, or some other arbitrary elements. It is then possible to bind IPv6 EIDs with IPv4 RLOCs and vice versa. This allows transporting IPv6 packets over an IPv4 network (or IPv4 packets over an IPv6 network) by deploying the xTRs at the border of heterogeneous networks, making LISP a valuable mechanism to ease the transition to IPv6.




Fig. 2.11 Illustration of LISP transmission between IPv4 and IPv6

Examples are as shown in Fig. 2.11, where IPv6 islands are connected via IPv4 infrastructure (on the left) and IPv4 islands are connected through IPv6 infrastructure (on the right). As LISP provides mechanisms for encapsulating IPv6 host packets using IPv4 RLOCs on the xTRs, LISP can then provide IPv6 communication between hosts/servers even though the intermediate network only supports IPv4. Similarly, xTR can encapsulate the IPv4 packets by adding the IPv6 RLOCs to achieve the IPv4 access through the IPv6 infrastructure. In addition, when the LISP infrastructure includes PxTRs, it can provide both IPv4 and IPv6 connectivity to non-LISP IPv4/IPv6 sites (legacy Internet). In summary, it can function in any of four following ways:

- IPv4-to-IPv4 over IPv4
- IPv4-to-IPv4 over IPv6
- IPv6-to-IPv6 over IPv4
- IPv6-to-IPv6 over IPv6

However, LISP does not translate between the IPv4 and IPv6 EIDs.

2.5.3 Mobility at client-side

Nowadays, the explosion of mobile devices that require the massive mobile Internet traffic results in the evolution of mobile network architecture. Keeping the communication during the roaming without any interruption becomes a key point. Thanks to LISP separating the identifier of a host and the locator of the attachment point, the connection between two hosts can be kept alive while roaming. Because the EID, which is used as the inner address, is permanent, whereas only RLOC (used as the outer address) is changed.

By using LISP mobility, it is possible 1) to allow TCP connections to stay alive while roaming; 2) to provide the shortest bidirectional data paths between a Mobile Node (MN) and Correspondent Node (CN); 3) to allow MN to communicate with another MN when

both are roaming; 4) not require fine-grained routes in the core network, nor the home-agent, foreign agent or other data plane network elements to support mobility; 5) there is no triangle routing of data packets as is found in Mobile IP [96]; 6) and there is no IPv6 new extension headers to avoid triangle routing [71]. LISP Mobile Node (LISP-MN) takes advantage of the LISP infrastructure so to overcome the limits imposed by Mobile IP. LISP mobility can be applied in 5 scenarios and the implementation of LISP-MN is achieved by OOR (Open Overlay Router) [23], which will be introduced Sec. 3.3.1 of Chapter. 3 in details:

1. LISP-MN to a LISP Stationary Node (LISP-SN);
2. LISP-MN to a Non-LISP-SN;
3. LISP-MN to another LISP-MN;
4. Non-LISP Site to a LISP-MN;
5. LISP Site to LISP-MN.

2.5.4 Mobility in Data Center

As the server virtualization and high availability across geographically dispersed data centers are more and more common, similar to the realization of LISP mobility at client-side, the migration of Virtual Machine (VM) within a data center or from a data center to another without dropping the running connections also becomes a potential use case of LISP. Two mechanisms at the state of the art can perform this operation. One is the aforementioned host-based LISP implementation called OOR, another method to handle VM mobility via LISP is actually implemented in some Cisco products, only partially documented in [20].

2.5.5 LISP SDN

As described before, LISP is inherently decoupling Control Plane from Data Plane. LISP moves all control functions onto Mapping System, while keeps data functions at the xTR level. The feature of LISP corresponds to the characteristic of Software Defined Network (SDN), which also decouples Network Control Plane from Data Plane, increases flexibility and development speed of features and functionalists. Thus, this decoupling entitles network operators to build a Software Defined Network (SDN) on top of LISP [104].

The OpenDaylight Project [24] is a collaborative open source project hosted by The Linux Foundation to promote Software Defined Network (SDN) and Network Functions Virtualization (NFV). It implements the LISP Control Plane functions as shown in Fig 2.12. The data plane elements consist of virtual switches and physical device interfaces, at the aim of constructing the LISP tunnel to transfer the packets. The LISP Plugin is the South-

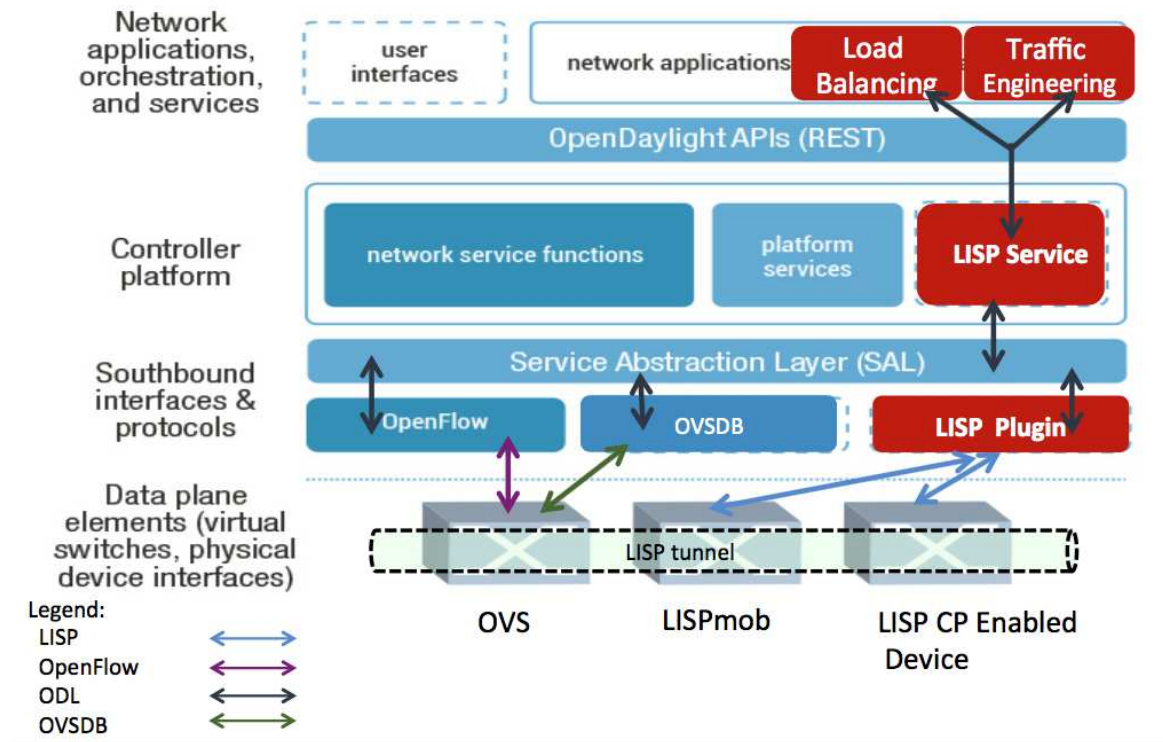


Fig. 2.12 LISP in OpenDaylight (source: [14])

bound interface, which is used to look up LISP policies from LISP Service by sending Map-Request and receiving Map-Reply, and create the LISP tunnel over LISP overlay across the different kinds of data plane devices. The LISP Service is actually implemented as LISP Mapping System to store the various user-defined policies. The only difference is that in OpenDaylight, it is the MS itself that replies to the Map-Requests, instead of forwarding them to any xTR. It is like every mapping has the "proxy Map-Reply" flag (P-bit) set to response on behalf of xTRs. The OpenDaylight APIs (REST) are used to set the policies, such as: the policies for Load Balancing and Traffic Engineering, obtained from the north bound into LISP Service. The policies are defined as the network applications or services by the users.

Network functions such as subscriber-management, content-optimization, security and quality of service, are typically delivered using proprietary hardware appliances embedded into the network [36]. But the next generation network functions are being implemented as pure software instances running on standard servers. They are unbundled virtualized components of capacity and functionality. LISP-SDN based flow-mapping, dynamically assembles these components to whole solutions by steering the proper traffic in the right sequence to the correct virtual function instance.

Chapter 3

Related Work

This chapter first presents the current LISP status, including: two available LISP platforms, the only LISP monitor, two LISP manual management tools, some private and open-source implementations, and the latest research achievements, from the LISP network evolution, the LISP mapping system assessment, the evaluation of interworking between LISP-site and legacy Internet, to the analysis of LISP mobility. Besides introducing the existing works, we also find some limitations in each aspect. Thus, we point out the missing works in the LISP monitor and implementation, as well as in the various sides of LISP performance evaluation. Most of them we have solved. The works that we conducted to complete LISP studies are presented in the following chapters, i.e., from Chapter. 4 to Chapter. 7. Further, we list some other LISP challenges and issues at the end of this chapter. Although these fields are not considered in this dissertation, they are also important to help researchers improve LISP technology.

3.1 LISP platforms

This dissertation is focus on the evaluation of LISP performance in the realistic environment, thus the large scale experiments are the basic and necessary methods. Since no matter the previous studies or the works described in this dissertation leverage on the experimental LISP networks, we introduce first two existing LISP platforms: LISP Beta Network and LISP-Lab platform, before discussing the current research results.

3.1.1 LISP Beta Network

The first global LISP testbed called LISP Beta Network [11] [44], has been deployed since 2008. Started as the only LISP testbed, the LISP Beta network had a steady growth and




Fig. 3.1 International LISP Beta Network High Level Topology – October 2017 [12]

achieved world-wide experimental deployment [106] with the purpose to gain real-life experience on LISP. Initiated by Cisco, the members of the LISP Beta Network are not only major companies and operators, but also academics, research laboratories, and startups offering LISP related services. Participants of this network are located in 34 different countries, with the highest concentration in Europe and North America.

The LISP Beta Network is composed of the mapping system, LISP routers (i.e., xTRs), mobile nodes, Re-encapsulating Tunnel Routers (RTR [46]) and Proxy Tunnel Routers (PxTR [76]). On March 14th 2012, the mapping system used by the LISP Beta Network has switched from LISP+ALT [57] to the more flexible LISP-DDT [58]. Changing the mapping system has considerably reduced the configuration and maintenance burden, while a slight performance degradation has been observed [106]. According to the latest architecture (October 2017) [12] shown in Fig. 3.1, LISP Beta Network has 3 DDT Root, 2 DDT TLD, 7 available MR/MsEs and 6 PxTRs. More precisely: 2 MR/MsEs and 2 PxTRs are in US, 3 MR/MsEs and 3 PxTRs are in Europe, 2 MR/MsEs and 1 PxTR are in Asia. It mainly uses two EID addressing spaces, which belongs to Cisco: 153.16.0.0/16 for IPv4 and 2610:00D0::/32 for IPv6 [106], but there are also some other EIDs beyond them.

3.1.2 LISP-Lab platform

Due to a monolithic control by one single commercial actor, however, LISP Beta network has limitations to explore the innovative services. Thus, some other projects, such as ANR LISP-Lab project [15], are under implementation to develop new enhanced features.

LISP-Lab project aims at building an open platform, based on the LISP architecture, providing the environment to perform high quality research and support the design, development, and thorough assessment of new services and use-cases. The range of technical tasks planned in the LISP-Lab project, from cloud networking, to access technology, through inter-domain connectivity, traffic engineering, and mapping management, boosting innovation beyond the LISP technology itself. LISP-Lab platform is solely deployed by the open source LISP implementation, using OpenLISP [25], opened to the external users from 2015. It is coordinated by a French consortium: two academic institutions (UPMC, TPT), two Cloud Networking SME (Alphalink, NSS), two network operators (Renater, Orange), two SMEs on Access/Edge Networking (Border 6, Ucopia) and one Internet eXchange Point (Rezopole). Ten international external partners also joined the platform. At the moment of writing, LISP-Lab consists of 3 MRs/MSeS and 2 PxTRs located in France, as well as 13 worldwide xTRs, and has already been interconnected to the LISP Beta Network with an Open-LISP DDT root [59].

3.2 LISP monitoring tools

As the researchers need leverage LISP tools to conduct the experiments on the aforementioned LISP platforms, this section introduces the currently existing tools supervising LISP status or troubleshooting it for some specific purposes. We will explain their working mechanisms and figure out their shortcomings. Also, we will state the necessity and urgency to propose a new comprehensive one.

3.2.1 LISPmon

LISPmon is the only monitoring platform that supervises the current public LISP status, i.e., the mappings between the EID-prefixes and the RLOCs. This platform is developed by the Advanced Broadband Communications Center of UPC (Universitat Politècnica de Catalunya). It scans the whole IPv4 addressing space everyday, normally beginning at 7:00 a.m. (UTC), and queries them by sending the Map-Request to one operator preferred MR of the LISP Beta network (i.e., 7 MRs are candidates to be selected). If this MR does not respond, the operator will choose another one as a replacement, and start the queries

again, to guarantee that the reception of LISP Map-Replies changes smoothly between two consecutive days. The available EID-RLOC mapping information is published once per day, ever since the beginning of 2010.

3.2.2 LIG

LISP Internet Groper (LIG) [52] is a LISP-specific manual management tool, used to obtain the mapping information for a certain EID by querying the LISP mapping system. It can be run by all devices that implement LISP, including xTR, PxTR, MR, MS, as well as by a host system at either a LISP-capable or non-LISP-capable site. A possible syntax for a LIG command could be:

$$lig < destination > [source < source >][to < map - resolver >]$$

Where *<destination>*: is either a Fully Qualified Domain Name (FQDN) or a destination EID for a remote LISP site; *source <source>*: is an optional source EID to be inserted in the 'Source EID' field of the Map-Request; *to <map-resolver>*: is an optional FQDN or RLOC address for a MR. When LIG is launched, a Map-Request is sent for a destination IP address, and the returned Map-Reply is displayed as follows:

- The EID-prefix for the site that the queried destination EID matches.
- The locator address of the responder.
- The RLOC tuples *<RLOC, Priority, Weight>*.
- A round-trip-time (RTT) estimate for the Map-Request/Map-Reply exchange.

The remaining work for LIG is about processing of Negative Map-Replies: LIG should be able to clearly indicate how and why a Negative Map-Reply is received. Negative Map-Replies could be sent in the following cases: the LIG request was initiated for a non-EID address or there was rate-limiting on the responder.

3.2.3 RIG

LISP-DDT Referral Internet Groper (RIG) [48] is also a LISP-specific manual management tool, but it is not like LIG, which can be served on top of any MDS to retrieve RLOCs used for packet encapsulation. RIG is used to recursively find all the MSes serving an EID-prefix, specifically within a LISP-DDT mapping database framework. It performs the same operation as that of a MR. More specifically, when RIG command is run, a Map-Request is sent into the DDT hierarchy for a destination EID, and is forward based on the EID-prefix and delegation referrals contained in the Map-Referral messages until the DDT MS

containing the mapping information is reached. It can be run on the same devices as LIG. A simple syntax for a RIG command could be:

$$\text{rig} < \text{eid} > \text{to} < \text{ddt} - \text{node} > [\text{follow} - \text{all} - \text{referrals}]$$

Where $<\text{eid}>$ is the destination EID being queried; $<\text{ddt-node}>$ is the RLOC of any DDT node in the DDT hierarchy; when the keyword [follow-all-referrals] is used, all the referral RLOCs will be queried; if this keyword is not used, one of the referral RLOCs will be selected to descend a branch of the DDT hierarchy. The results of returned Map-Referral messages mainly contain the RLOCs of child nodes storing the destination EID mapping information, as well as RTT indicating the time spends for getting the referral DDT nodes.

3.2.4 Missing work

LIG and RIG are the manual management tools, it is used to troubleshoot for some specific purposes by the experimenters rather than monitor the LISP platforms. LISPmon as the first and only LISP monitor satisfies the preliminary LISP requirements at the initial stage, but faces some limitations as the quick development of LISP. LISPmon just publishes the mapping information as the daily once per day. All the published results are queried to one MR of LISP Beta Network and captured only by one VP. Since LISP Beta Network has multiple MRs over the world and join of LISP-Lab platform later, the current monitoring mechanism of LISPmon shows the incompleteness. To move LISP forward, it is necessary to fully understand the performance of every LISP network entity. Beyond only publishing the mapping information between EID-prefix and its RLOCs, automatically conducting the comprehensive experiments from the various aspects will help the LISP experimenters improve their research speed and enrich their study fields. Thus, we filled this gap by proposing a new LISP monitor at large scale, which will be explained in Chapter. 5.

3.3 LISP implementations

The most widely used are the implementations on Cisco routers and on AVM's FRITZ!Box home routers [5]. Both can be used to test LISP and connected with the LISP Beta Network. However, the source codes of their implementations are not available and thus cannot be extended for LISP research purposes. Thus, in this section, we focus on other LISP implementations, which can help the researchers verify their improvements so to move forward LISP technology. Among these implementations, several are proprietary and some are open-source. At end, we analyze their advantages but also provide the limitations.

3.3.1 Existing implementations

OpenLISP

OpenLISP [25] [100] is an open source implementation of LISP protocol. The LISP Data Plane, such as: LISP-Cache and LISP-Database, as well as the encapsulation/decapsulation functions are implemented in the kernel of FreeBSD [7] Operating System. LISP Control Plane is written in C in the user space of FreeBSD and Linux. Everything related to the Control Plane is meant to run in the user space. OpenLISP does not provide any specific control plane protocol implementation associated to mapping distribution protocol. Instead, it implements a new type of sockets, called the Mapping Sockets providing an API that can be used by any users space process. OpenLISP was implemented by researchers from Universite catholique de Louvain and T-Labs/TU Berlin. At the moment of writing, Pierre et Marie Curie University maintains the Lip6-lisp [10], which is an extended implementation of OpenLISP and it supports fully featured LISP Control-Plane.

OOD

Open Overlay Router (OOD) [23], the successor of LISPmob [41], is an open source LISP and LISP mobile node implementation written in C for Linux, Android (only on rooted devices) and OpenWrt that uses LISP or VXLAN [82]/GRE [63] as an overlay for SDN. LISPmob was initially developed by Cisco. By interoperating with LISP Beta Network, LISPmob enables mobile hosts to change its network attachment point without losing connectivity, while maintaining the same IP addresses. Now OOD is maintained by Barcelona Tech University. It focus on supporting for network function virtualization (NFV), e.g., by SFC, and integration with OpenDaylight [24]. The extended feature set makes OOD attractive for application.

PyLisp

PyLisp [28] is a small but functional implementation of LISP in Python. There is a packet of PyLisp implementation on Github, called pylisp [27] provides the means to parse and create LISP data and control messages, as well as command line tools to perform actions like asking a map resolver for a mapping. The functions of PyLisp are not complete, so the intention to the later stage is the implementations for map server, map resolver and DDT node.

jLISP

jLISP [111], proposed in 2016, is an open source implementation of LISP in Java and runs in the user space. It is platform-independent and object-oriented. jLISP is compatible with the LISP RFCs and is modular, which facilitates to explicit programs for every LISP component, such as: xTRs, RTRs, NTRs, and the mapping-system.

LISP on OMNeT++

OMNeT++ [22] is an extensible, modular, component-based C++ simulation library and framework, primarily for building network simulators. A LISP simulator is proposed in [74], which integrates LISP into the INET framework [8] for OMNeT++. This simulator implements basic LISP functions, interworking mechanism and LISP-MN architecture. In addition, their work also support NAT traversal. Since this simulation model is proposed in 2012, the mapping system is LISP-ALT instead of nowadays LISP-DDT. In addition, the simulator OMNeT++ is open source, but this LISP extension is not. This would be an obstacle for this implementation that is widely used now.

Lispers.net

Lispers.net [17] is a Python implementation aiming to implement the complete feature set of the open IETF LISP protocols and interoperate with all modern proprietary and open source implementations available today. It can not only run on Ubuntu, Fedora, CentOS, and Debian Linux distros as well as MacOS and Raspbian, but also in a docker container. The implementation supports the Data-plane components: ITR, ETR, and RTR, and the Control-plane components: MR, MS, and DDT-Node for 10 different EID address types and 7 different RLOC address types. The first version was released in November 2013. Additional non-LISP-specific functions and behaviors are also added, which make lispers.net a general overlay controller. Until December 27th 2017, the latest version of lispers.net is Beta Release Version 0.430. Although the goal of Lispers.net is the world's richest feature set of LISP, it is not open source.

LISP controller in ONOS

The LISP controller in Open Network Operating System (ONOS), which is an open source SDN OS for service providers [62], is one of South Bound Interfaces (SBI) for SDN controllers. The LISP controller in ONOS is proposed to address the OpenFlow issues, such as: scalability, performance, heterogeneity, inter-operability with legacy networks and protocols. An open source SDN controller development project, called OpenDaylight [24]

sub-project, is developing the LISP mapping system in SDN environments. The LISP controller centralizes LISP controls into a single entity. It supports LISP as an SDN SBI to work with other protocols such as OpenFlow, and provides APIs for LISP applications. The LISP controller brings many benefits, such as: simplified management, a global view of LISP networks, short information convergence time, and easy development support of LISP based applications.

3.3.2 Missing work

At the moment, LISP has some implementations for the different Operation Systems, including either the proprietary source or the open source. However, researchers mostly prefer to rely on simulation than real testbeds, because the simulation environment simplifies the studied problem and allows researchers to concentrate on the most critical issues [88]. So, simulation has a dominant capability over the other evaluation tools, when the researchers want to quickly test new features on a larger scale for different network scenarios. Among the recent network simulators, such as ns-2 [38], ns-3 [64], OMNET++ and OPNET [42], basic LISP features has already been extended on OMNET++. However, it is not open source, which poses the issues for the researchers to easily and flexibly conduct the simulations. In addition, ns-3 is more popular with the researchers than the others [102] and there is an open source project¹ implementing the fundamental LISP on it in 2016. Whereas this project is based on ns-3.24 (the version at that moment), and it does not support features related to LISP mobility. Thus, we decide to leverage on this project, but first adapt to ns-3.26, and realize the extension of LISP mobility. Such contribution will be introduced in Chapter. 7.

3.4 LISP mapping system evaluation

3.4.1 Current studies

The performance of LISP Mapping System is evaluated in [106] and [44] from different aspects. As the LISP Beta Network changed its mapping system from LISP+ALT to LISP-DDT on March 14th 2012, the authors measure the delay of resolving mapping (the latency between sending the Map-Request and receiving the Map-Reply) during this transition period. They find the delay to get the mapping by using LISP+ALT is smaller and more stable than LISP-DDT. However, the manageability of the later is better than the former, since

¹Source code available at:
https://bitbucket.org/Lionel_Agbodjan/tfe__towards_a_lisp_simulator

by using LISP+ALT, the Map-Requests normally need traverse the entire BGP overlay and maintaining the ALT also causes many cumbersome overheads especially as the network growing. Although using LISP-DDT slightly degrades the resolving mapping performance, but the delays are rather small and the mapping system is able to cope with a significant growth of the addressing space. The resolution delay is slightly various from the vantage points and the MRs, but the European infrastructure is unreliable due to the loss of many queries.

3.4.2 Missing work

The previous measurements of the MDS are mainly depended on the delay, including the different resolving latency that each MR needs to use, and comparing which VPs have longer delay to query the mappings. However, the authors have never taken look at the contents of mapping information, i.e., the type of mappings and locators etc.. In addition, there is no work to compare whether the consecutively received Map-Replies change over time. If yes, how they change? If there are any patterns to follow? There are also no comparisons on whether the Map-Replies sent by the different MRs for the same EIDs are identical. Nor to the Map-Replies received by the various VPs. All these evaluations will be discussed in Chapter. 4.

3.5 LISP network evolution

3.5.1 Current studies

After deploying LISP in LISP Beta Network from 2008, more and more participants attended to the development of LISP. The LISP network is steady flourishing. Thus, the paper [106] looks at how the LISP network evolves. The authors observe that the number of LISP mapping was as low as 20 in January 2010, but it has consistently grown to reach 80 mappings in 2012. The number of LISP mappings is four times to the initial number, which indicates that LISP network has a regular growth. Within two months of the second part of 2011, the number of negative mappings was doubled, i.e., from about 150 to about 300. The authors suppose that is due to a more fragmented EID-space. The paper also observes some malformed mappings, which belongs to the LISP mappings but have an empty RLOC set. However, this kind of mappings is never more than 0.5% of the total number of mappings within one day. Further, the paper presents that from January 2010 to May 2012, the percentage of mappings using two or less RLOCs has increased. Moreover, there exist mappings with four RLOCs in 2012, whereas they were not present in January 2010.

3.5.2 Missing work

Previous works respectively evaluate the LISP network evolution in 2010 and 2012. However later, LISP mapping system is changed from LISP+ALT to LISP-DDT in 2013. Besides, as mentioned in Sec. 3.1.2 of Chapter. 3, LISP-Lab platform is open since 2015. It brings new LISP-sites over the world and interconnects with LISP Beta Network. Further, LISP Beta Network itself also updated their xTRs and changed its network architecture in 2016. Thus, the number of LISP mappings, the fragment of EID-space, the distributions of RLOCs are also changed over time. It is interesting to have a look at them in order to know how the LISP network evolves in recent years. This work will be mentioned in Chapter. 5.

3.6 LISP interworking with legacy Internet

3.6.1 Impact on traffic between LISP-site and the legacy Internet

LISP has no impact on traffic to nowadays legacy Internet if neither source nor destination is LISP. However, it has significant impact on the traffic between LISP-site and non-LISP-site. First, to support the interworking between LISP-site and the legacy Internet, the PxTRs should be set up. The configuration on PxTR does not cause any special technical issue, but using PxTR potentially causes path stretch between source and destination. It makes the packet exchange experience a larger delay and potentially leads to higher packet loss rate. Thus, placing the PxTR near the source/destination of traffic allows the communication between the non-LISP site and the LISP site to have the least path stretch (i.e., the least number of forwarding hops when compared to an optimal path between the sites) [76].

When deploying the PxTRs in the experimental infrastructure, it is better to think about first where to deploy the PxTRs, since the placement in topology has a significant impact on the path stretch. Besides, as the number of PxTRs has a direct impact on the traffic load and the recovery success in case of the failure of a PxTR, we need consider how many PxTRs are needed. Further, we need to think if it is necessary that all the PxTRs advertise for the whole EID space, or it only needs that each of them just announces for the different segmented EID-prefixes.

3.6.2 Current studies

As more and more deployments of LISP, it becomes necessary that the LISP networks communicate with the legacy Internet. The evaluation of the interworking performance is presented in reference [44]. It chooses 200 VPs on the PlanetLab (a global research network

that supports the development of new network services from 2003) [26] located in North America, Asia and Europe by the order of diversity of number of VPs. It also selects 116 EIDs, which have at least 1 corresponding RLOC in the mapping from the daily of LISPmon project [19] published on May 17th 2013. The experiment is conducted from November 4th to 7th 2013. The authors find that the interworking generally has a negative impact compared to without using LISP. The results show that there are 70% of the EID-prefixes have at least 20% of increase in the delay (i.e., Round-Trip Time (RTT)). One PxTR behaves the worst that 95% of the EID-prefixes using it suffer from a delay increase. Based on BGP, the selection of PITRs is driven by the economical relationship between autonomous systems, instead of their geographical positions. As a result, even if the LISP Beta Network has 9 available PxTRs at the moment, there are only 4 are used in the experiment. Further, the most VPs use PxTRs in Asia although they are located in Europe or America. The most popular PxTR is selected 79% of times.

3.6.3 Missing work

The only evaluation of PxTR used for the interworking between LISP and legacy Internet was conducted in 2013. At that time, the unique LISP platform was the LISP Beta Network. Thus, the measurement was only based on its PxTRs and the VPs on the research network PlanetLab. After the open of LISP-Lab platform in 2015, there is no evaluation about the PxTR anymore. We do not know the recent performance of PxTR on LISP Beta Network and we have no knowledge about if the PxTR of LISP-Lab functions well. Further, the previous assessment was conducted within an experimental/research environment, there is no experience about the interworking performance in the realistic network, especially for the interoperability with the current popular websites on the world. In the Chapter. 6, we will explore all these issues of LISP interworking performance.

3.7 LISP mobility

3.7.1 Current studies

To support IP mobility, the conventional LISP should be extended. The first mobility-compatible LISP extension is LISP-MN [49][92]. LISP-MN is an end-host based IP mobility protocol which requires updating the software of the mobile terminal. In other words, the end-host that uses LISP-MN can be regarded as a small LISP-site.

The main drawbacks of LISP-MN are the double encapsulation and the need to modify software in end-host. Double encapsulation increases mapping latency, overhead and MTU

issues. The modification at end-host hinders the deployment of LISP-MN. To overcome these drawbacks, lots of research efforts are made by LISP community.

The performance of LISP-MN is investigated in a qualified way in [85]. The authors illustrate the performance degradation due to double encapsulation and the triangle route problem under some conditions. They also propose mechanisms such as location-aware LISP mobility node, local mapping systems to improve the performance of LISP-MN. A simple model that help analyze the handoff procedure in IP mobility is proposed by [99]. Besides LISP-MN, the proposed model can be also used for the handoff analysis for other IP mobility protocols such as MIPv6.

To avoid the modification to end-host required in LISP-MN, a network-based and LISP-based IP mobility protocol: LISP-ROAM is proposed in [60]. Their solution is that the network assigns the same IP address regardless of their network attachment point under the cooperation between some additional/modified network components (DHCP servers, authentication services, LISP xTR and servers). The authors implement LISP-ROAM on top of LISPmob and give an experimental evaluation of the performance of LISP-ROAM. As an effort to avoid double encapsulation in LISP-MN, LISP-NEMO is proposed in [115]. The basic idea of LISP-NEMO originates from NEMO [70]. It introduces a new network element called mobile router (MR) that manages and represents the mobile network. To mitigate the problems in LISP-NEMO and LISP-ROAM, LISP-HNM is proposed in [112]. The authors introduce an access register protocol and an active notification scheme to conventional LISP so that the latter can support both end-host and network mobility.

The packet loss due to IP mobility is also a concern. To reduce the packet loss, Isah et al. [68] propose an improvement to LISP-MN. The basic idea is that those packets generated during handover procedure are buffered to servers close to the MN's new location and forwarding them to the MN on handover completion.

3.7.2 Missing work

To achieve the seamless handover, LISP can be implemented on the border router or directly on the host to support the mobility. However, these proposals are only limited at the theoretical level. There are no experimental results show which solution is better based on some metrics, such as: the smaller handover delay or the less packet loss. No evaluations respectively compare the advantages and the disadvantages of each solution. What's more, no matter which solution is adopted to conduct the mobility, SMR and Map-Versioning are the candidates to revoke the remote ITR to update its mapping cache. Thus, it is interesting to make a comparison between these two mechanisms, for example modeling the overhead

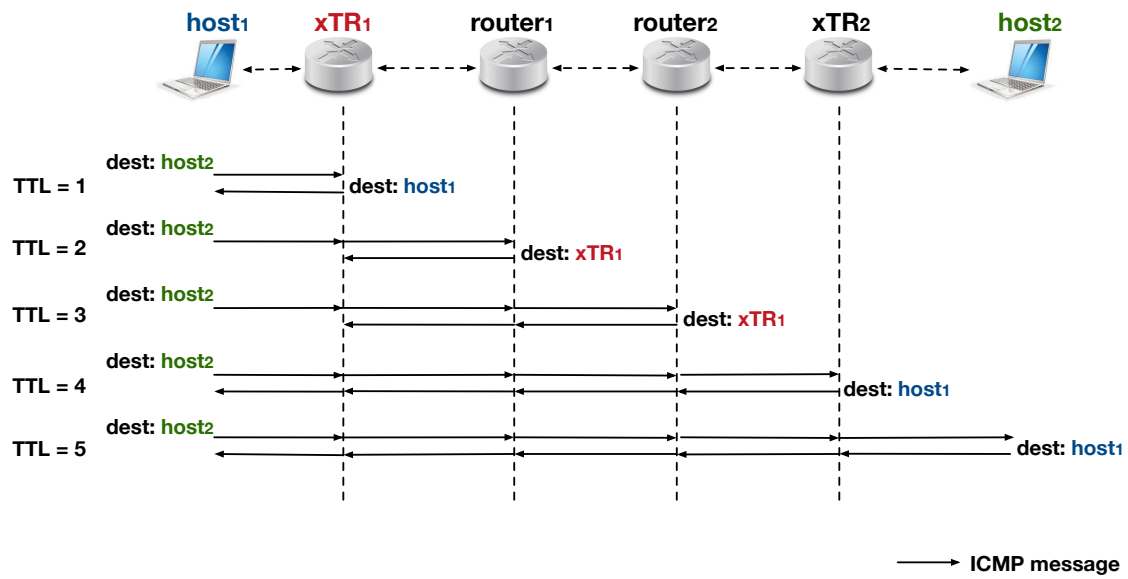


Fig. 3.2 Illustration of LISP troubleshooting

traffic according to the number of MNs and CNs of each mechanism. The Chapter. 7 will analyze some scenario of LISP mobility.

3.8 Troubleshooting problem

A big issue in the LISP experimentation is the difficulty of troubleshooting. The most frequent used troubleshooting tool is *traceroute*, but it can not provide the exact path that the packets pass between the xTRs or PxTRs. The experimenters only know how many hops between them. When the host traceroutes a destination, it uses each EID as the source/destination IP address of ICMP. However, when the packets pass to the ITR, it encapsulates the packets, so the following routers think that the ITR is the source of ICMP and will give back the ICMP message to the ITR. In this way, the host, which is the real source launching the ICMP receives no response. Until the packets reach to ETR or PETR, they decapsulate the packets to know the real source of ICMP and answer to the EID of host.

An example of LISP troubleshooting is shown in Fig. 3.2. When *host1* traceroutes *host2*, *host1* puts its own EID and the one of *host2* respectively as source and destination address. The first ICMP message with TTL (time-to-live)=1 expires at *xTR1*, then it sends back the ICMP Time Exceeded message to *host1*. Thus, the *host1* knows that its next hop is *xTR1*. However, when *xTR1* receives the ICMP message with TTL=2, it decreases the TTL value, and encapsulates the ICMP message with its own RLOC as source address and forwards to

the next hop $router_1$. Since the ICMP message that the $router_1$ receives showing xTR_1 is the source instead of $host_1$, the $router_1$ drops the packets and replies back to xTR_1 . In this way, the $host_1$ cannot receive any ICMP Time Exceeded messages until the xTR_2 is tracerouted. When the xTR_2 receives the ICMP message, it denpsulates the packets and finds that the real traceroute sender is $host_1$. Thus the xTR_2 replies to $host_1$. Same to $host_2$, the received ICMP message shows that $host_1$ is the traceroute origin and itself is the destination, so it returns an ICMP Echo Reply message $host_1$.

To the conventional Internet topology, *traceroute* uses the returned ICMP Time Exceeded messages to build a list of routers that packets traverse until the destination is reached. However, due to the troubleshooting problem LISP facing to, the list of routers between xTRs is hidden to the traceroute sender. It can only know how many hops it needs to arrive the ETR. Thus, when there is a problem in the LISP network, it is hard for the operators to pinpoint the reason as they have only a partial view of the network. The experimenters can leverage on the LISP Cache/Database or querying the MRs to observe the potential problem in the network, but the knowledge remains limited.

This issue brings inconveniences to our experiment, which will be discussed in Chapter. 6. In that experiment, we can only find how many hops between our VPs and the PxTRs, but are not able to specify the exact routing paths.

3.9 Other LISP Related studies

LISP is still in infancy. To completely move it forward, there are many other aspects needed to be taken into consideration, such as: the caching, the LISP security, the MTU issue, and so on. As they are also important and necessary to LISP, we list them here. However, since they are not related to this dissertation, they will not be discussed in the following chapters.

3.9.1 LISP Cache performance

Current studies

A thorough analysis of LISP Cache is presented in [66], [72] and [73]. Their works complement and validate each other. The papers show that although the cache time largely reduces from initial 24-h to as small as 60s, the LISP cache maintains a very high efficiency, with more than 99% hit-ratio. As a result, the cache highly reduces the size, and the traffic volume overhead introduced by both the traffic encapsulation and the control traffic to request missing mappings is limited to 6%. Around 70% of cache missing, it is generated by UDP protocol regardless the port number. The authors argue that is mainly caused by P2P applica-

tions or any other distributed applications generating the same spread traffic pattern. They also compare the the original LISP specification, called vanilla LISP with the symmetric LISP model, proposed by Saucez et al., and show that the later model provides higher security. Besides, they evaluate the relationship between the size of cache and the number of users. It shows that when the number of users is increased to three times, the size of cache is only double. It means that the average size of the LISP Cache is not linear relationship with the number of hosts.

Missing work

As querying the mapping information to the MDS largely increases packet transmission delay, to store the necessary mappings in the LISP Cache can significantly alleviate it. However, as limited capacity of LISP Cache and dynamic changes of mapping information, it still requires smart algorithms to determine which mappings should be reserved or removed for better performance, especially when there are many expired mappings caused by mobility [53].

3.9.2 LISP security

LISP security impact

As LISP is born to address the scalability issues of Internet, it is not designed for the security by nature. The possible security threats with that LISP are faced are given in [107]. Although after the authentication, the MSes only accept the registrations of their compromised ETRs, the compromised ETRs may de-aggregate their EID-prefixes in order to register more EID prefixes than necessary to their MSes, so to cause overwhelmed Map-Requests. The compromised ETRs may also overclaim the prefix they own in order to influence the route followed by Map-Requests for EIDs outside the scope of their legitimate EID-prefix. Besides, the Map-Register stage is vulnerable to masquerading and content poisoning attacks, for example: the ETRs register the fake mapping information to the MSes [33] [89] [32]. By sending the Map-Requests at a high rate can overload the nodes in the MDS and cause the flood of Map-Replies as well. Some bits in the LISP control plane messages also have some security issues, for example: the Locator-Status-Bits can cause a DoS (Denial-of-Service), a replay, a packet manipulation, or a packet interception and suppression attack; the Map-Version number can mount a DoS, an amplification, or a spoofing attack; the routing locator reachability is able to mount a packet manipulation and an amplification attack.

A set of security mechanisms (LISP-SEC) that provides authentication, integrity and anti-replay protection for LISP are introduced in [83]. A number of vulnerabilities should be considered before deploying LISP in large scale [101].

LISP security proposal

The reference [32] analyses the security and functionality of the LISP mapping procedure using a formal methods approach based on Casper/FDR tool. The authors point out several security issues in the protocol such as the lack of data confidentiality and mutual authentication. The reference [33] figures out that the Map-Registration stage is vulnerable to masquerading and content poisoning attacks. Therefore, it presents a new security method for protecting the LISP Registration stage. The proposed method uses the ID-Based Cryptography (IBC), which allows the MDS to authenticate the source of the data.

3.9.3 The maximum transmission unit (MTU) issue

LISP encapsulation adds 36 bytes for IPv4 (i.e., 20 bytes of outer header + 8 bytes of UDP header + 8 bytes of LISP header) and 56 bytes for IPv6 (i.e., 40 bytes of outer header + 8 bytes of UDP header + 8 bytes of LISP header) to the original packet size. As with other protocols that encapsulate or tunnel traffic (e.g., IP Security (IPsec) [113] and generic routing encapsulation (GRE) [51]), it is possible that the resultant encapsulated packets would exceed the permissible MTU somewhere along the transit path. LISP provides both stateful and stateless mechanisms for handling potential MTU issues [6], with the main goals being (i) preventing packets from being dropped, and (ii) preventing the need for ETRs to perform packet reassembly prior to LISP de-encapsulation.

Based on real Internet experiments, [66] and [73] show that although the encapsulation overhead is not negligible, it is very limited (in the order of few percentage points of the total traffic volume). In addition, no major issue due to the MTU is observed on the LISP Beta Network.

Chapter 4

Evaluation of LISP mapping system

4.1 Introduction

As presented in Chapter 2, to retrieve the routing information (RLOCs) for EIDs, LISP relies on a pull model by explicitly querying the Mapping Distribution System (MDS). The performance of the mapping system is of paramount importance for LISP. Thus, we leverage on *LISP Beta Network*, where LISP has been deployed in the wild since 2008 [11], to conduct measurement campaigns about the mapping system. Such a worldwide real playground has been used to assess the state of LISP deployment and evolution over the years [106]. Yet, to our best knowledge, nobody at the moment has evaluated the mapping system from the following two standpoints: Stability and Consistency. We thus propose a much deeper study and a more thorough evaluation, focusing on the exploration of all the mapping systems on LISP Beta Network through active LIG (LISP Internet Groper [52]) measurement. The temporal comparison shows that the mapping system is quite stable over time and the crosswise comparison indicates that the mapping system is consistent between network entities. However, some instability and inconsistency cases are also identified through evaluation, which indicates that the mapping system is still in infancy as well. Hence, apart from proposing a brand new taxonomy in order to classify them, we carry out an in-depth investigation of such cases to provide hints on how to improve the performance of LISP.

The remainder of this chapter is organized as follows. Sec. 4.2 presents an overview of the experiment campaigns and introduces the metrics that we use to evaluate the performance of mapping system. Sec. 4.3 and Sec. 4.4 respectively describe the mapping system stability and consistency results in details, as well as the instability and inconsistency cases by using the proposed taxonomy, and investigate in depth the causes from different viewpoints. Finally, Sec. 4.5 provides some findings in other scopes and concludes this chapter.

4.2 Methodology

In order to collect data to carry out our analysis, several measurements campaigns have been performed. Hereafter we provide an overview on the measurements and the metrics used in our evaluation.

4.2.1 Measurements Overview

To obtain a comprehensive and in-depth evaluation of LISP Beta Network Performance, we conducted experiments using a tool called LIG (LISP Internet Groper [52]). This management tool is used to query the LISP mapping system. Its usage is described in details in Sec. 3.2.2 of Chapter 3. When LIG is launched, a *Map-Request* is sent for a destination IP address, and the returned *Map-Reply* is then stored for further analysis. It exists three possible types of outcome: (i) *LISP Map-Reply*, where the reply contains mapping information to reach the LISP-enabled sites; (ii) *Negative Map-Reply*, where the reply states that the requested IP address belongs to a prefix of a non-LISP site; (iii) *No Map-Reply*, where no answer is received within some amount of time. A query is considered to be successful if either a LISP Map-Reply or a Negative Map-Reply is received. While a query is considered to be failed if No Map-Reply is received within 5 seconds.

The experiments are conducted from July 2nd to 18th 2013. From a set of different vantage points (VPs), we used LIG to send *Map-Request* messages to all the Map-Resolvers of the LISP Beta Network, for all selected destination IP addresses every 30 minutes. The set of VPs was composed by part of academic networks, commercial Internet, PlanetLab and spread across Europe and USA. The destination IP addresses consisted of the first address within each prefix regardless of the type (i.e., the EID-prefixes or regular prefixes), derived from the LISPmon project [19]. The latter is a LISP Beta network monitoring site that periodically crawls the IP addressing space, through the MDS, to discover all the prefixes used as EID addressing space. During this measurement campaign, 5 VPs and 13 MRs/MSeS were available and 613 EIDs were queried in total.

We analyzed the Map-Replies associated with EIDs instead of all regular IP addresses. In the remainder, we use the term *experiment round* to identify a complete *Map-Request/Map-Reply* exchange by LIG for all the selected destinations to all the available MRs over all existing VPs. Although No Map-Reply occasionally appears in some experiment rounds, since it does not contain any mapping information, we just focus on LISP Map-Reply and Negative Map-Reply to conduct the analysis.¹

¹We use "*simultaneously*" or "*at the same time*" as synonym of "*at a same experiment round*".

Table 4.1 Map-Reply Differences

Category	Mismatch
Map-Reply Type	Negative Map-Reply and LISP Map-Reply
RLOC set	RLOC number RLOC address RLOC priority RLOC weight RLOC state

4.2.2 Metrics

As indicated in Sec. 2.1.2 of Chapter 2, the LISP Control plane is in charge of offering the mapping information that consists of an EID-prefix associated with a list of $\langle RLOC, Priority, Weight \rangle$ tuples. Theoretically, unless configuration has changed, if the same EID-prefix is queried, Map-Replies should always contain the same mapping information, no matter from which VP the query is sent, and no matter which MR is queried. During our measurements, however, we found that the Map-Replies are not always identical as they are supposed to be. The reasons causing the Map-Reply changes are summarized in Table 4.1.

The Map-Reply differences can be classified into two broad categories: the first is a difference in terms of *Map-Reply Type*, which means that a mix of Negative Map-Reply and LISP Map-Reply occurs; the other type of difference is in the *RLOC set*, which means that at least one of the attributes associated with RLOC set is different. More specifically, *RLOC number* indicates a difference in the number of RLOC tuples received in the Map-Reply. The others are differences in a RLOC tuple, which consists of RLOC address, RLOC priority, and RLOC weight (presented in Sec. 2.1.2 of Chapter 2).² The RLOC state is derived from the *R-bit* in Map-Reply message, which is set when the sender of a Map-Reply has a route to the RLOC, meaning that from its viewpoint the RLOC is up and running. If at least one of the six parameters shown in Table 4.1 conveyed in a Map-Reply is different from the previous one (over time) or the others (between different VPs or MRs) at the same experiment round, the Map-Reply is marked as different. Based on the criteria of the Map-Reply difference, we propose two metrics to evaluate in more details the performance of the LISP mapping system: (i) *Stability* and (ii) *Consistency*.

A mapping is considered to exhibit *stability* if the Map-Replies for one specific EID obtained from one MR observed at a given VP (briefly denoted as a tuple $\langle EID, MR, VP \rangle$),

²The terms *RLOC address* and *RLOC* are used interchangeably in this chapter.

over time, are identical. A classification taxonomy for unstable cases is proposed later in Sec. 4.3.

A mapping exhibits *consistency* if: *i*) Map-Replies for a specific EID from all MRs observed at a given VP (denoted as a group $\langle EID, *, VP \rangle$) are identical at the same time (this kind of consistency is referred as *Consistency by MR*); *ii*) Map-Replies for a specific EID-MR pair in all VPs (denoted as a pair $\langle EID, MR, * \rangle$) are identical at the same time (referred as *Consistency by VP*). A classification taxonomy for non-consistent cases is proposed later in Sec. 4.4.

4.3 Mapping System Stability Evaluation

In this section, we study the stability of the mapping system over time. More precisely, we consider that an EID is *stable* if all mappings retrieved for this EID are identical. If it is not the case, we say there is *instability*.

Overall, when we consider the whole dataset, 91.49% of observations are stable, which indicates that over the full set of observations, changes are infrequent and the MDS is stable. When the stability is considered on a per vantage point basis, we observe, as shown in Fig. 4.1, that the stability is independent of the vantage point since no significant deviation from the overall average can be observed (dashed line in Fig. 4.1). Nevertheless, while these results demonstrate that the mapping system is generally stable, in the following of this section, we investigate the reason why about 8.51% of the observations exhibit instability.

Fig. 4.2 complements Fig. 4.1 by showing, for each tuple $\langle EID, MR, VP \rangle$, the frequency $f_{\langle EID, MR, VP \rangle}$ at which mappings are different, by using the following definition:

$$f_{\langle EID, MR, VP \rangle} = \frac{\# \text{Instabilities}}{\# \text{Experiments}} \times 100\%. \quad (4.1)$$

Thus, if the instability frequency $f_{\langle EID, MR, VP \rangle}$ is 0%, it means that such a tuple is always stable during the whole experiment, i.e., its Map-Replies are always the same. On the contrary, a tuple $\langle EID, MR, VP \rangle$ with instability frequency of 100% indicates that its current Map-Reply is never identical to the previous ones, i.e., its Map-Replies change at every experiment round.

Fig. 4.2 shows the CDF of the mapping changes frequency (i.e., instability frequency). It shows the probability where the instability frequency equals to zero is 91.49%, which is coherent with the result in Fig. 4.1 (dashed line) and means that the mapping is globally stable. Apart from the stability case, the others (100.0% – 91.49%) belong to the mapping changes case, i.e., instability. This CDF also shows when the instability frequency equals to




Fig. 4.1 Overall percentage of stability observed at five different VPs.

1%, the CDF reaches to 98.53%, which indicates that among the instability cases, 82.77% ($\frac{98.53-91.49}{100.0-91.49}\%$) of all mapping changes events have an instability frequency less than 1%. We can deduce that mapping changes, if they occur, are rare events. The existence of instability frequency around 50% means that, in some cases, even though it is rather rare, the mapping changes occurs almost one time every two experiment rounds.

While Fig. 4.2 indicates that mappings can potentially change frequently, Fig. 4.3 shows the CDF of the stability duration, i.e., for how many experiment rounds a mapping remains stable. We observe that the majority of instabilities has a stability duration of 1 experiment round. This fact shows that the mapping change occurs during each measurement period (i.e., 30 minutes, the interval between two experiment rounds). At the other side of the curve, stability lasts up to 800 experiment rounds are observed, indicating that the mapping changed at the very beginning of or the end of the measurements we performed. In the dataset we identified both cases for such situations. To further understand the instability, we classify it into four categories:

- *New Deployment*: the Map-Replies of a tuple $\langle EID, MR, VP \rangle$ mix two different types: Negative Map-Reply and LISP Map-Reply. (It is caused more often by turning on/off the xTRs than the real deployment of LISP sites. Since the consequence of Map-Reply changes between Negative Map-Reply and LISP Map-Reply, we make this case named New Deployment.)




Fig. 4.2 CDF of instability frequency. X-axis is instability frequency, indicating the ratio of instability occurrence number for one $\langle EID, MR, VP \rangle$ tuple over total experiment number.

Table 4.2 Percentage of Instabilities by Category

New Deployment	Statistical outliers	Mobility	Reconfiguration
72.36%	7.23%	0%	20.41%

- *Statistical outliers*: the Map-Replies of a tuple $\langle EID, MR, VP \rangle$ change more frequently than usually seen. In our measurements, outliers corresponds to mappings that change at least 3 times during a day.
- *Mobility*: the Map-Replies of a tuple $\langle EID, MR, VP \rangle$ are explicitly tagged with the mobility bit as specified in RFC6830 [47].
- *Reconfiguration*: all cases that fit neither into the New Deployment case nor the Statistical outliers case.

Tab. 4.2 presents the percentage of observations for the four categories. New Deployments dominate the total unstable cases with a percentage of 72.36%, followed by Reconfiguration with 20.41%, and Statistical outliers with 7.23%. Throughout the dataset, we did not observe any Mobility event, probably due to the fact that the mobility specifications were not yet implemented at the time of measurements. Therefore, we omit the Mobility category in the rest of this chapter.




Fig. 4.3 CDF of stability duration (in number of experimental rounds). X-axis is stability duration, indicating during how many experiment rounds that the mapping for one $\langle EID, MR, VP \rangle$ tuple is stable.

New Deployments are the most common instabilities and we used spectral analysis with Fourier transforms and auto-correlation to understand if such events are correlated between themselves. This analysis shows that New Deployments are equally spread over the whole measurement period, i.e., their occurrence per day is relatively constant. We performed the same study on the Reconfiguration category and Statistical outliers and reached the same conclusion, i.e., the absence of correlations between events.

Since New Deployments are incidental events that can happen at any time, no particular duration is observed for this category of instabilities. Moreover, as New Deployments account for 72.36% of instabilities, it significantly impacts the duration distributions shown in Fig. 4.3, explaining the relatively smooth evolution of the CDF. On the contrary, Statistical outliers are characterized by very frequent changes and they bias the distribution around the lowest duration values, explaining the relatively high number of observations for small duration.

To characterize the mapping changes happening during instabilities, we derived the following *dissimilarity* metric $d(m_i, m_j) = 1 - \frac{|m_i \cap m_j|}{|m_i \cup m_j|}$, based on the Jaccard similarity coefficient, where m_x is the RLOC set of mapping x . A dissimilarity of 1 indicates that no RLOCs are shared between the two compared mappings, while a dissimilarity of 0 indicates that the RLOC sets are the same. Fig. 4.4 shows the CDF of the mappings' dissimilarity for each



Fig. 4.4 CDF of the dissimilarity among different types of instabilities.

unstable $\langle EID, MR, VP \rangle$ tuple. In our experiment, the dissimilarity results are discrete (as the discrete points shown in the Fig. 4.4) in a range from 0.33 to 1.0. As one can expect, the most likely dissimilarity is 1 in 90.09% of the total number of instabilities. These observations mostly come from New deployment case where before turning on the xTR, no RLOC is provided with the Negative Map-Reply (RLOC set is empty), hence, switch to a LISP Map-Reply implies a full change in the RLOC set. As 90.09% is larger than the 72.36% of New deployments share, it indicates as well that complete change of the RLOC set have also been observed during Reconfiguration and Statistical outliers. For the rest, no particular trend can be observed showing that changes of RLOC set can actually take any form.

4.4 Mapping System Consistency Evaluation

In this section, we discuss the consistency of the mapping system respectively by MR and by VP. More precisely, we consider that a tuple of $\langle EID, *, VP \rangle$ exhibits consistency, if all the Map-Replies from different MRs, for a certain EID, observed at a given VP are identical at the same time (same experiment round). This kind of consistency is denoted as *consistency by MR*. Similarly, if all the Map-Replies sent by a specific MR, for a certain EID, and observed at different VPs are identical at every experiment round, then we consider that

a tuple of $\langle EID, MR, * \rangle$ is consistent, and denote it as *consistency by VP*. Otherwise, for cases not falling in the above two, we simply say that there exists *inconsistency*.

Overall, throughout the whole dataset, we found that 86.3% of observations of all $\langle EID, *, VP \rangle$ tuples are consistent, which indicates that non-identical Map-Replies from different MRs are uncommon. Similarly, consistency by VP is observed in 90.48% of observations of all $\langle EID, MR, * \rangle$ tuples, showing that the Map-Replies received at different VPs are usually identical. In the following, Sec. 4.4.1 will go deeper in the analysis of consistency by MR, proposing as well several types of inconsistency; then Sec. 4.4.2 will provide thorough analysis of consistency by VP.

4.4.1 Consistency Evaluation by MR

The overall percentage of consistency by MR is 86.3% (see the dashed line on Fig. 4.5), meaning that across the whole experiment traces the inconsistency among the Map-Replies sent by different MRs is not common. Yet, we can also observe from Fig. 4.5 that different VPs show some deviations from the overall consistency by MR, reflecting the fact that while certain VPs receive the same Map-Reply from different MRs, some other VPs receive inconsistent responses. Such deviation is very small (no more than 2.77%) among different vantage points, indicating that it remains a relatively rare event. These results show that the mapping system is consistent among different MRs in general, but the inconsistency exists at the same time with a percentage around 13.7%. In the following of this section, we go deeper this issue, exploring the reasons why the few observations exhibit inconsistency by MR.

We compute the CDF of the inconsistency frequency by MR: $f_{\langle EID, *, VP \rangle}$, i.e., the probability of inconsistency frequency over all $\langle EID, *, VP \rangle$ tuples, by using the following definition:

$$f_{\langle EID, *, VP \rangle} = \frac{\# \text{ Inconsistencies by MR}}{\# \text{ Experiments}} \times 100\%. \quad (4.2)$$

For a $\langle EID, *, VP \rangle$ tuple, an inconsistency frequency of 0% means that all Map-Replies from all of the different MRs for one EID are the same at each experiment round throughout the whole measurement campaign. For a $\langle EID, *, VP \rangle$ tuple, an inconsistency frequency of 100% means that at each experiment round during the whole measurement campaign, there is at least one MR replying a Map-Reply different from the replied by other MRs. The resulting CDF of inconsistency frequency is shown on Fig. 4.6. The value 0% of inconsistency frequency has a probability of 82.38%, indicating that through all the dataset, there are 82.38% of $\langle EID, *, VP \rangle$ tuples that are always consistent by MR at whichever VP. Yet, it is worth noticing that this value is lower than the overall consistency of 86.3% shown



Fig. 4.5 Overall consistency by MR at 5 different Vantage Points.

in Fig. 4.5. The reason of such difference lays in the fact that, when we analyze the 0% inconsistency frequency, we just take into account those common EIDs exhibiting a consistency by MR in all VPs. Hence 82.38% represents the subset of the EIDs considered in the former percentage and is the lower bound, guaranteeing that 82.38% of $\langle EID, *, VP \rangle$ tuples always have consistent responses, regardless of the VP. In addition, the inconsistency frequency concentrates at a very low range, since a sharp increase occurs between 0% and 2% with a growth of 16.15% (98.53% - 82.38%). It indicates that the $\langle EID, *, VP \rangle$ tuples classed as inconsistent do actually receive inconsistent responses rarely, hence, demonstrating that the mapping system is generally consistent. Yet, in very few cases (only 0.5%), the inconsistency between different MRs occurs all the time.

So far, the analysis indicates that the mapping system is consistent by MR in general, but the inconsistency indeed exists. To have a closer look to the inconsistency, we further cluster it into two cases:

- **Map-Reply Type:**, The type of Map-Reply provided by the 13 MRs for a specific $\langle EID, *, VP \rangle$ tuple is different, meaning that some MRs provide LISP Map-Reply while others return a Negative Map-Reply. The difference of received Map-Reply type leads the xTRs to use a different way to send the packets, i.e., xTRs receiving a LISP Map-Reply sends LISP-encapsulated packets by using the obtained RLOCs, while xTRs receiving a Negative Map-Reply will forward the conventional packets (i.e., without LISP-encapsulation).

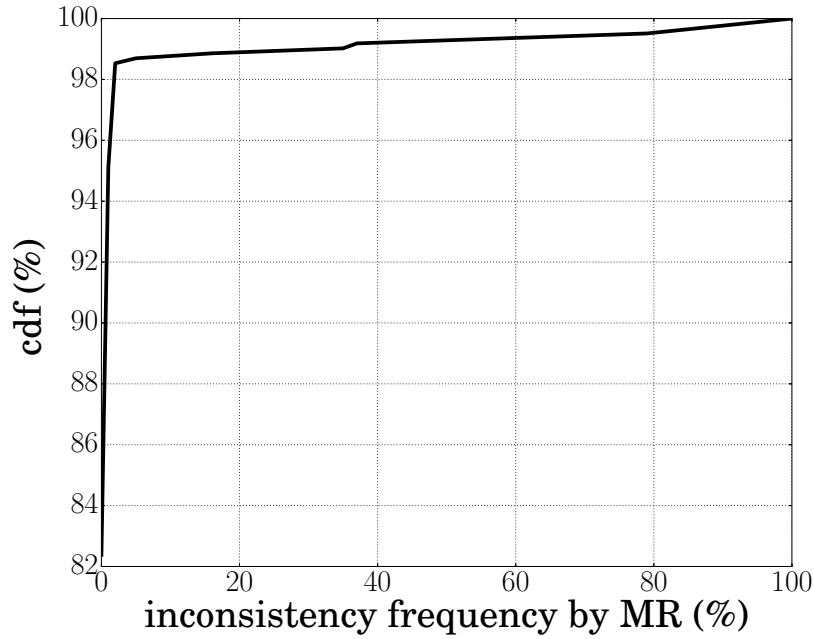


Fig. 4.6 Cumulative Distribution Function of inconsistency frequency by MR. X-axis is inconsistency frequency by VP, indicating the ratio of inconsistency occurrence for each $\langle EID, *, VP \rangle$ tuple over total experiment number.

- **RLOC set:** All the Map-Replies returned by the 13 MRs for a $\langle EID, *, VP \rangle$ tuple belong to LISP Map-Reply type, but the associated attributes (cf. Tab. 4.1 from Sec. 4.2) are different. For instance, some Map-Replies may contain $RLOC_1$, while some others convey $RLOC_2$ (i.e., the RLOC address is not identical); or some Map-Replies have only one RLOC, whereas others have multiple RLOCs (i.e., the RLOC number is different), and so on. The difference in such Map-Replies could lead the requesting xTRs to select different destination RLOCs.

The large majority of inconsistencies, around 98.34%, are Map-Reply Type, probably because there is an update of mapping information in the mapping system for the New Deployment case. Not all of the 13 MRs simultaneously update the mapping information, meaning that a convergence time is needed so that the global mapping information maintained in all parts of the mapping system is synchronized. If VPs query the mapping system during this convergence period, some MRs will still provide Negative Map-Replies, while other (already updated) MRs will forward the Map-Requests towards the corresponding ETRs, which in turn will send the updated LISP Map-Replies, and vice versa. This situation contributes to the inconsistency frequency around 1%, which means that there is only once or twice inconsistency between Map-Replies from different MRs for one specific EID, since changing the status of xTR is not frequent.




Fig. 4.7 Overall consistency by VP from the different Map-Resolvers.

Just few inconsistencies, only 1.66%, are caused by the RLOC set. This scenario usually occurs for those EIDs whose RLOCs change frequently over time (the case of Statistical outliers instability). So, because of the convergence time of the mapping system, it is very difficult to guarantee that every MR can provide the same LISP Map-Reply at every experiment round. This reason for changes in the RLOC set is most probably due to traffic engineering policies. Similarly as the previous case, when different MRs provide a different RLOC set for a specific EID, the requesting xTRs may end up using different RLOCs to reach the same destination EID.

4.4.2 Consistency Evaluation by VP

The overall percentage of consistency by VP for the whole experiment dataset is 90.48%, presented by the dotted line in Fig.4.7, which indicates that the Map-Replies received at different VPs from the same MR are highly consistent. Yet, we can also observe that different MRs show some deviations from the overall consistency by VP, reflecting the fact that certain MRs send the same Map-Reply to all VPs, but some do not. The deviation among different MRs is however limited (less than 12.23%). Therefore, we can certainly state that the mapping system shows generally a high degree of consistency by VP but presents a non-negligible amount of inconsistencies (9.52%). In the rest of this section, following the same




Fig. 4.8 Cumulative Distribution Function of inconsistency frequency by VP. X-axis is inconsistency frequency by VP, indicating the ratio of inconsistency occurrence for each $\langle EID, MR, * \rangle$ tuple over total experiment number.

approach as for the consistency by MR, we go deeper this issue, exploring the reasons why some observations exhibit inconsistency by VP.

Fig.4.8 shows the CDF of the inconsistency frequency by VP: $f_{\langle EID, MR, * \rangle}$, i.e., the probability of inconsistency frequency for all $\langle EID, MR, * \rangle$ tuples, by using the following definition:

$$f_{\langle EID, *, VP \rangle} = \frac{\# \text{ Inconsistencies by VP}}{\# \text{ Experiments}} \times 100\%. \quad (4.3)$$

Like for the explanation for Fig. 4.6, here an inconsistency frequency by VP of 0% means that all Map-Replies received by all of the VPs from a same MR are totally consistent at every experiment round during the whole experiment campaign, while an inconsistency frequency of 100% means that at least one of VPs received an inconsistent Map-Reply compared to the other VPs at every experiment round. The probability of consistency by VP (when X-axis is 0% in Fig. 4.8) is similar to that by MR, with a value of 81.57%. It can be noticed that this percentage is slightly lower than the overall percentage of consistency by VP in Fig. 4.7. The reason is that here the inconsistent occurrence with 0% only includes those EIDs for which all VPs always show consistency, no matter which MR is actually queried. A sharp increase with a growth of 16.97% (98.53% - 81.56%) is found when X-axis is 1%, which means that although VPs sometimes receive inconsistent Map-Replies

at several experiment rounds, the occurrence remains pretty low, hence showing that even by VP the mapping system exhibits a very high degree of consistency. Differently from the largest inconsistency frequency by MR, which goes up to reaches 100%, the largest inconsistency frequency by VP is only 66%. This phenomenon indicates a fact that the Map-Replies received by different VPs are not inconsistent all the time, at least there are one third of observations show that their Map-Replies are totally same. Thus the consistency by VP is higher than consistency by MR.

To further explore the inconsistency by VP, we then use same metrics as in Sec. 4.4.1, with the exception that the queried target changes from $\langle EID, *, VP \rangle$ to $\langle EID, MR, * \rangle$ tuples.

The most inconsistent case, 89.06%, is due to a mix of Negative Map-Reply and LISP Map-Reply, i.e., demand for a same $\langle EID, MR, * \rangle$ tuple, at some VPs receive Negative Map-Reply but at others receive LISP Map-Reply. This is caused by the New Deployment case, but differently from the cause for inconsistency by MR, the reason for inconsistency by VP is probably because of Map-Requests sent by different VPs can not be processed exactly at the same time, due to queuing in the MR. Thus, if the queries are sent during the period of mapping information update, some VPs, whose Map-Requests are processed first will have different Map-Reply Type from the latter ones. Inconsistencies caused by different RLOC sets, i.e., different VPs receive LISP Map-Replies with different RLOCs set from the same MR, also exist, contributing for around 10.94%. It is probably caused by a combination of the Statistical outliers instability case and the queuing mechanism in MR. As the RLOC sets provided by a MR for some certain EIDs frequently change, and the Map-Requests sent by different VPs need to be processed one by one, it is very possible that the first processed Map-Requests are replied by the different RLOC sets compared to the latter ones. The difference between Map-Replies is probably due to traffic engineering policies, but the further experiments are needed to explore the deeper reasons for such inconsistency.

4.5 Conclusion

During the evaluation of LISP mapping system performance, besides the observations about the stability and consistency that we state in details in the previous sections, we also investigate some interesting LISP-related results. We provide them here so that they can be used to improve LISP performance. LISP relies on mappings, which allow LISP-sites to be reachable through several locators. We identified that 8.97% of the LISP-sites leverage this functionality in order to be multihomed to the Internet, with an average of 3 distinct locators. Interestingly, when we compare the locators with their Autonomous System Num-

ber we found that around 50% of the locators in multihoming situations belong to different ASes, which indicates that LISP is used to transparently interconnect LISP-sites by the intermediate of multiple ISPs.³ As the control-plane (i.e., MS/MR and mapping system) and the data-plane (i.e., ETR/ITR) are decoupled in LISP, Map-Replies can originate from the LISP-sites directly or from shared utilities constituting the mapping system (i.e., MS/MR). This particularity of LISP is already used in the current deployment, where only 88.8% of the non-negative Map-Replies are originated directly from their LISP-sites, with the remaining 11.2% being originated by delegated MSes replying on behalf of the LISP-sites. This allows reducing the control traffic seen by LISP-sites and help them to scale better.

In addition to mappings, LISP also relies on encapsulation, meaning that with LISP traffic can be transported over IPv4 or IPv6 without requiring end-hosts to be dual-stack. Interestingly, we identified that no less than 8.84% of the locators are IPv6. However, we have not observed cases where a LISP site was reachable strictly using IPv6 while 82.79% of sites are reachable only through IPv4. This discrepancy tends to confirm that network operators still lack of confidence being reachable only through IPv6. As a result, the 17.21% of LISP-sites explicitly using IPv6, always include at least one IPv4 address in order to be reachable.

As a solution to cope with the increase of BGP routing table, traffic engineering, and scalability issues, LISP has been widely deployed on the LISP Beta Network experimental testbed for ten years. To ensure accurate and efficient communications, LISP mapping system should be stable and consistent. However, there is no study evaluating the stability and consistency of LISP mapping system. For such purpose, we continuously measured the LISP Beta Network for seventeen days to assess the mapping system from these two aspects. Measurements show that the LISP mapping system is stable and consistent over time. Nevertheless, instability and inconsistency are observed and for studying them in details we developed a new taxonomy. All in all, instability and inconsistency are rare events. At last, we observe utilization of multi-homing and dual-stack (i.e., IPv4 and IPv6) during the whole experiment and plan to use these results to further improve the LISP performance.

³We used the Team Cymru BGP database to associate IP addresses to ASNs (see <http://www.team-cymru.org/IP-ASN-mapping.html>).

Chapter 5

LISP-Views: LISP mapping system monitor

Proposed nearly ten years ago, LISP is still at an early stage. Currently only two experimental platforms have already been deployed in the Internet: LISP Beta Network and LISP-Lab. However, only the LISP Beta Network is monitored with LISPmon that partially monitors the mapping system once a day. To accompany the growth of LISP, a dynamic and complete monitoring system is required. Therefore, we propose LISP-Views, a dynamic versatile large scale LISP monitoring architecture. LISP-Views allows to automatically conduct comprehensive and objective measurements. After running LISP-Views in the wild for several months and comparing the monitoring results with LISPmon, we confirm that LISP-Views provides more detailed and accurate information. We observe the different behaviours between every network entity within mapping system, and also explore the current LISP performance for further improvements.

In the remainder, Sec. 5.2.1 introduces the limitations of LISPmon and why we propose a new LISP monitor. Sec. 5.2.2 describes our proposed LISP monitoring architecture in details. Sec. 5.3 validates LISP-Views by comparing with LISPmon. Sec. 5.4 provides the snapshot of what kind of further analysis can be done with our proposal. Finally, Sec. 5.5 concludes the chapter.

5.1 Introduction

To promote the development of LISP and boost the related research, large scale flexible platforms are necessary. Two LISP-related platforms have been interconnected so far. The experimental LISP Beta Network testbed [11] is deployed since 2008, and the LISP-Lab

platform [15] is open to external experimenters since 2015. Currently, a unique LISP monitoring system called LISPmon [19] supervises the global MDS and publishes the mapping information daily. However, it is known that the mapping information sometimes changes frequently within a day and that the elements constituting the MDS are not always consistent [81]. We hence propose a dynamic versatile LISP monitoring architecture, namely LISP-Views, to overcome these limitations. LISP-Views automatically explores the whole MDS every 2 hours and stores the detailed mapping information, so to facilitate the experimenters to evaluate the LISP comprehensive performance.

We used a one-month long set of traces produced by LISP-Views to evaluate its performance and accuracy, including comparing it with LISPmon. We show that LISP-Views is more accurate than LISPmon since it monitors all the MDS elements in parallel. In addition, thanks to its detailed reporting, LISP-Views allows to assess high level metrics of LISP deployments such as reliability, latency, or configuration issues.

5.2 Proposed Monitoring Architecture

5.2.1 Motivation

In order to move LISP forward, we need to deeply understand the behaviour of the different LISP network entities and since the MDS reflects the status of a LISP network as it stores all the mapping information, it is essential to be able to monitor the MDS. LISPmon was the first step towards a systematic LISP monitoring. However, it monitors the MDS just from one vantage point (VP) once per day and only queries one MR. Upon MR issues, LISPmon must be manually re-configured to monitor another MR. Yet the mapping information may be unstable and inconsistent between MRs, i.e., the mapping information sometimes dramatically changes within a day, and the Map-Replies from the different MRs may not coincide at a given time [81]. It is similar to the world-wide distributed BGP looking glass servers, which do not always provide the same responses for an IP address as the whole routing system may not have converged or because of routing policies. From such point of view, LISPmon has strong limitations since it is not able to detect the changes of mapping information within one day, and is not able to show the differences among MRs. Thus, we propose a new versatile LISP monitoring architecture, called LISP-Views, to monitor public LISP deployments, as well as to enable further performance evaluation of LISP defined by the users themselves. In fact, LISP-Views not only can be used to monitor LISP, but also can be used in the non-public networks, such as VxLAN [82].

LISP-Views is an open source implementation and has been designed to fulfill the following objectives:¹

1. LISP-Views queries to all the working MRs existing in current MDS, in parallel, while LISPmon just queries one, aiming at building a complete view of the current LISP status.
2. LISP-Views periodically monitors all the MRs with arbitrary intervals,² while LISPmon just does it daily, aiming at providing information about the mapping evolution at smaller time granularity.
3. LISP-Views supervises the whole MDS without any manual process, automatically reacting to failing components (e.g., unresponsive MRs).
4. LISP-Views obtains the mapping information from all the MRs of the LISP Beta Network as well as LISP-Lab platform, whereas LISPmon prefers to leverage the MR of the former one.
5. LISP-Views is flexible and configurable, with users able to define different monitoring jobs and get the various measurements, whereas LISPmon just publishes the mapping list daily.
6. By design LISP-Views can be extended to be used as a monitoring facility for the internal deployed VxLAN.³
7. LISP-Views is able to be deployed on multiple VPs over the world, while LISPmon publishes the results depending on one VP.

5.2.2 Description of LISP-Views

The architecture of the proposed LISP-Views monitoring tools is depicted in Fig. 5.1. LISP-Views consists of several modules with different functions. The *Measurement*, *Report*, and *Raw Data* modules are deployed on a centralized server while, the *Crawler*, *Sonar*, and *Controller* modules can be deployed on several different VPs. As for now, LISP-Views is just deployed on one VP in Paris, France, where the monitor is developed and the centralized module is deployed. Nevertheless, we are planing to implement the distributed version of LISP-Views. For this reason, we will not discuss further how to deploy LISP-Views on multiple VPs. All the modules are implemented in *Python* and described as follows:

Sonar module: the main module with two functions: 1. sending the LISP encapsulated Map-Requests and receiving the Map-Replies to/from all the existing MRs based on the

¹Source code available on Github: <https://github.com/SeleneLI/LISP-Views>

²This interval can be changed by the experimenters to accord their necessity. But in this chapter, the interval is set the shortest according to the process capacity of our server.

³In this chapter we focus on LISP.




Fig. 5.1 Monitoring Architecture

standard [47]; 2. storing the received information in *Report* and *Raw Data* with different purposes.

The IP address that *Sonar* uses to query to MRs is selected either from the output of *Crawler*, or from the recorded information in the previously produced *Report*. The reasons are explained in the corresponding modules.

Crawler module: scans all the existing IPv4 addressing space using *Sonar*. If the queried IP (e.g., 192.0.2.1) has no Map-Reply, *Crawler* increments the IP by one as the next queried IP (i.e., 192.0.2.2) and lets *Sonar* send the Map-Request. If *Sonar* obtains Negative or LISP Map-Reply, since the Reply contains a prefix (mentioned in Sec. 2.1), *Crawler* sets the first IP beyond the returned prefix. For instance if the returned prefix is 192.0.2.0/25, the next queried IP is 192.0.2.128. Scanning the whole IPv4 address space takes very long time, and is mainly caused by time wasted waiting for no Map-Reply. Indeed, no Map-Reply means to wait 3 seconds, and the next queried IP is just increased by one, instead of skipping a block of IP addresses like in the case of Negative/LISP Map-Reply.

Report module: contains the collected EID-prefixes in a list, as well as a list of MRs that answered. As crawling the address space may take long time but we aim to obtain the status of MRs as frequently as possible, *Sonar* sends Map-Requests for the EIDs recorded

in *Report* only to MRs that have previously responded. Thus, it decreases the possibility to receive no Map-Reply, so that to get LISP status within a shorter time.

Raw Data module: contains all the detailed information of Map-Replies for each MR, such as Map-Reply type, RLOCs, required EID-Prefix, Round Trip Time (RTT) and the returned source for each round specify with *TimeStamp* (regardless the source of *Sonar*). So, it can be used to perform thorough performance analysis of MDS. Moreover, since the *Raw Data* is stored according to the MR, it is possible to track the performance of the MRs individually.

Measurement module: provides the composition of requested measurements (i.e., select different measurement plug-ins) based on the analysis of *Raw Data* and *Report*.

REST API: is connected to *Measurement module* so that the users can launch a custom experiment by setting the experiment time, the monitored MRs, and obtain the different aspects of LISP status. We are currently implementing the REST API, so for this paper the execution of measurements was performed using command lines.

Controller module: synchronizes all the modules in LISP-Views and also specifies the start and stop time of producing both *Report* and *Raw Data*. The interval of generating *Report* and *Raw Data* can be changed according to the hardware processing capability. However, since the inconsistency between MRs exists, the interval of producing *Report* and *Raw Data* for every MR differs from each other. Thus, the interval set in Control module should cover the slowest MR.

5.3 LISP-Views Validation

5.3.1 Methodology

In order to validate and evaluate the LISP-Views monitoring tool, we used raw data and report collected during one month (from 0:00 September 4th to midnight of October 4th 2016), by deploying LISP-Views on one VP, which is an xTR of LISP-Lab platform.

The interval to produce reports was 6 hours, and the interval to produce raw data was 2 hours. Unfortunately, both MRs of LISP-Lab platform just responded to the Map-Requests at the first time and then stopped. They were not able to handle large number of queries, because the Map-Requests fill the request queues, which are not deployed in manner to sufficiently follow up the queries, resulting a drop in the MRs. Per-se this is a success, because this bug was unknown, and the OpenLISP coders are working on a fix. All the results of evaluation used in this section only depend on the 6 working MRs of LISP Beta Network (3 in Europe, 1 in US and 2 in Asia) and the other 6 MRs were down at the moment




Fig. 5.2 Comparison between LISPmon and LISP-Views. The upper sub-figure shows the number of LISP Map-Reply from LISPmon generally at 7:00 and LISP-Views exactly at 8:00 over days. The bottom sub-figure indicates the LISP Map-Reply and overall Mapping from LISP-Views over time with an interval of 2 hours.

of conducting the experiment. The aim of the work is to validate LISP-Views by assessing if it provides at least the same results as LISPmon.

5.3.2 LISP-Views vs. LISPmon

In this section, we compare LISPmon and LISP-Views, so to assess if subset of information comparable between the two monitoring platforms are identical, hence validating LISP-Views. As we indicated in Sec. 5.2.1, at any time LISPmon just queries one MR, on LISP Beta Network once per day and generally begins at 7:00 a.m.. The queried MR is changed sometimes when it has issues. LISP-Views, however, keeps sending Map-Requests to all the 6 MRs with an interval of 2 hours everyday, so to retrieve the actual real status of each MR.

Fig. 5.2 shows the number of LISP Map-Replies received from LISPmon and LISP-Views over time. The upper sub-figure makes a comparison between the 2 monitors. The line with point is obtained from the daily LISPmon publications, indicating the number of LISP Map-Replies returned by one MR in an uncertain time monitoring round every day. To facilitate the comparison, we pick up the results from LISP-Views aft 8:00 a.m., which is the nearest experimental round to LISPmon. As presented in the curve with triangle symbols, LISP-Views reflects the fact of a combination of the number of LISP Map-Reply

containing the different EID-prefixes from all the MRs at a fixed experiment time. In most cases, LISP-Views receives around 20 LISP Map-Replies more than LISPmon, reflects the fact that the MRs are not coherent at any time due to the existence of convergence time for synchronizing the mapping information among them. The other 5 days, where LISP-Views receives very few LISP Map-Replies, show that there were issues to normally receive LISP status at that time, whereas LISPmon always presents a smoother trend but hides the reality.

The lack of stability can be better appreciated in the bottom sub-figure of Fig. 5.2. Such figure depicts the number of LISP Map-Reply (in dashed line) and the overall Mapping (LISP and Negative Map-Reply, in solid line) from LISP-Views for every 2 hours monitoring round during the whole experiment. Although the number of Map-Replies presented in the figure is still the combination of results from 6 MRs when the returned EID-prefixes are different, it shows that the MRs are not quite stable. Sometimes the MRs even provide 0 or less than 10 LISP Map-Replies, but in the next experiment round, the number is recovered. Besides, we evaluate the total number of successful queries every 2 hours. It also presents the behavior of instability, where it generally oscillates between 400 and 500, but it occasionally drops lower than 300. Both solid and dashed lines have almost same changing trend, indicating that the change of overall Map-Replies is mainly affected by LISP Map-Replies. Monitoring rounds where the number of LISP Map-Replies approaches 0 but where Negative Map-Replies are still received indicate issues at the MRs, especially for answering the LISP Map-Reply.

We compared the number of received LISP Map-Replies within a whole day between LISPmon and LISP-Views over 30 days. Shown in Fig. 5.3, as LISPmon only publishes one record each day, the number (blue points) is exactly identical to those in Fig. 5.2. LISP-Views, however, provides a combination of LISP Map-Replies not only from all MRs but also from 12 monitoring rounds within one day. We observe that our monitor architecture receives more LISP Map-Replies than LISPmon in all days with a large difference. 83.3% of the time LISP-Views receives more than 20 LISP Map-Replies. The maximum number of LISP Map-Replies that our proposed monitor obtains is 55, while the maximum value for LISPmon is 43. Differently from the upper sub-figure of Fig. 5.2, where LISP-Views occasionally receives nearly 0 LISP Map-Replies, the total number within a whole day is always more than 40, which illustrates that the MRs may not work normally in some rounds but they are able to recover fast.

Since LISP-Views repeats querying simultaneously to every MRs, it is able to report on the status of each MR at anytime. It provides more complete mapping information compared to LISPmon and is also able to highlight sporadic issues with MRs.




Fig. 5.3 Comparison between the LISP Map-Reply of LISPmon and LISP-Views over days

5.4 Dissecting LISP with LISP-Views

After validating LISP-Views in the previous section, this section presents several examples of how LISP-Views can be used to dissect LISP and obtain in-depth results. The experimental dataset used is the one in Sec. 5.3.1.

Fig. 5.4 shows the *Reliability* of each MR during our data collection, by calculating the percentage of successful queries over the total number of Map-Requests. MR1, MR4, and MR5 give the highest reliability values, which are more than 60%. The lowest reliability values, about 20% are observed for MR2, MR3, and MR6. In general, the reliability of each MR is different and low compared to the years of 2012 and 2013 presented in [44], which coincides with the fact that there was a change of the MRs architecture on LISP Beta Network that year, and an updated-software was tested on MRs as well. The low value of reliability is caused by MRs having an unstable behaviour, as shown in Fig. 5.2, the number of Map-Reply changes heavily over time, and especially the number of LISP Map-Reply sometimes drops to 0.

In order to understand the behavior of each MR in terms of latency, we analyze the median RTT obtained from our dataset. The RTT here refers to the Round Trip Time from sending out the Map-Request until receiving the Map-Reply. Fig. 5.5 shows that the best performance come from MR1, MR3, MR4, and MR5; since the overall RTTs are much

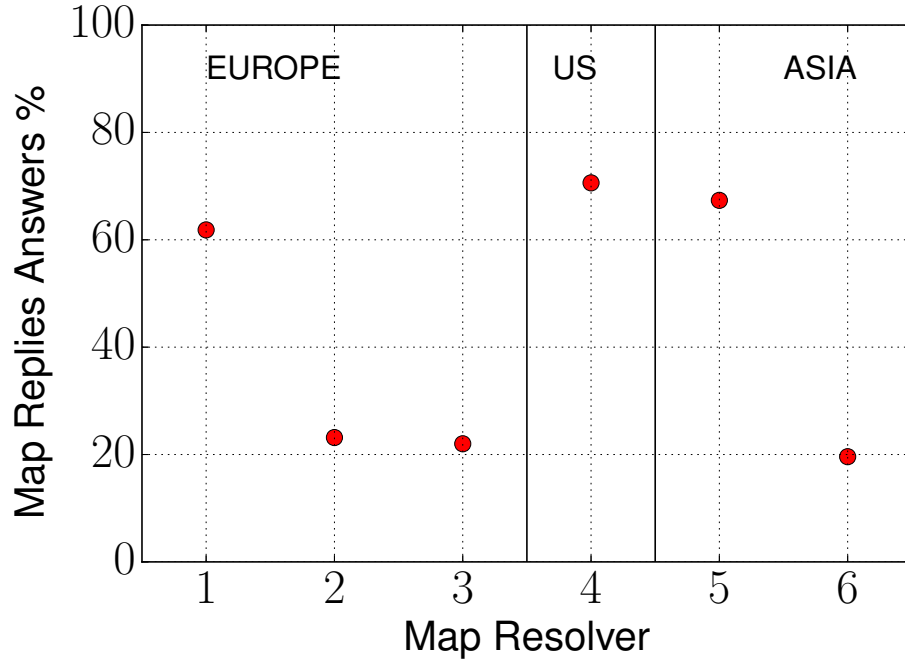


Fig. 5.4 Reliability of each MR

lower than the others. For MR3, the number of LISP Map-Replies is much higher than the number of Negative Map-Replies, probably because its embedded Map-Server registers less EID-prefixes hence requiring Map-Requests to be forwarded to a remote MS. Then, the Map-Request is forwarded to the xTR, where the remote MS registers and the xTR gives back the Map-Reply. Compared to the Negative Map-Replies, which are normally returned by MR, LISP Map-Replies take longer time, especially in this case. However, MR2 and MR6 present very high latency, particularly for MR2 (located in Europe), but its RTT is even higher than MR5, which is located in Asia. It is mainly caused by a very high CPU usage of these two MRs preventing them sometimes to reply. If we focus on the RTT of LISP Map-Replies, only MR1, MR4, and MR5 provide the best behavior, which coincides with the result of Fig. 5.4 on reliability. The high RTTs of LISP Map-Replies are from the other 3 MRs, where the median RTT is around 1300 ms, is caused by the failed queries. Furthermore, it also explains why we can always receive the Negative Map-Replies, but the number of LISP Map-Replies is sometimes quite low in Fig. 5.2. Half of LISP Map-Replies are more than 1300 ms and partly even more than 3 s. Since our measurement timeout is set to 3 s it implies that some Map-Replies are dropped by our measurement unit. As the RTT of Negative Map-Replies is generally lower, we can receive more of them. This phenomenon may be very biased by the VP location. Thus, deploying LISP-Views on multiple VPs to get more reliability information is a paramount future work. In addition, Fig. 5.5 presents

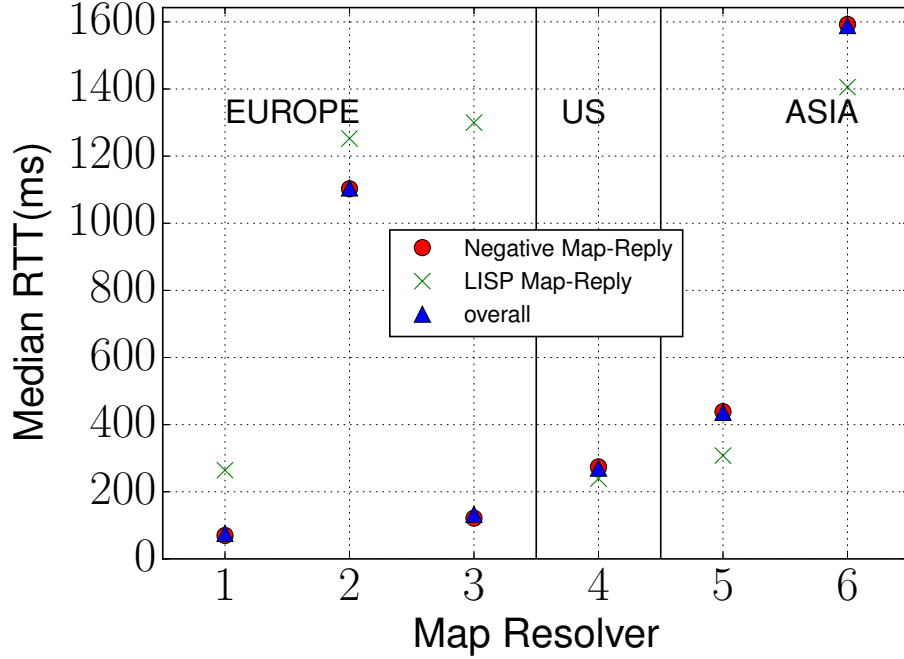


Fig. 5.5 Median RTT per MR (In the most time, the Negative Map-Reply and overall are overlapped.)

that the number of Map-Replies returning from the MRs located in Europe and US is indeed higher than the number of Negative Map-Replies, which is expected. But both Asian MRs return LISP Map-Replies faster than Negative Map-Replies. The reason of this observation is unclear and requires further exploration.

We also calculate the Empirical Cumulative Distribution Function (ECDF) of the RTT for the 6 MRs. Fig. 5.6 provides the ECDF of the RTT for Negative, LISP, and overall Map-Replies. As indicated in Sec. 2.1, the LISP MDS needs more time to solve a complete mapping (LISP Map-Reply) than the Negative Map-Reply. We observe that the Negative Map-Reply is indeed faster than the LISP Map-Reply until the RTT reaches 1000 ms, then the behavior changes, i.e., the LISP Map-Reply sometimes becomes faster. What's more, we find that 67.03% of RTTs within 500 ms, 70.33% less than 1000 ms and 90.88% don't exceed 1500 ms for overall Map-Replies. In details, for the Negative Map-Reply, we found that 67.8% of RTT values within 500 ms, for the LISP Map-Reply that 59.7% of RTT values less than 500 ms. This figure almost presents a bi-modal distribution, where 500 ms is a peak and 1300 ms is another one. These two high occurrences of RTTs are exactly the most two frequent latency in Fig. 5.5.

The dataset obtained with our monitoring architecture also shows the information about *Mapping Source*. We explore the source answering the Map-Replies according to different




Fig. 5.6 ECDF RTT comparison between LISP and Negative Map-Reply

Table 5.1 Percentage of Mapping Source

Map-Reply Type	Map Resolver	xTR	Other
Negative Map-Reply	98.79%	-	1.21%
LISP Map-Reply	0.14%	88.37%	11.49%

types (LISP or Negative). In the case of LISP Map-Replies, we expect the source of replies to come either from one of the ETRs or the queried MR. On the contrary, for Negative Map-Replies, replies should just come from the queried MR. Tab. 5.1 presents the percentage of observations for the two types of Map-Replies. For the Negative Map-Reply, 98.79% come from the queried MR and 1.21% come from the other sources. Further, the other sources are actually the other MRs without query and it happens just for a fixed EID-Prefixes, i.e., if we send a Map-Request for one of these EID-Prefixes to a dedicated MR, the Map-Reply always comes from a fixed specific MR. As a conclusion, all the Negative Map-Replies are answered by MRs. For the LISP Map-Reply, we observe that 0.14% come from the queried MR, 88.37% come from one of their ETRs, but 11.49% come from the other sources in different locations. This unexpected behaviour needs further investigation.

The following type of measurement is the distribution of the size of RLOC set, i.e., how many RLOCs are associated to one EID prefix on average. As previously observed, the percentage of mappings using two or less RLOCs has increased between 2010 and 2012 [106].




Fig. 5.7 CDF of the number of RLOCs

Fig. 5.7 shows that this trend keeps going on, i.e., more mappings use fewer RLOCs and the maximum number of RLOC is 4. An interesting point is that although LISP is a good candidate to support the increasing multi-homing in Internet, more than 60% LISP users are not multi-homed and among them the majority only has 2 RLOCs. Not only the number of RLOCs of each site does significantly change, but also the RLOCs themselves remains stable. We measured that the stability reaches 99.8% for all the dataset, i.e., once an EID-prefix – to – RLOCs mapping is decided, it rarely changes. Surprisingly, we have not found any mobile LISP sites.

All the aforementioned observations are based on the experiment made between September and October 2016 as described in Sec. 5.3.1. Later LISP-Views detects that MR2 and MR6 in Fig. 5.5 with very high overall RTTs are down, and three new other MRs are up. After confirming with the operators of the LISP Beta Network, these two MRs are indeed definitely down. The change of the architecture of MRs also proves the accuracy of LISP-Views, while the LISPmon presents a smooth change in its daily reports, hence hiding these facts. Since February 2017 LISP-Views publicly publishes preliminary daily reporting online [16]. Development efforts are still ongoing to provide a complete and production-level website. Fig. 5.8 is a capture from the website, it is one type of measurement about the number of EID-prefix quantified by the different size of RLOC set. The shown result is an union of all MRs during a whole day, to present the most complete mapping information of




Fig. 5.8 EIDs by Quantity of Associated RLOCs [16]

the actual LISP deployment. The line on the top indicates the number of LISP Map-Replies, i.e., the Map-Replies with at least 1 RLOC, which coincides to the results shown in Fig. 5.2 and is mainly affected by the EID-prefix with 1 RLOC. Moreover, the composition of the different sizes of RLOC set is also rather identical to the one in Fig. 5.7. The only difference is that in the previous dataset there is no EID-prefix with 3 RLOCs, while in the latest dataset it appears, but not very stable. In addition, the line with 3 RLOCs is almost complementary to the line with 4 RLOCs. It is probably caused by one LISP-site that has 4 xTRs but one being always down, or the LISP-site having 4 interfaces on a same xTR among which one is down. Besides, the number of observed EIDs per day heavily drops two times in the figure. Since the two valley don't drop to zero, we know that the problem doesn't come directly from the MDS. Instead, the problem comes from issues that occurred within the network of the LISP-Views server itself. This observation highlights the urgent need of distributing VPs.

5.5 Conclusion

Very little is known about the behavior of LISP in operational environments and it still lacks of troubleshooting tools. Motivated by the only LISP monitor deployed so far, named LISPmon, which records the current LISP status only from a specific MR just once per day, we propose a more dynamic LISP monitoring architecture, namely LISP-Views, so to deepen the understandings on LISP and to ease day-to-day operations and troubleshooting.

As LISP-Views aims at being deployed in large scale and dynamic networks, we make a comparison with LISPmon to validate the former one by comparing their behavior during one full month.

It demonstrates that LISP-Views provides more information by discovering more mapping information from all MRs and more complete mapping information of each MR. Furthermore, with our proposed monitoring platform, more mapping system performance metrics such as reliability, latency, and configuration issues can be assessed, which helps for further LISP improvements. LISP-Views is still in the first phase, so the implementation of a REST API, the deployment on multiple VPs and the test of IPv6 behavior are still ongoing and future work.

Chapter 6

Assessing LISP interworking performance through RIPE Atlas

LISP is currently under standardization in IETF and is deployed in the wild at the same time, thanks to two testbeds: the LISP Beta Network and the LISP-Lab project. The interworking mechanism is proposed to ensure the communication between LISP-speaking sites and legacy Internet. The performance of LISP interworking with legacy Internet is of paramount importance to promote the adoption of LISP in future Internet. The existing research efforts to evaluate LISP interworking lack of large scale experimental results. The objective of this chapter is to fill this gap by providing a latency evaluation and routing path measurement about LISP interworking mechanism in real networks. The work is based on two experiments (one is conducted in 2015, the other one is launched in 2016) by using RIPE Atlas, which is the largest existing Internet measurement infrastructure for both IPv4 and IPv6. Experimental results show that LISP introduces additional latency, especially for close destinations, but negligible for intercontinental long-distance destinations. The additional latency highly depends on the position of the new network element being in charge of the communication between LISP-sites and the legacy Internet. Although LISP introduces some stretch, its performance is generally stable and its latency is not very high compared to natively forwarding without using LISP.

In the remainder, Sec. 6.2.1 and Sec. 6.2.2 introduce the necessary platform that our experiment leveraged on. Sec. 6.3 describes the methodology we used to conduct the experiment. Sec. 6.4 and Sec. 6.5 respectively present the IPv4 ping results obtained from 2015 and 2016. Sec. 6.6 shows the IPv6 ping results and Sec. 6.7 indicates the traceroute results for both IPv4 and IPv6. Finally, Sec. 6.8 concludes the chapter.

6.1 Introduction

To improve LISP and promote the deployment of LISP, implementations and large scale flexible experimental platforms are indispensable. At the moment of this writing, two LISP platforms are deployed world-wide. One is the experimental LISP Beta Network testbed [11] deployed in 2008, and the other one is LISP-Lab platform [15] open to external experimenters since 2015. Both of them have all LISP-required network entities and have been inter-connected. Although LISP has been evaluated in terms of scalability [106], stability [81], LISP evolution [77], and delay resolving the bindings between EIDs and RLOCs ([106], [44]), LISP is never analyzed from the aspect of the interworking with the legacy Internet at large scale. The performance of such interoperation is very important since it determines the deployment speed of LISP networks [53]. Hence, it is necessary to assess how such platforms integrate with the legacy Internet, evaluate the performance, and offer realistic experience to improve the testbeds themselves and provide hints to move the LISP technology forward.

We conduct two experiments: one spans 6 hours in 2015 and the other one lasts 15 days in 2016. The results confirm that the use of proxies to connect LISP and non-LISP sites introduce negative effects, which are important for the nearby destinations but can be ignored for the intercontinental long-distance destinations. It also shows that the selection of the proxy location is very important, since having them close either to the sources or to the destinations can decrease a lot the negative stretch. Although LISP introduces some overhead, the performance is quite stable for IPv4. However, the same conclusion does not hold for IPv6. Furthermore, we observe that the interworking performance of LISP-Lab is more reliable than LISP Beta Network.

6.2 Related Work

6.2.1 RIPE Atlas

RIPE Atlas [29] is the largest Internet measurement infrastructure, consisting of a global network of more than 9000 probes all over the world that measure Internet connectivity, reachability, and provide an understanding of the state of the Internet in real time. From early 2013, the hardware of the probe is a modified TP-Link wireless router (model TL-MR 3020) with a small USB thumb drive in it, but this probe does not support WiFi. Atlas also has more than 200 worldwide anchors, which are enhanced probes, offering more process capacity and sufficient bandwidth to support the larger number of measurements. Thus, the anchors are normally more stable, and can be used as reference. All the probes and anchors whenever

are set up, they automatically start executing a set of pre-defined measurements, called built-in measurements. The set of built-in measurements contain *ping*, *traceroute*, *DNS*, *SSL* and some *HTTP* types, mostly towards well-known targets such as DNS root servers, but also towards some of the RIPE Atlas infrastructure components. Besides, the probes and anchors also provide the user-defined active measurements with the same measurement types. Volunteers all over the world host these small hardware devices to actively measure Internet performance. To avoid malicious attacks, credits are required to launch experiments and the amount of required credits varies according to experiment types. In addition, the maximum allowed number of measurements towards the same target is 10. RIPE Atlas provides a set of RESTful API with which experiment campaign parameters (e.g., type of measurement, query source and destination, the duration and interval of experiment etc.) are passed to probes and measurement traces can be retrieved. Based on the available API, one can schedule experiment campaign in an automatic manner.

6.2.2 Alexa

Alexa [1] [2] is a website created by Amazon which provides commercial web traffic data, global rankings, and other information on 30 million websites. The analytic such as website rankings is based on the data collected by a toolbar developed by Alexa and installed within users web browser. The toolbar provides functions such as popup blocker, a search engine, etc. In early 2015, Alexa stated that there had been 10 million downloads of the toolbar. Alexa ranks sites based primarily on tracking a sample set of Internet traffic—users of its toolbar for the Internet Explorer, Firefox and Google Chrome web browsers. Due to its huge sampling space, the website ranking that it published is widely used to evaluation the popularity of websites.

6.3 Measurement Methodology

The mechanism of LISP interworking is presented in Sec.2.2 and Coras et Al. [44] describe how the use of PxTRs introduces a stretch in the path between LISP and non-LISP networks. Thus, it is necessary to conduct a large scale experiment in the real world to quantify this overhead. The objective is to answer three questions through such an experiment: how much is the negative impact? Under which conditions such stretch has an important impact on the performance? Conversely, under which situation such overhead is so small that it can be ignored? The follows describe how we conduct our experiments.

Table 6.1 Different configurations of probes in 2015

Name	using LISP	network type	probe/anchor
LISP-Lab	yes	academic	probe
mPlane	no	academic	probe
rmd	no	industrial	probe
FranceIX	no	industrial	anchor

6.3.1 For Dataset2015

As the purpose of our experiment campaign is to fully obtain the knowledge about the LISP interworking performance with legacy Internet, we deployed a probe (RIPE Atlas probe number is #22341) with IPv4 and IPv6 address on the LISP-Lab platform inside an academic institute in Paris, France to conduct the LISP-enabled active measurements. For IPv4, it uses both PETR and PITR of LISP-Lab to communicate with the legacy Internet, and the connection between the xTR of probe and PxTR is via a MPLS VPN. However, the MPLS tunnel did not support IPv6 at that time. Thus, we configure the ITR of the LISP-Lab probe natively forward the packets towards the Internet core for IPv6 targets (i.e., PETR is not used for IPv6). As a result, the IPv6 packets outgoing from this probe are not encapsulated into LISP packets and natively forwarded, but the returned packets still pass through the PITR of LISP-Lab platform.

The mPlane probe (#13842) resides in an academic network and uses the conventional routing. It allows to compare with non-LISP academic networks. While the rmd probe (#16958) resides in industrial network and also uses the conventional routing. It is chosen in order to compare with non-LISP and non-academic networks. Further, a stable probe with much more measurement capacity is necessary as a reference with all other probes. To this end, the only anchor in Paris named FranceIX (#6118) is selected. It resides in industrial network and does not use LISP. Thus, in our experiment, there are in total 5 probes used as sources to conduct the active measurements. Thus, in this experiment, there are in total 4 probes used as sources to conduct the active measurements. Tab. 6.1 shows the different configurations of probes.

We are now in a setup phase, where we start with a reduced number of destinations so to first setup the automated experiments. In this first experiment, the selected 4 probes ping to the top 50 Alexa sites every 10 minutes during 6 hours. However, from the experiment we find that there are 14 websites resolve to the same IPv4 addresses. We filtered them, so the assessment presented in Sec. 6.4 are analyzed by the results of 36 top popular websites. With the collected dataset, we evaluate the various performances leveraged on RTT.




Fig. 6.1 Locations of probes and anchor

Table 6.2 Location of LISP Beta Network PxTRs in 2016

Number	Continent	Country
1	Europe	Netherlands
2	Europe	Denmark
3	Europe	Norway
4	America	US
5	America	US
6	Asia	Japan

6.3.2 For Dataset2016

The new experiment still leverages on the LISP-Lab probe and RIPE Atlas infrastructure, but provides the following new contributions:

- Add another LISP probe connected to the LISP Beta Network as experimental source, to compare LISP-Beta Network and LISP-Lab.
- Enlarge the number of destinations to the first 500 most popular websites on the Alexa ranking [1].
- Consider the performance of IPv6 besides IPv4.
- Use *traceroute*, in complement to latency measurements, in order to have a deeper understanding of the observed behavior.
- Extend the experimental span to 15 days so to study potential periodicity of traffic.

Table 6.3 Configuration of two LISP probes

Probe	PETR in Lyon	PITR
LISP-Lab (IPv4)	Via MPLS Tunnel	Lyon (Via MPLS Tunnel)
LIP6 (IPv4)	Via MPLS Tunnel	LISP Beta Network
LISP-Lab (IPv6)	not used	Lyon
LIP6 for (IPv6)	Via BGP Routing	LISP Beta Network

In order to ensure comprehensiveness and accuracy, we need some other probes geographically close to each other having both IPv4 and IPv6 address as comparative items. Ideally, they are connected through both the (non-) academic network provider and (non-) LISP platform. Thus, we selected 4 other probes located in Paris and their locations are shown in Fig. 6.1. The LIP6 probe (#2403) also connects to academic network behind both LISP Beta Network and LISP-Lab platform. When this probe communicates with the hosts in legacy Internet, it uses PETR of LISP-Lab by configuration and chooses one of 6 PITRs of LISP Beta Network according to the BGP behavior. Depending on the BGP announcement to the legacy host, a different PITR is used. The more precise location of PITRs on LISP Beta Network is shown in Tab. 6.2: 3 are in Europe (Netherlands, Denmark and Norway), 2 are in US, and 1 is in Asia (Japan). Both LISP probes have a MPLS tunnel from their ITRs to the PETR at Lyon for IPv4, while for IPv6, the ITR of LIP6 probe is configured to use normal BGP routing to reach the PETR in Lyon (i.e., PETR is used but without MPLS VPN). The use of the MPLS tunnel IPv4 packets experience a shorter path compared to the IPv6 BGP-based one. Thus, theoretically the IPv4 packets sending from the LIP6 probe should arrive to PETR faster than those of IPv6. The configuration of two LISP probes are listed in Tab. 6.3.

Since the rmd probe is not configured with an IPv6 address, we replace it by the Gandi probe (#3141), which also resides in industrial network and also uses the conventional routing for both IPv4 and IPv6. Although it is a little further to the LISP-Lab probe than the rmd probe, it is the nearest probe having both IPv4 and IPv6. Since the mPlane probe and the FranceIX anchor satisfy all the requirements of this experiment, we still keep them as the references. Tab. 6.4 shows the different configurations of probes.

What we care about is the delay performance and the path of LISP interworking with legacy Internet. Thus, we rely on the ping and traceroute tools provided by RIPE Atlas. As destinations of our measurements we selected the 500 most popular websites (according to worldwide website ranking provided by Alexa [24]), which reply to IPv4 ping and traceroute. Among them, 122 websites are configured with IPv6 address. Thus, in our experiment,

Table 6.4 Different configurations of probes in 2016

Name	using LISP	network type	probe/anchor
LISP-Lab	yes	academic	probe
LIP6	yes	academic	probe
mPlane	no	academic	probe
Gandi	no	industrial	probe
FranceIX	no	industrial	anchor

500 IPv4 and 122 IPv6 addresses are the destinations of ping and traceroute. We develop a Python script to schedule the experiment campaign by using the API provided by RIPE Atlas. The parameters of experiment campaign are as follows: the experiment campaign lasts 15 days (from December 15th to 29th 2016). The 5 chosen probes ping and traceroute to 500 IPv4 and 122 IPv6 addresses. The intervals of ping and traceroute measurements are respectively 30 minutes and 60 minutes. We want to evaluate LISP interworking performance as frequently as possible, but each probe sequentially launches 622 traceroute measurements, which last more than 30 minutes. To avoid the heavy traffic burden and guarantee that all the measurements in a same experimental round can be finished before next round, we set the sampling interval of traceroute at 60 minutes. As a summary, the results presented in this journal come from the following 4 experiment campaigns:

- 5 probes ping to 500 destinations during 15 days with an interval of 30 minutes (IPv4).
- 5 probes ping to 122 destinations during 15 days with an interval of 30 minutes (IPv6).
- 5 probes traceroute to 500 destinations during 15 days with an interval of 60 minutes (IPv4).
- 5 probes traceroute to 122 destinations during 15 days with an interval of 60 minutes (IPv6).

From the experiment, we find that some websites actually resolve to the same IPv4 (34 over 500) or IPv6 (47 over 122) addresses. Further, for some destinations, at least one probe (mainly the LIP6 probe) does not get any responses of *ping* during the whole experiment (42 for IPv4 and 10 for IPv6). In this paper, we filtered the anycast destinations and only consider the destinations having the responses. Thus, after cleaning the traces for the above reasons, all the results of *ping* measurements used in the following subsections consists on the 5 chosen probes as the experimental sources and 424 IPv4 and 65 IPv6 responding addresses as the destinations. Differently from the ping measurements, for the *traceroute* dataset, all the successful *traceroute* responses from each probe are kept for further analysis.




Fig. 6.2 cdf of average RTT between different probes

6.4 IPv4 Ping Results from Dataset2015

Fig. 6.2 shows the cumulative distribution function (CDF) of average RTTs towards the selected 50 targets. Since the experimental destinations are located all over the world, the range of observed RTTs varies from few milliseconds (ms) to nearly five hundreds ms. When the RTT values are in the small range, especially when RTT is less than 50ms, the latency from LISP-Lab probe is always higher than the three others and the difference is mainly around 10ms. This latency difference is caused by the stretch introduced by the proxy technology (cf. Sec.2.2), since every packet sent between LISP-Lab probe and the legacy Internet has to pass through the LISP-Lab PxTR, which is located in Lyon (approximately 400km far away from the probe). As the RTT increases, actually the difference (surprisingly) decreases. When the RTT around 200ms is reached, all probes show basically the same RTT. Around such RTT values, the destinations concentrate in North America. Going further in the high range RTT values, i.e., more than 350ms, when destinations are mainly located in Asia, the probe in the LISP-Lab domain actually shows the lowest RTT. Thus, it indicates that the network connection between the LISP-Lab PxTR to the (asian) intercontinental destinations has a better performance that the stretch delay from LISP-Lab probe to PxTR can be ignored. We wanted to use traceroute to further investigate the reasons, but it cannot be natively used, since the LISP-encapsulation prevents its correct function and we are exploring new ways to find out what is happening.

Table 6.5 Correlation coefficient to FranceIX

	LISP-Lab	mPlane	rmd	FranceIX
Coefficient	0.9733	0.9784	0.9646	1.0

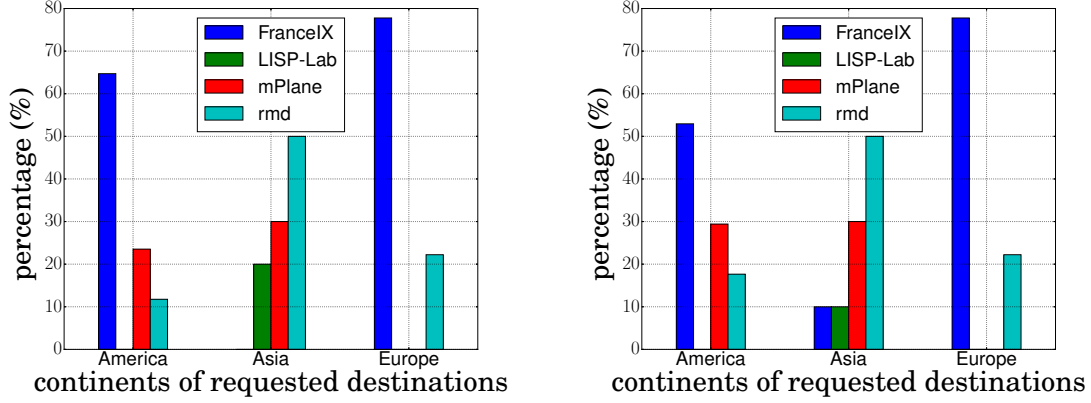


Fig. 6.3 Smallest mean RTT (left) and smallest median RTT (right) grouped by continent.

Further, we tried to quantify the percentage of times that one probe's RTT is the smallest compared to three others, since the destinations that the different probes contact to are exactly same. The result shows that in 52.8% of the time, FranceIX shows the smallest RTT to reach the destinations. It is normal that its RTTs are the smallest, since FranceIX is an Internet Exchange Point (IXP), hence well connected, and also acts as one of the anchors of RIPE Atlas, thus with a more powerful hardware. While, only in the 5.6% of the cases, the LISP-Lab probe has the smallest RTT. In this small percentage the main contribution comes from Intercontinental destinations, i.e., when the RTT values belong to high range.

Since FranceIX shows the best RTT most of the time, we tried to evaluate the correlation between the RTT of the other 3 probes and FranceIX. The purpose is to see if RTT measurements are correlated or totally independent. We compute the correlation coefficient for every RTT series of different probes to all destinations with the ones of FranceIX. Tab.6.5 shows the results. The absolute value of coefficient is 1 in the case of a perfect direct linear relationship, whereas 0 in the case of no correlation at all. The correlation coefficient of LISP-Lab is 0.9733 (> 0.8), showing a very high correlation with FranceIX, like the other 2 probes. This means that the LISP while certainly introducing some overhead, has quite stable performance that does not deviate from normal network operation.

Fig. 6.3 depicts the percentage of times that one probe's RTT is the smallest compared to three other probes grouped by continents where the selected targets are located. When the destinations are in Europe and America, FranceIX is the fastest most of the time. This is reasonable, since FranceIX is an Internet Exchange Point (IXP), acting as Atlas' anchor,

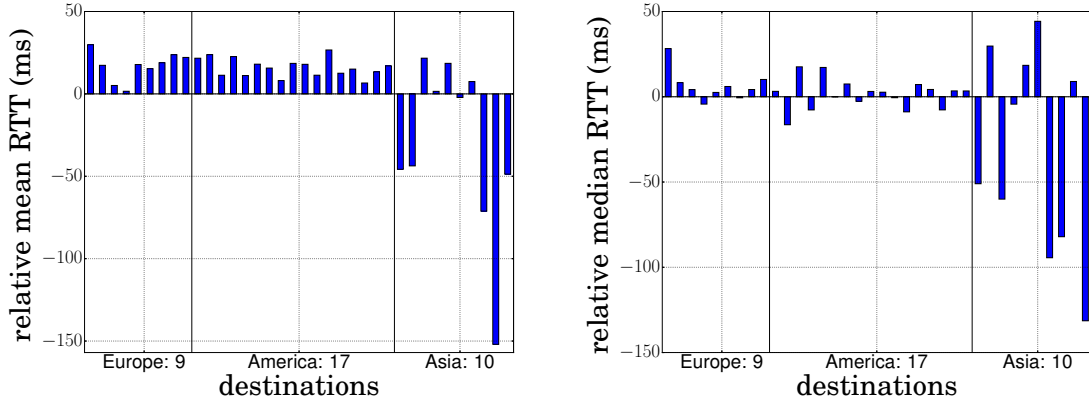


Fig. 6.4 Relative mean (left) and median (right) RTT clustered by different continents.

thus it is well connected with a more powerful hardware. Whereas, the percentage of LISP-Lab is always 0. Its higher RTT is caused by the proxy stretch, since traffic between LISP-Lab probe and the legacy Internet has to pass through the LISP-Lab PxTR, which is in Lyon (approx. 400km away from the probe). When the targets are in Asia, LISP-Lab becomes the fastest with a percentage of 20% (mean RTT) and 10% (median RTT). It indicates that such connection from the LISP-Lab PxTR is faster, so that the stretch can be ignored. The performance of FranceIX is not very stable to the Asian destinations, being the fastest 0% in average, but it is 10% looking at the median RTT. It shows that FranceIX sometimes has extremely high RTT values to Asian destinations.

We also evaluate whether the performance of LISP-Lab is as stable as FranceIX. We define a metric called *Relative RTT* ($rRTT$) for each destination as:

$$rRTT_{LL}(d) = RTT_{LL}(d) - RTT_F(d) \quad (6.1)$$

where d is the destination, subscriptions LL and F respectively refers to LISP-Lab and FranceIX. The results clustered by continents are shown in Fig. 6.4. The left one shows the mean RTT, while the right one shows the median RTT. For the European and American targets, LISP-Lab is a little slower than FranceIX but with a stable behavior. On the contrary, for half of the Asian destinations, LISP-Lab is significantly faster than FranceIX. It shows that the network connection between LISP-Lab PxTR and Asian destinations has better performance. Comparing the two subfigures, there is no negative values at all in Europe and America area in left figure, but there are some in the right one. It indicates that LISP-Lab RTTs to these destinations are very unstable and the variance is quite high.




Fig. 6.5 CDF of median RTT between different probes (IPv4)

6.5 IPv4 Ping Results from Dataset2016

With the collected dataset from IPv4 ping experiment, we assess the performance of LISP interworking with legacy Internet in terms of round-trip time (RTT).

Fig. 6.5 shows the Empirical Cumulative Distribution Function (ECDF) of the average RTT toward the selected IPv4 destinations. Because of the experimental targets are world-wide, the range of observed RTTs varies from few milliseconds (ms) to nearly 500ms. As an anchor, the FranceIX probe outperforms other probes in most cases, hence, showing the smallest RTT. With RTT range $[0, 200]$ ms, the latency from two LISP probes, LISP-Lab and LIP6, are respectively 25ms and 60ms slower than three other non-LISP probes. Such a RTT performance degradation (also called stretch) is caused by the path length stretch introduced by the proxy technology (as mentioned in Sec. 2.1), since every packet sent between LISP probe and the legacy Internet has to pass through a PxTR. Although both LISP probes introduce the stretch, LIP6 probe is even 50ms slower than the LISP-Lab probe. Given that both LISP-Lab and LIP6 probes pass through the same PETR located in Lyon (approximately 400km far away from the probe), the performance degradation of LIP6 probe is due to the PITR selection for traffic from legacy Internet to LISP network. The PITRs used by LIP6 probe belongs to the LISP Beta Network and there are in total 6 available PITRs. The PITRs announce the LISP Beta prefixes using BGP to their AS peers. So different destinations select different PITRs depending on where the destinations are and




Fig. 6.6 Smallest median RTT grouped by continent (IPv4).

lead to the reply goes through the different paths. For example, an Asian destination normally uses the PITR in Japan instead of selecting one in Europe. As listed in Tab. 6.2, LISP Beta Network has not deployed PITR in France, while the PITR of LISP-Lab is in the same country to its probe. Thus, the return path that we measured for the LIP6 probe is longer than the LISP-Lab probe. As a result, the closer the destination is located to the probes, the bigger difference on RTT between two LISP probes.

Within RTT range $[200, 500]$ ms, all probes except LIP6 show basically the same RTT. It indicates that LISP-Lab has a better performance in scenario of long distance network connection. The reason is that stretch delay from LISP-Lab probe to PxTR (in Lyon) can be ignored. Fig. 6.5 confirms this situation, within RTT range $[200, 500]$ ms, the difference between two probes becomes significantly smaller than that in range $[0, 200]$ ms.

The aforementioned results and analysis coincide with the Fig. 1 in [79], which just cover comparison related to LISP-Lab probe. This indicates that the LISP PxTR performance is stable: the measurement results do not change with the different destinations at the different time.

Even with stretch brought by PxTR, LISP probes still outperform non-LISP probes in some cases according to the measurement results. To identify these cases, worldwide destinations are classified by continent. Fig. 6.6 depicts the percentage of times that one probe has the smallest RTT compared to the other probes grouped by continents where the selected targets are located. The experimental destinations are located in 4 continents, i.e., Europe,

Table 6.6 Correlation coefficient to FranceIX (IPv4)

	Gandi	mPlane	LISP-Lab	LIP6	FranceIX
Coefficient	0.9647	0.9707	0.9766	0.8547	1.0

North America, South America, and Asia. The *ping* measurement is repeated 720 times as the experiment lasts 15 days and its interval is 30 minutes. The RTTs of the FranceIX anchor and mPlane probe are the smallest most of the time, especially for the European and South American destinations. Only for the Asian destinations the RTT of LISP-Lab probe experiences sometimes the smallest latency for 8.82% of the targets, while the LIP6 probe experiences the smallest latency in 0.98% of the cases. Among the North American destinations, there is just one destination for which the RTT of the LISP-Lab probe is the smallest, while there is no destination for which LIP6 experiences the smallest latency. Fig. 6.5 shows that for RTT above 200ms LISP-Lab performs like the other probes, so that we can question whether it actually has the best performance in some cases, and for which destinations. Fig. 6.6 shows the geographic location of destinations for which LISP-Lab actually is the best performing probe. In particular, the LISP-Lab probe has no additional delay compared to other non-LISP probes for the intercontinental transmission to the Asian destinations. It means that the negative effect of LISP stretch delay is evident for communications with European and American destinations, but can be neglected for Asian intercontinental destinations. This phenomenon is similar to the result shown in [78], which shows only in 5.6% of the cases that the LISP-Lab probe experiences the smallest latency and all these destinations are in Asia. It proves again that the negative effect of LISP PxTR can be ignored for the intercontinental transmission to the Asian destinations (when considering the European vantage point).

As FranceIX is an enhanced probe with more measurement capacities presenting more stable performance, we try to evaluate the correlation in term of RTT between FranceIX and other probes. The purpose is to see if RTT measurements are correlated or totally independent. The method is to compute the correlation coefficient for every RTT series of different probes to all destinations with the ones of FranceIX. The results are shown in Tab. 6.6. The absolute value of coefficient is 1 in the case of a perfect direct linear relationship, whereas 0 in the case of no correlation at all, i.e., independent. The correlation coefficient of LISP-Lab is 0.9766 (> 0.8), showing a very high correlation level with FranceIX, and even higher than the correlation coefficient of all the other probes. The correlation coefficient of LIP6 is 0.8547, shows a high correlation with FranceIX as well. Both of 2 LISP probes having a high correlation with FranceIX means that while certainly introducing some overheads, LISP has quite stable performance that does not deviate from normal network operation.




Fig. 6.7 IPv4 Relative median RTT clustered by different continents for LISP-Lab (left) and LIP6 (right)

Such result has two sides for the performance of LISP. Good because LISP shows the stability. Bad because it means that LISP experiences the same performance and failures like the normal Internet.

Another interesting point is to assess whether the two LISP probes have a stable performance as FranceIX according to the destinations. To this end, we leverage on the same metric of *Relative RTT* ($rRTT$) (formula 6.1), defined in Sec. 6.4 for each destination as:

$$rRTT_{LL}(d) = RTT_{LL}(d) - RTT_F(d) \quad (6.2)$$

$$rRTT_{L6}(d) = RTT_{L6}(d) - RTT_F(d) \quad (6.3)$$

where d is the destination, subscriptions LL , $L6$ and F respectively refers to LISP-Lab, LIP6 and FranceIX. The results clustered by continents are shown in Fig. 6.7, where the relative RTT for LISP-Lab probe is on the left and the relative RTT for LIP6 probe is on the right. The positive relative RTT indicates FranceIX faster, while the negative values indicating that LISP-Lab or LIP6 are faster. In the left side of Fig. 6.7, for the European and American targets, the values are between 10ms and 20ms in most cases (i.e., 31.37%), showing that LISP-Lab is a little slower than FranceIX but with a stable behavior. On the contrary, for some Asian destinations, LISP-Lab is significantly either faster (in 20.6% of the cases) or slower than FranceIX and the largest difference reaches 150ms. It shows as well that for the European and American targets, LISP-Lab probe is stable compared to FranceIX, since the difference is mainly caused by the transmission delay between Paris (location of probe) and Lyon (location of PxTR). But when considering the Asian destinations, LISP-Lab does not show a very stable behavior, where the difference might be very large, since the path

Table 6.7 Reliability of each probe (IPv4)

	Gandi	mPlane	LISP-Lab	LIP6	FranceIX
Reliability (%)	99.66	99.65	99.43	77.78	99.62

diversity for long-distance intercontinental transmission is higher, with every path having very different performance. The right figure of Fig. 6.7 shows that for LIP6 probe the difference is almost always positive. For the European destinations, the average relative RTT is around 130ms and much higher than those in the other 3 continents. All European targets show that the LIP6 probe is significantly slower than the FranceIX anchor. For the American destinations, most relative RTT decreases to around 50ms but some of them are around 150ms. For the Asian targets, the relative RTT presents big differences. Some still stay at around 150ms, some drop to below 50ms, and there are 5 destinations even show negative relative RTTs. The reasons causing the relative RTTs for LIP6 probe varying a lot are similar to that for LISP-Lab but not totally the same. The same point is that the longest transmission of traffic in our experiment are the destinations located in Asia. Since the stretch can be ignored in the case of intercontinental transmission, both of two LISP probes are not always slower than FranceIX to Asian targets. The different performance between LISP-Lab probe and LIP6 probe is caused by the location of each PxTR. Although the PxTR of LISP-Lab is not beside to its probe, but at least they are in a same country. While the LIP6 probe uses the same PETR to LISP-Lab probe but its PITR is not in France, and is even quite far away. As a result, for the shorter transmission, i.e., to the European destinations, the LIP6 probe introduces extremely higher RTTs than the non-LISP probes and even higher than LISP-Lab probe. Thus, the location of PxTR is very important. The PxTR near to either the sources or the destinations can obviously decrease the stretch. At least the probe having a PxTR in the same country shows a better performance compared to the PxTR just in the same continent.

We also want to evaluate the *Reliability* when having the response to each ping measurement. We measure reliability, in our experiment, by simply calculating the percentage of replies over the total number of requests. If every ping measurement is successful, i.e., there is a response having the RTT value for a probe at every experiment round, its reliability is 100%. As shown in Tab. 6.7, except for the LIP6 probe with 77.78%, all the other probes are more than 99%. The lower reliability of LIP6 probe indicates that sometimes its ping measurement is not successful. Since it is much lower than the others, to make sure that the LIP6 probe works well and there is no congestion or misconfiguration on the probe or any of its connected routers, we conduct another experiment letting the LIP6 probe use a normal public IPv4 address pinging the same 500 destinations. In this case, the LIP6 probe




Fig. 6.8 ECDF of median RTT between different probes (IPv6)

is non-LISP-speaking. The results show that the reliability of LIP6 probe is 98.91%, which is very close to the other probes. Thus, we eliminate the possibility of the problem on the LIP6 probe itself. As the LIP6 probe uses the same PETR to the LISP-Lab probe, while the reliability of LISP-Lab is very high, the difference is caused by the PITRs. The LISP-Lab probe uses the PITR of LISP-Lab platform but the LIP6 probe leverages one of LISP Beta Network. On one hand, LIP6 probe has longer transmission distance causing more risk of losing the packets and thus leads to more losses. On the other hand, we can conclude that the performance of LISP-Lab PxTR is more reliable than that of LISP Beta Network.

Since the experiment in 2015 just lasts 6 hours and shows no periodicity. Thus, this experiment campaign is extended to conduct with a span of 2 weeks to assess whether the RTT measurements have periodicity. Two methods are used to check: Fast Fourier transform (FFT) and Auto-correlation. However, the analysis results also show that none of RTT series from probes to any destinations exhibits periodicity. It means that the traffic does not periodically fluctuate with a certain interval.

6.6 IPv6 Ping Results

In this section, LISP interworking performance of IPv6 is evaluated with the same metrics as those used in Sec. 6.5.




Fig. 6.9 IPv6 Relative median RTT clustered by different continents for LISP-Lab (left) and LIP6 (right)

The ECDF of average RTTs to the selected IPv6 destinations is shown on Fig. 6.8, which is generally similar to the Fig. 6.5 for IPv4. The FranceIX anchor still shows the best performance with the smallest delay for almost all the destinations. The RTT of LISP-Lab probe is always higher than the other non-LISP probes with RTT range $[0, 30]ms$ and the difference is around 7ms. The RTTs are almost same when the values are more than 30. The phenomenon of IPv6 performance is similar to IPv4 that the stretch introduced by PxTR is significant when the range of RTT value is small, but can be ignored when the range of RTT value becomes large.

However, in the experiment for IPv6 targets, the traffic produced by LISP-Lab probe are natively forward to the destinations instead of going to PETR at Lyon first, thus the difference later than the other non-LISP probes becomes smaller compared to those for IPv4. But the latency still exists, caused by the returning traffic that still need to pass through a PITR. Since the latency decreases, it is easier to be ignored for IPv6. That is the reason why the stretch can be ignored when RTT values are higher than just 30ms for IPv6 while it needs to be higher than 200ms for IPv4. While the LIP6 probe is always the slowest compared to the other probes, the difference is even 160ms to the non-LISP probes and 150ms to the LISP-Lab probe. The reason having such a big difference is not only caused by the bidirectional traffic should pass through the PxTR, but also caused by the route between xTR of LIP6 probe to the PETR at Lyon, which has no tunnel indicating the used routing path longer than the case of having an MPLS tunnel. Similar to IPv4, when RTT values are higher than 250ms, the difference decreases to 40ms in most cases.

To further explore why two LISP probes are slower, especially to understand why the LIP6 probe has an extremely large latency compared to the others in detail, we leverage

on the Relative RTT between both LISP probes and the FranceIX anchor grouped in four continents as mentioned in Sec. 6.5. In the left figure of Fig. 6.9, which is the relative RTT between LISP-Lab and FranceIX, all the values are positive, indicating that the LISP-Lab probe is slower than the FranceIX anchor no matter to which destination. For the European and North American targets, the relative RTTs are mostly between 5ms and 17ms. But the average relative RTTs decrease to 7ms for the 25 Asian destinations. It proves the fact that for the nearby transmission, the delay introduced by PxTR is significant, while it effects less for the intercontinental transmission to the Asian destinations. What's more, the absolute relative RTT values for IPv6 are smaller than those for IPv4, because there is only returning traffic passing PITR for IPv6 so to reduce the latency. Thus, natively forwarding is faster than using PETR, but not so much. From the Fig. 6.9 we also know that there are 32 destinations located in Europe and just 7 in North America. It explains the reason why the curve of CDF for LISP-Lab sharply increases when the RTT is in the small range and mixes with the other non-LISP lines so quickly. It is because a half of destinations are not far away from the probes. Thus, a high percentage (around 50%) consists of small RTT values and these RTTs of LISP-Lab have a consistent delay compared to the other non-LISP probes. The mix part is mainly produced by the Asian destinations, to which the RTT values are large and the delay introduced by only the PITR of LISP-Lab is less significant. The relative RTT for the LIP6 probe is shown on the right of Fig. 6.9. The relative RTTs are all positive and much higher, indicating that the LIP6 probe is always quite slower than the FranceIX anchor to whichever target. Especially for the European destinations, the values are around 170ms. But to the North American and Asian targets, the relative RTTs are obviously smaller. As mentioned in Sec. 6.3.2, the outgoing packets from LIP6 probe still need to go to PETR first and especially the path between its xTR and PETR has no MPLS VPN tunnel. Besides, the main reason that the higher relative RTT values for IPv6 than IPv4 is caused by the LIP6 probe using the PITRs of LISP Beta Network. Different from IPv4 [3], there are only two ASes of IPv6 PITRs and both of them are located in US [4]. As a result, no matter from which destination, all the returning packets need to pass the PITRs in US first and then forward back to the LIP6 probe located in Paris. By consequence, the relative RTTs to the North American destinations are extremely smaller than those to European destinations. Although the traffic returning back from the Asian destinations also pass through the PITRs in US, the distance between the source and destination is quite far, thus it does not effect a lot to the relative RTTs for Asian targets.

Further, in the 84.61% of cases the FranceIX anchor has the smallest RTT in the IPv6 experiment. The LISP-Lab probe and LIP6 are not the fastest to any destinations, even to the Asian targets. It shows that the LISP performance of IPv6 is a little bit worse than

Table 6.8 Correlation coefficient to FranceIX (IPv6)

	Gandi	mPlane	LISP-Lab	LIP6	FranceIX
Coefficient	0.9565	0.968	0.967	0.3734	1.0

Table 6.9 Reliability of each probe (IPv6)

	Gandi	mPlane	LISP-Lab	LIP6	FranceIX
Reliability (%)	98.43	99.97	99.98	82.24	99.99

IPv4. Since FranceIX has the smallest RTTs in most situations and as an anchor, its higher measurement capacities lead to its higher stability, we also use the correlation between the RTT of the other 4 probes and FranceIX to see if RTT measurements of each probe are as stable as FranceIX. As shown in Tab. 6.8, the correlation coefficient of LISP-Lab is 0.967, which is almost one, indicating although LISP over IPv6 has higher latency than FranceIX caused by the introducing of PxTR, the performance is still stable. However, the correlation coefficient of LIP6 probe is just 0.3734. It is higher than 0.2, but much lower than 0.8, showing that the LIP6 probe is not totally independent from FranceIX, but it has quite low correlation. In fact, the IPv6 RTT series of LIP6 fluctuate a lot over time during the experiment and are not stable.

The Reliability of every probe for IPv6 is almost the same of IPv4 as shown in Tab. 6.9 except for the LIP6 probe. The higher reliability of LISP-Lab probe than the LIP6 probe confirms again that even for IPv6, the PxTR of LISP-Lab is still more reliable than that of LISP Beta Network, although the latter one has 6 world-wide PxTRs in total, from which only 2 can be used for IPv6. The reliability of LISP-Lab probe for IPv6 is a little bit higher than itself for IPv4, showing that using PxTR may decrease the reliability but the effect is very small. The reliability of the LIP6 probe for IPv6 is also higher than itself for IPv4, indicating that the IPv6 performance of PxTR on LISP Beta Network is better than IPv4, mainly thanks to the use of PxTRs in US. Thus from this comparison, we get a conclusion that the two PxTRs in US are more reliable than the other 4 PxTRs of LISP Beta Network.

The periodicity check for IPv6 is also conducted, but no probes show that their RTT measurements to any destinations have periodicity, either. It indicates that the latency of traffic has no relationship with the specified experiment time, i.e., the RTT measurements generally have the same results regardless the experiment time and duration. Thus, the results of these experiments are not occasional phenomenon with the special results, they are reproducible instead.




Fig. 6.10 Distribution of IPv4 Relative Hops Number for LISP-Lab (left) and LIP6 (right)

6.7 Traceroute-related Results

In this section, we use hops number obtained via traceroute from sources (i.e., probes) to destinations as the metric to further evaluate the LISP performance and also further investigate the reasons for the results presented in the ping-related experiments. Since the traceroute can only present the performance of the outgoing path, i.e., the routing path from probes to destinations, we focus on the use of the PETR and existence of a VPN between the xTR and the PETR.

When the IPv4 targets are in Europe and America, there are very few cases that the hops number of LISP-Lab or LIP6 is the smallest. Precisely, there is 7.65% of cases that LISP-Lab has the shortest path by hops number to the Asian destinations and 6.63% of cases that LIP6 has the shortest path in the experiment. This result coincides to the percentage of the smallest RTT shown in Fig. 6.6. This consistent relationship between smallest RTT and shortest path exists as well for the IPv6 experiment, where there is no destination at all for which the hop number or the RTT are the smallest (for both LISP-Lab and LIP6 probes). For IPv6, it is FranceIX or Gandi that always has the shortest path to all the targets.

As FranceIX is always the most stable probe and has the shortest path in most cases, we define *Relative Hops Number* (rHN) clustered by different continents to look at the difference between the hops number of LISP probes and FranceIX. The definition is as:

$$rHN_{LL}(d) = HN_{LL}(d) - HN_F(d) \quad (6.4)$$

$$rHN_{L6}(d) = HN_{L6}(d) - HN_F(d) \quad (6.5)$$




Fig. 6.11 IPv4 Relative hops number clustered by different continents for LISP-Lab (left) and LIP6 (right)

where d is the destination, subscriptions LL , $L6$ and F respectively refers to LISP-Lab, LIP6 and FranceIX. The positive value indicates that the hops number of FranceIX is smaller than LISP probes. Fig. 6.10 provides an IPv4 distribution of relative hops number for LISP-Lab on the left and for LIP6 on the right. This figure shows the probability of each relative hops number. For the LISP-Lab probe, the most common relative hops number is 12 with a percentage of 21.77%, hence remaining limited in most of the cases. From the ping-related experiment in Sec. 6.5 and Sec. 6.6 we know that using PxTR introduces some overhead. So the Relative Hops Number being generally positive is reasonable. The very large relative hops number just appears once for 19 and 25, so we regard them as outliers. Similar to the LISP-Lab case, for the LIP6 probe, the most frequent relative hops number is 13 in 19.53% of cases.

In order to understand where the most common relative hops number appears, we produce the relative hops number for each destination clustered by different continents in Fig. 6.11 to complete Fig. 6.10. The left figure is for the LISP-Lab probe, indicating that the relative hops number is almost the same for the European and North American destinations and contributes to the relative hops number with high probabilities in Fig. 6.10. The small relative hops number is mostly produced from the Asian targets. Whereas the outliers are all from the European destinations. The right figure of Fig. 6.11 is for the LIP6 probe, mainly presenting the same result as the LISP-Lab probe, except that the latter rarely has the relative hops number higher than 16, but LIP6 has more.

With the same evaluation metrics of IPv4, the distribution of Relative Hops Number and the Relative Hops Number clustered by different continents for IPv6 are respectively shown in Fig. 6.12 and Fig. 6.13. The left figure of Fig. 6.12 presents that the most common relative




Fig. 6.12 Distribution of IPv6 Relative Hops Number for LISP-Lab (left) and LIP6 (right)

hops number between the FranceIX anchor and the LISP-Lab probe is 7 in 56.72% of cases, more than a half, much higher than the probability of other relative hops numbers. The range of relative hops number is smaller compared to IPv4 and the most common relative hops number is also smaller, since the traffic does not need to pass the PETR, but is just natively forwarded so the overhead is reduced. The left figure of Fig. 6.13 reveals the truth that 7 relative hops numbers are produced by the majority of targets regardless where the destinations are, instead of like the one of IPv4, where the relative hops number for European and American destinations are more than the Asian's. As the relative hops number from PETR is only 3, but 7 for natively forwarding, it is likely that the PxTR has shorter paths to most destinations, leads to the decreasing of hops number for the LISP-Lab probe. The relative RTTs higher for European and American destinations than those for Asia shown in left figure of Fig. 6.9 are mainly caused by the returning path, since it is natively forward for outgoing and there is no difference by continents. For IPv6, the LIP6 probe has extremely high relative hops number in a range from 12 to 24. The majority of relative hops number is 23 in 19.57% of the cases and followed by 13 with a percentage of 17.39%. The higher relative hops number is caused by not using the VPN between xTR and PETR compared to the IPv4 case. The right figure of Fig. 6.13 visualizes the 23 relative hops number coming from the European destinations and 13 coming from the Asian targets. The form coincides to the relative RTT for IPv6 shown in the right hand part of Fig. 6.9, indicating that the high relative RTTs are caused by the high relative hops number.

Concerning the RTT performance of the LISP-Lab probe to Asian destinations, which sometimes shows even better than FranceIX, we analyze the AS-path (Autonomous System path) from FranceIX and LISP-Lab to Asian destinations that the packets traverse. Since all the IPv4 packets from LISP-Lab probe are sent to PETR at Lyon first, so the AS-path




Fig. 6.13 IPv6 Relative hops number clustered by different continents for LISP-Lab (left) and LIP6 (right)

discussed in the followings is actually from PETR to the Asian destinations. By comparing the AS-path, we find that FranceIX and LISP-Lab very often take different paths, with only the last 1 or 2 hops in common. Even further, by looking at the geographic location of each hop of *traceroute*, it can be observed that FranceIX and LISP-Lab take the path even in the different directions in most of the cases. In particular, there are 11 destinations out of 117 in total (i.e., 9.4%) for which the packets sent by LISP-Lab pass through US, while the packets sent by FranceIX pass through Eastern Europe. There are 68 destinations (i.e., 58.1%), which show exactly the opposite situation, i.e., packets sent by FranceIX pass through US, while packets sent by LISP-Lab pass through Eastern Europe. This last case is when LISP-Lab has smaller RTT (compared to FranceIX), almost all of the times (with only few exceptions).

6.8 Conclusion

In this chapter, we conduct a six-hours as well as a two-weeks real network measurements with RIPE Atlas to provide a comprehensive performance evaluation of LISP interworkings. Concretely, we provide a first thorough sight on the performance of LISP PxTR. Since the results of experiment in 2016 coincides to the one in 2015 and are more comprehensive. Thus, we conclude this chapter by mainly presenting the observations of the Dataset 2016.

In our large scale measurement campaign, we take into account 5 probes as sources, 500 IPv4 and 122 IPv6 addresses as destinations, conducting ping and traceroute experiments. We find that the PxTR indeed introduces the negative effects for the destinations located in Europe and America, but the negative impact of PxTR can be ignored for the intercontinental

long-distance transmission to Asia destinations. From the experiment, the results show that the position of PxTR is very important. The PxTR either near to the sources or the destinations can decrease the latency a lot. Generally speaking, LISP is stable compared with the reference anchor FranceIX, except for the IPv6 performance of LISP Beta Network. Further, the performance of LISP-Lab PxTR is more reliable than the one of LISP Beta Network, although the latter has 6 worldwide PxTRs used for IPv4 and 2 located in US used for IPv6, whereas LISP-Lab has only 1 PxTR for both IPv4 and IPv6. Compared to leveraging on PxTR, natively forwarding without using LISP decreases the latency, but not much. The traceroute experiment shows that introducing PxTR of course brings more hops, but if the PxTR is well configured, so to always have peers to the destinations, there are only 4 more hops compared to the packets being natively forward by xTR without encapsulating with LISP.

In our experiment, as LIP6 probe does not receive any *ping* responses for 42 IPv4 and 10 IPv6 destinations, we plan to cluster these destinations by geography and explore the reasons. Besides, we try to modify the configurations so that PxTR of LISP-Lab allows to receive the IPv6 traffic from LISP-Lab probe. The aim is to evaluate thoroughly IPv6 performance of LISP-Lab platform. Furthermore, as the degradation on performance observed in our paper is mainly due to the PxTR selection by the destinations. It will be interesting to explore more the interaction between BGP anycast and LISP interworking.

Chapter 7

ns-3 Implementation of LISP and LISP-MN

The *Locator/Identifier Separation Protocol* (LISP) reconstructs the current IP addressing space to improve the scalability and no interrupt mobility issues. LISP Mobile Node (LISP-MN) is based on the basic LISP functionality to provide seamless mobility across networks. The basic LISP architecture is deployed on LISP Beta Network and LISP-Lab platform to offer the researchers a realistic experimental environment, but both do not support LISP-MN. Some simulation models with LISP extensions are implemented on various simulators, but are unfortunately not open source. Providing a free and flexible simulation model with the basic LISP architecture as well as the extensions so to help researchers quickly test new LISP behaviors motivates our work. This paper introduces the implementation of the basic LISP architecture model and LISP-MN on the simulator ns-3. It also provides the evaluation results in mobility scenario to validate the model and shows when the current proposal of LISP-MN is behind a LISP-site has a very high delay during the handover procedure.

The rest of chapter is organized as follows: Sec. 7.1 reviews LISP architecture, highlighting the LISP mobility mechanism; Sec. 7.2 ... Sec. 7.3 analyzes the design and implementation of our prototype, and afterwards, Sec. ?? presents preliminary evaluation results of our implementations. We finish with a discussion and some ideas for future work.

7.1 Introduction

The *Locator/Identifier Separation Protocol* (LISP) [47] is initially proposed to solve the scalability and flexibility issues of current Internet architecture, since it decouples the Routing Locators (RLOCs) and Endpoint Identifiers (EIDs). This allows that the BGP routing

tables in Internet core just announce the globally routable RLOCs whereas EIDs are only locally used within the subnets. LISP is under standardization at the IETF around ten years and as it develops, more advantages are found, such as: seamless mobility at terminal or in the Data Center, IPv6 transition and traffic engineering [110]. Among them, the continuous communication without interruption during the handover of terminals becomes a hot topic recently.

To test the various features of LISP, two LISP testbeds are built: LISP Beta Network and LISP-Lab platform. But both of them are only the basic LISP architectures and do not support mobility functionality. Some platform-specific LISP implementations, such as: OpenLISP [100], Open Overlay Router (OOR) [18] and Cisco's implementation [13], can provide realistic LISP evaluations but lack the flexibility to test the new features. A LISP simulator is a good choice for the researchers. To the best knowledge, LISP is first implemented on ns-3 [21], which is widely used by academic researchers and educational use in [?], but mobility support is not included as well. While within OMNET++ platform, LISP and mobility features are implemented [74]. However, their source code is not available online. To promote the development of LISP and as ns-3 gradually shows the potential to take the momentum in network simulator domain [102], it highlights the importance to implement LISP and its mobility extensions under ns-3 and motivates our work.

In this paper, we introduce the implementation of basic LISP functions and mobility features on ns-3 by both modifying the existent ns-3 modules and integrating the new modules. Our simulator is validated by achieving a mobility scenario, in which a terminal changes the attachment points by keeping the communication with the remote node. All the handover procedures are transparent to the terminal, however, the delay to receive the packets again presents high.

7.2 Related work

ns-3 [21] is a popular and free discrete-event network simulator for networking research. To be closer to the real implementation (in a real Operating System), easily include C-based implementation codes, ease debugging and reduce the cost on maintaining in a long term, C++ is prioritized to be the unique programming language for ns-3. Besides, ns-3 offers the possibility to visualize the simulation instance so to allow the users to visually confirm the packets flow as they expect.

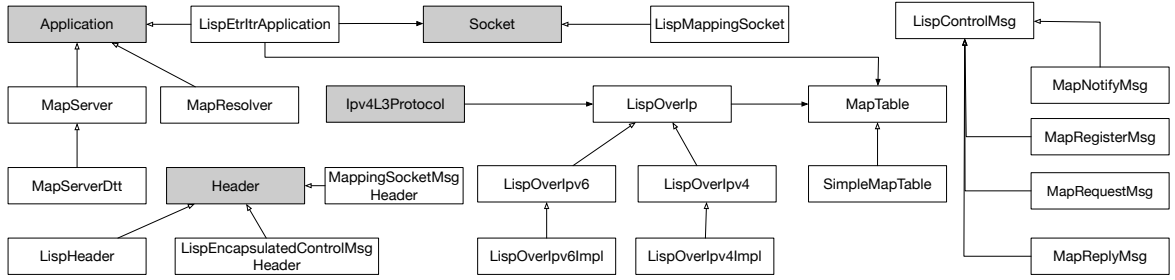


Fig. 7.1 UML diagram of LISP/LISP-MN implementation. The solid arrow refers to a composition relation, while the blank one refers to an inheritance relation.

7.3 LISP/LISP-MN Implementation based on ns3

It should be noted that this work is based on a prior implementation work by Lionel. However, his implementation has no support for LISP MN.

Our implementation is under ns-3.26 and based on LISP [?] and LISP-MN standards [?]. The main classes are shown in Fig. 7.1 in form of UML diagram. The blank blocks refer to the classes that we added into ns-3, while darker blocks are classes already in ns-3. As a design choice, we implement LISP/LISP-MN functionalities by modifying and extending the *internet* module of ns-3, instead of creating a new independent module. The justification of this design is that LISP/LISP-MN and legacy *internet* module have an interdependent relationship. However, this kind mutually dependent relationship between modules is not supported by ns-3. Inspired from design of OpenLISP [108], the Data Plane implementation is in "kernel space" (i.e. ns-3's *TCP/IP stack*) and Control Plane is implemented in "user space" (i.e. ns-3 *Application*). The communication between LISP Data and Control Plane is developed via a dedicated socket (i.e. *LispMappingSocket*) that inherits from ns-3 *Socket* class. It should be noted that our implementation only supports IPv4 at time of this writing. The IPv6 support (i.e. the implementation related to IPv6 such as *LispOverIpV6Impl*) is still in process.

7.3.1 Implementation of LISP Data Plane

A LISP-compatible node (terminal or router) should be capable of determining whether a packet should be passed to LISP-related procedure and retrieving the associated mapping information if necessary. To this end, a new class called *LispOverIp* and its extended classes (refer to Fig. 7.1) are added to ns-3 *internet* module. This class is in charge of checking whether necessary to do LISP-related operations (*NeedEncapsulation()*, *NeedDecapsulation()*), and encapsulating conventional IP packets (i.e., *LispOutput()*) as well as decapsulating LISP packets (*LispInput()*). It also contains a smart pointer pointing to the LISP database

and LISP cache. Both data structures (store the EID-RLOC mapping information) are represented by class *SimpleMapTable* that inherits from *MapTable*. The inheritance mechanism allows other users to implement their own implementation of LISP database and cache. To support LISP functionalities, the *Ipv4L3Protocol*, which is the IP layer implementation in ns-3, contains one *LispOverIp* object and *Ipv4L3Protocol*'s packet transmission and reception procedures are accordingly adapted.

To process outgoing packets, the adapted *send()* in *Ipv4L3Protocol* first verifies whether the *LispOverIp* object is present. If yes, some checks are then conducted to determine that this packet should be processed by *LispOutput()* (to encapsulate the packets) or by conventional packet transmission routine. For example, if both source and destination IP address of this packet belong to the same network, the LISP-related process (e.g., encapsulation) is skipped and this packet is processed as in a non-LISP network. Otherwise, EID-RLOC mapping information is searched from LISP cache and LISP database on LISP-MN node. In case of cache missing (i.e. destination RLOC is not found in the cache), the packet is dropped and *SendNotifyMessage()* in *LispOverIp* notifies the *LispEtrItrApplication* that runs on LISP-MN node via a *LispMappingSocket* socket. Once reception of cache missing event from LISP Data Plane (i.e. *LispOverIp* object), *LispEtrItrApplication* initiates a Map-Request message to LISP mapping system. Once reception of Map-Reply, the received EID-RLOC mapping is inserted into LISP cache. It should be noted that as an implementation choice, before the reception of Map-Reply message, all transmitted packets with the required RLOC as destination are dropped. One can also design a buffer to queue these packets and resend them once the required mapping information is received via Map-Reply. The advantage of such an implementation is to reduce the packet loss rate.

For an incoming packet, if the destination of this packet is the node itself, the packet is processed by *LocalDelivery* method in *Ipv4L3Protocol*. Before passing to transport layer, *LocalDelivery* checks if the packet should be decapsulated. If yes, it is passed to *LispInput* method, in which the packet is decapsulated and reinjected in the IP stack. If the received packet destination is not this node, this packet is processed by patched *IpForward* method. This packet may be ended up with LISP encapsulation procedure.

7.3.2 Implementation of LISP Control Plane

The implementation of LISP Control Plane at least should provide ITR/ETR, MR and MS. In practice, ETR and ITR functionalities are usually placed on a same router. In our implementation, they are included into class *LispEtrItrApplication*. A ns-3 node that runs *LispEtrItrApplication* is a LISP-compatible router. It should be able to communicate with *LispOverIp* on the same node (e.g. inform cache missing event) and other LISP-compatible

routers (e.g. Map-Request/Map-Reply). To support LISP-MN feature, *LispEtrItrApplication* also communicates with DHCP client application. For example, once a LISP-MN obtains an IP address from DHCP server, *LispEtrItrApplication* receives the corresponding EID-RLOC mapping and sends a Map-Register message [?].

A node that runs a *MapServer* application is the MS in a LISP-supported network. This class maintains a LISP database to store the EID-RLOC mapping information, learned from Map-Register message at the initialization stage. In current implementation, the role of MR is to receive the Map-Request message from xTR and forward it to the MS.

7.3.3 Integration of TUN net interface card

To support mobility, LISP-MN actually can be regarded as a small LISP-Site, in which xTR functionalities and DHCP service are implemented, as well as configured address of MR and MS. As a LISP-MN node, it has a static permanent EID and dynamic RLOC assigned by the DHCP server. To differentiate with conventional RLOC of xTR interface, such kind of RLOC is referred to as the local RLOC (LRLOC). Different from conventional LISP node, at least two net interface cards (NIC) are installed into LISP-MN. One is *WifiNetDevice*, the other is a TUN type card. The DHCP client application runs on LISP-MN's *WifiNetDevice* and thus the LRLOC is allocated to this card. The permanent EID is assigned to *Virtual-NetDevice* net card. We modify the node's routing table so that each packet's inner header contains IP address on *TunNetDevice* and outer header contains IP address of *WifiNetDevice* as source address.

7.3.4 Integration of DHCP

To support mobility within conventional LISP node, a modified DHCP client application is integrated into ns-3 node. To be compatible with LISP functionality, DHCP client application is modified. Once the DHCP client receives an allocated IP address (i.e. LRLOC), it notifies the *LispEtrItrApplication* (i.e. LISP Control Plane) by sending a dedicated message that contains the EID-LRLOC mapping. *LispEtrItrApplication* is in charge of populating the received mapping entry into LISP database. During the mobility process, when wireless link is down, the DHCP client flushes the LISP-MN database and populates the database again once reception of new LRLOC.

7.4 Evaluation

We validate our LISP/LISP-MN implementation by conducting a simulation and providing a preliminary performance evaluation of handover delay. The simulation topology is shown in Fig. 7.2.




Fig. 7.2 LISP mobility simulation scenario for double encapsulation

7.4.1 Simulation Setup

In our simulation, a LISP-MN with permanent EID 172.16.0.1 is initially placed in network 10.1.1.0/24. An *echo* application on LISP-MN sends one packet per second to a remote stationary node CN with EID 10.3.3.2, and the LISP-MN moves into network 10.1.7.0/24 at speed of 7.07m/s. The distance between xTR_1 and xTR_2 is 170m. LISP-MN node uses Wi-Fi to connect to xTR_1. At a certain moment during the moving, the Wi-Fi link between LISP-MN and xTR_1 is down, which triggers the handover procedure. Afterwards, LISP-MN connects to xTR_2 and reestablishes the communication with CN node. The total simulation time is set to 45s and the DHCP procedure delay is set to 1s. We conduct many times of simulations with the various beacon interval of Wi-Fi channel in the range of 0.05s to 2s.

7.4.2 Results

As shown in Fig. 7.3, when MN sending the packets to CN during the simulation, it needs double encapsulation and the packet flow sequences are as follows. The traditional IP packets with EID 172.16.0.1 as source address and EID 10.3.3.2 as destination address are encapsulated by adding Local RLOC 10.1.1.1 as outer source address and RLOC 10.1.5.2 as outer destination address after MN querying the mapping information to MR. The LISP

packets are encapsulated and forwarded to xTR_1. The latter gets the mapping information from MR and encapsulates the packets again by adding the RLOC 10.1.2.1 as the outer source address and RLOC 10.1.5.2 as the outer destination address, then sends the packets on Internet core. The xTR_3 needs dencapsulate the packets twice after it receiving and verifying the packets, and finally sends to CN. Once LISP-MN lost its Wi-Fi connection with xTR_1, it needs a DHCP procedure (consisting of DHCP Discover, DHCP Offer, DHCP Request and DHCP ACK) with xTR_2 to get a new LRLOC 10.1.7.1 and then triggers LISP SMR to xTR_3. After xTR_3 getting new mapping information of <10.1.7.1, 10.1.3.1> and <172.16.0.1, 10.1.7.1>, the connection between LISP-MN and CN is re-established again.

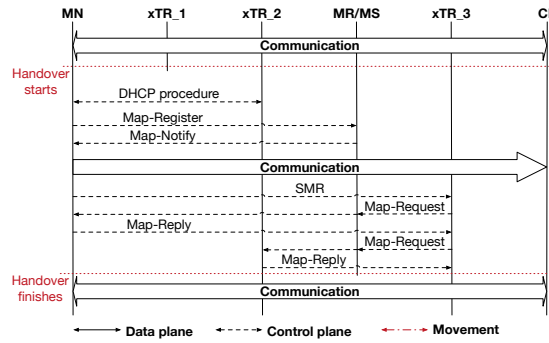


Fig. 7.3 Schema for LISP-MN mobility

The overall handover delay in this paper is defined at the moment that LISP-MN sending DHCP Discover message to xTR_2 and end up with xTR_3 receiving the last Map-Reply from xTR_2. Precisely, it is composed by three parts: the Wi-Fi association delay, the DHCP related delay and LISP SMR delay:

$$D_{overall} = D_{Wi-Fi}(BI) + D_{DHCP} + D_{SMR}$$

where D is the delay, BI is Beacon Interval, subscriptions $Wi-Fi$, $DHCP$ and SMR respectively refers to Wi-Fi association, DHCP procedure and LISP SMR.

After several executions of simulation program, we observe that the overall handover delay changes by the various beacon intervals, in particular the Wi-Fi association delay depends on the different beacon intervals, whereas LISP SMR procedure always cost around 3s. To get the lower bound of overall handover delay, we can ignore the Wi-Fi association delay when the beacon interval is 500ms, and the latency due to DHCP procedure is always 1s. Thus, adopting LISP-MN to conduct the host-based mobility takes at least 4s. Compared to current most stable solution for host-based IP mobility management MIPv6, which latency including L2 and L3 in a real Wi-Fi testbed is around 3.68s [?], LISP-MN

has a higher delay caused by the double encapsulation mechanism introduced by LISP-MN behind LISP-Site.

During handover, CN can successfully receive packets from LISP-MN right after DHCP procedure being accomplished, but LISP-MN cannot receive the packets from CN until LISP SMR procedure is also finished. Thus, during DHCP procedure, all bi-directional transmitted packets are lost. To improve the performance, [112] proposes a network-level LISP-MN solution, but has not validated their proposals neither in simulation nor in testbed. Our ns-3 implementation can be used to realize them.

7.5 Conclusion

As a promising technology for the future Internet architecture, LISP attracts more and more attention. There exist some LISP implementations, but they do not support LISP-MN or they are proprietary. Further, although measurements on LISP-testbeds can provide real time performance, due to the complicated topological structure, it is somewhat like a black box test which hinders us to find the exact explanation for some results. This highlights the importance to have an open source simulator for LISP in particular to support LISP-MN functionality. In this paper, we present our implementation for LISP/LISP-MN within ns-3, since the latter is a largely accepted simulator in networking research. The simulation results show that our implementation works well, and reveal the current LISP-MN proposal with a double encapsulation that has an high level delay during handover procedure. Our simulator can be a perfect choice to test the improvements of LISP-MN.

There are two possible directions to support IP mobility in LISP: host-based (i.e. LISP-MN) and network-based (i.e., xTR) mobility. We can compare the performance between LISP double encapsulation described in this paper with only host supporting LISP and only router supporting LISP leveraging our proposed simulator. As Map-Versioning [67] is another Mapping Cache update mechanism, we can also compare the performance between it and SMR that we present in this paper by our simulator.

Chapter 8

Conclusion and Future Work

In this chapter, we summarize our major contributions and discuss future research directions.

8.1 Major Contributions

8.1.1 Evaluation of LISP Mapping System

Describe the main observations of Chapter. 4.

8.1.2 LISP-Views: LISP Mapping System Monitor

Describe the main observations of Chapter. 5.

8.1.3 Assessing LISP interworking performance through RIPE Atlas

Describe the main observations of Chapter. 6.

8.1.4 ns-3 Implementation of LISP and LISP-MN

Describe the main observations of Chapter. 7.

8.2 Discussion & Future work

- Evaluation of LISP Mapping System (Chapter. 4)
 - Evaluate the latest *Stability* and *Consistency* performance of LISP mapping system.

- Compare to the results in Chapter. 4 and get LISP developing trend.
 - Synchronized mechanism should be proposed to the MDS.
- LISP-Views: LISP Mapping System Monitor (Chapter. 5).
 - The implementation of a REST API for LISP-Views.
 - The deployment of LISP-Views on multiple VPs.
 - Test of IPv6 behavior for LISP-Views.
- Assessing LISP interworking performance through RIPE Atlas (Chapter. 6).
 - Evaluate LISP IPv6 performance for LISP-Lab probe.
 - Compare LISP IPv6 performance between LIP6 probe connecting to PETR with MPLS or not.
- ns-3 Implementation of LISP and LISP-MN (Chapter. 7).
 - Comparison the performance between only host supporting LISP, only router supporting LISP and both of them supporting LISP (LISP double encapsulation).
 - Comparison the performance between SMR and Map-Versioning.

References

- [1] Alexa. <http://www.alexacom/>.
- [2] Alexa top sites ranking. <https://www.alexacom/topsites>.
- [3] As3943 ipv4 route propagation. http://bgp.he.net/AS3943#_graph4.
- [4] As3943 ipv6 route propagation. http://bgp.he.net/AS3943#_graph6.
- [5] AVM's FRITZ!Box Configuration and Operation. See https://en.avm.de/fileadmin/user_upload/EN/Manuals/FRITZ_Box/Manual_FRITZ_Box_3490-en.pdf.
- [6] Cisco Locator/ID Separation Protocol (LISP) Q&A. See https://www.cisco.com/c/en/us/products/collateral/ios-nx-os-software/locator-id-separation-protocol-lisp/qa_c67-582925.pdf.
- [7] FreeBSD (website). <https://www.freebsd.org/>.
- [8] INET Framework (Homepage). See <https://inet.omnetpp.org/>.
- [9] Internet live stats. Accessed on Jan. 7, 2018. [Online]. Available: <http://www.internetlivestats.com/internet-users/>.
- [10] LIP6-LISP (github). <https://github.com/lip6-lisp>.
- [11] LISP Beta Network. See <http://www.lisp4.net>.
- [12] LISP Beta Network Latest Architecture. See <https://www.lisp4.net/images/lisp-ddt.pdf>.
- [13] Lisp cisco ios implementation (website). <http://lisp.cisco.com>.
- [14] LISP for SDN and NFV (OpenDaylight Submmit). See https://events.linuxfoundation.org/sites/events/files/slides/LISP_ODLSummit_2014.pdf.
- [15] LISP-Lab. See <http://www.lisp-lab.org/>.
- [16] LISP-Views (website). See <http://hunter.lisp-views.org/>.
- [17] Lispers.net. See <https://www.lispers.net/>.
- [18] LISPmob (website). <https://www.openoverlayrouter.org/lispmob/>.
- [19] LISPmon. See <http://www.lispmon.net>.

- [20] Locator ID Separation Protocol (LISP) VM Mobility Solution, (White paper). See https://www.cisco.com/c/dam/en/us/products/collateral/ios-nx-os-software/locator-id-separation-protocol-lisp/lisp-vm_mobility_wp.pdf.
- [21] ns-3. See <http://https://www.nsnam.org/>.
- [22] OMNet++ (Homepage). See <https://www.omnetpp.org/>.
- [23] Open Overlay Router (website), the successor of LISPmob. See <https://www.openoverlayrouter.org/>.
- [24] OpenDaylight (website). See <https://www.opendaylight.org/>.
- [25] OpenLISP (website). <http://www.openlisp.org/>.
- [26] Planetlab (website). <https://www.planet-lab.org/>.
- [27] pylisp (Github). See <https://github.com/steffann/pylisp>.
- [28] pylisp (Reference). See <https://www.biostat.wisc.edu/~annis/creations/PyLisp/pylisp.html>.
- [29] RIPE Atlas. <https://atlas.ripe.net/>.
- [30] Smartinsights. Accessed on Jan. 7, 2018. [Online]. Available: <http://www.smartinsights.com/mobile-marketing/mobile-marketing-analytics/mobile-marketing-statistics/>.
- [31] J. Adan. Tunneled inter-domain routing (tidr). *Work in Progress*, 2006.
- [32] M. Aiash. A novel security protocol for resolving addresses in the location/id split architecture. In *International Conference on Network and System Security*, pages 68–79. Springer, 2013.
- [33] M. Aiash, A. Al-Nemrat, and D. Preston. Securing address registration in location/id split protocol using id-based cryptography. In *International Conference on Wired/Wireless Internet Communication*, pages 129–139. Springer, 2013.
- [34] R. Atkinson, S. Bhatti, and S. Hailes. Evolving the internet architecture through naming. *IEEE Journal on Selected Areas in Communications*, 28(8):1319–1325, 2010.
- [35] H. Balakrishnan, K. Lakshminarayanan, S. Ratnasamy, S. Shenker, I. Stoica, and M. Walfish. A layered naming architecture for the internet. In *ACM SIGCOMM Computer Communication Review*, volume 34, pages 343–352. ACM, 2004.
- [36] S. Barkai, D. Farinacci, D. Meyer, F. Maino, V. Ermagan, A. Rodriguez-Natal, and A. Cabellos-Aparicio. Lisp based flowmapping for scaling nfv. Technical report, Internet-Draft draft-barkai-lisp-nfv-10, Internet Engineering Task Force. Work in Progress, 2017.
- [37] S. Bhatti and R. Atkinson. Identifier-locator network protocol (ilnp) architectural description. 2012.

- [38] L. Breslau, D. Estrin, K. Fall, S. Floyd, J. Heidemann, A. Helmy, P. Huang, S. McCanne, K. Varadhan, Y. Xu, et al. Advances in network simulation. *Computer*, 33(5):59–67, 2000.
- [39] S. Brim, N. Chiappa, D. Farinacci, V. Fuller, and D. Lewis. Lisp-cons: A content distribution overlay network service for lisp. draft-meyer-lisp-cons-04, 2008.
- [40] A. Cabellos and F. Coras. Internet engineering task force (ietf) d. saucez request for comments: 7834 inria category: Informational l. iannone. 2016.
- [41] A. Cabellos, A. R. Natal, L. Jakab, V. Ermagan, P. Natarajan, and F. Maino. Lispmob: mobile networking through lisp. *LISPmob white paper*, 2011.
- [42] X. Chang. Network simulations with opnet. In *Simulation Conference Proceedings, 1999 Winter*, volume 1, pages 307–314. IEEE, 1999.
- [43] F. Coras, A. Cabellos-Aparicio, and J. Domingo-Pascual. An analytical model for the LISP cache size. In *Proc. IFIP Networking*, May 2012.
- [44] F. Coras, D. Saucez, L. Iannone, and B. Donnet. On the performance of the LISP beta network. In *Proc. IFIP Networking*, June 2014.
- [45] S. Deering and R. Hinden. Rfc 2460: Internet protocol, 1998.
- [46] V. Ermagan, D. Farinacci, C. White, J. Skriver, D. Lewis, and F. Maino. Nat traversal for lisp. 2016.
- [47] D. Farinacci and et al. The locator/ID separation protocol (LISP). RFC 6830, Internet Engineering Task Force, January 2013.
- [48] D. Farinacci, A. Jain, I. Kouvelas, and D. Lewis. Locator/id separation protocol delegated database tree (lisp-ddt) referral internet groper (rig). Technical report, 2017.
- [49] D. Farinacci, D. Lewis, D. Meyer, and C. White. Lisp mobile node, draft-ietf-lisp-mn-01. Technical report, institution=IETF Internet draft, Apr. 2017.[Online]. Available: <https://tools.ietf.org/html/draft-ietf-lisp-mn-01>, 2017.
- [50] D. Farinacci, T. Li, S. Hanks, D. Meyer, and P. Traina. Rfc 2784-generic routing encapsulation (gre). *IETF, March 2000*, 2000.
- [51] D. Farinacci, T. Li, S. Hanks, D. Meyer, P. Traina, and R. Glenn. Generic routing encapsulation. *IETF (The Internet Engineering Task Force) Request for Comments RFC*, 2784, 2000.
- [52] D. Farinacci and D. Meyer. The locator/ID separation protocol Internet groper (LIG). RFC 6835, Internet Engineering Task Force, 2013.
- [53] B. Feng, H. Zhang, H. Zhou, and S. Yu. Locator/identifier split networking: A promising future internet architecture. *IEEE Communications Surveys & Tutorials*, 2017.
- [54] A. Ford, C. Raiciu, M. Handley, and O. Bonaventure. Tcp extensions for multipath operation with multiple addresses. Technical report, 2013.

- [55] V. Fuller. LISP-DDT: A new mapping database for LISP, June 2012. Nanog 55 Talk.
- [56] V. Fuller and et al. Locator/ID separation protocol (LISP) map-server interface. RFC 6833, Internet Engineering Task Force, January 2013.
- [57] V. Fuller, D. Farinacci, D. Meyer, and D. Lewis. Locator/ID separation protocol alternative logical topology (LISP+ALT). RFC 6836, Internet Engineering Task Force, January 2013.
- [58] V. Fuller, D. Lewis, and A. Ermagan, V. Jain. LISP Delegated Database Tree. Internet Draft, Work in Progress draft-ietf-lisp-ddt-01.txt, Internet Engineering Task Force, March 2013.
- [59] V. Fuller, D. Lewis, V. Ermagan, and A. Jain. Lisp delegated database tree. *draft-ietf-lisp-ddt-09*, Internet Engineering Task Force, 2017.
- [60] A. Galvani, A. Rodriguez-Natal, A. Cabellos-Aparicio, and F. Risso. Lisp-roam: network-based host mobility with lisp. In *Proceedings of the 9th ACM workshop on Mobility in the evolving internet architecture*, pages 19–24. ACM, 2014.
- [61] A. García-Martínez, M. Bagnulo, and I. Van Beijnum. The shim6 architecture for ipv6 multihoming. *IEEE Communications Magazine*, 48(9), 2010.
- [62] Y. Han, S. Ryu, Y.-J. Suh, and J. W.-K. Hong. Design and implementation of lisp controller in onos. In *NetSoft Conference and Workshops (NetSoft), 2016 IEEE*, pages 417–422. IEEE, 2016.
- [63] S. Hanks, D. Meyer, D. Farinacci, and P. Traina. Generic routing encapsulation (gre). 2000.
- [64] T. R. Henderson, M. Lacage, G. F. Riley, C. Dowell, and J. Kopena. Network simulations with the ns-3 simulator. *SIGCOMM demonstration*, 14(14):527, 2008.
- [65] M. Hoefling, M. Menth, and M. Hartmann. A survey on mapping systems for locator/identifier split internet routing. *IEEE Communications Surveys and Tutorials*, 15(4):1842–1858, November 2013.
- [66] L. Iannone and O. Bonaventure. On the cost of caching Locator/ID mappings. In *Proc. ACM SIGCOMM CoNEXT*, December 2007.
- [67] L. Iannone, D. Saucez, and O. Bonaventure. Locator/id separation protocol (lisp) map-versioning. Technical report, 2013.
- [68] M. Isah, S. Simpson, Y. Sani, and C. Edwards. Towards zero packet loss with lisp mobile node. In *Computing, Networking and Communications (ICNC), 2017 International Conference on*, pages 265–271. IEEE, 2017.
- [69] L. Jakab, A. Cabellos-Aparicio, F. Coras, D. Saucez, and O. Bonaventure. LISP-TREE: a DNS hierarchy to support the LISP mapping system. *IEEE Journal on Selected Areas in Communications*, 28(8):1332 – 1343, October 2010.
- [70] S. Jeon, J. Jang, and Y.-H. Kim. Network mobility support in the proxy mobile ipv6 domain. *Network*, 2010.

- [71] D. Johnson, C. Perkins, and J. Arkko. RFC 3775. *Mobility support in IPv6*, 2004.
- [72] J. Kim, L. Iannone, and A. Feldmann. A deep dive into the LISP cache. In *Proc. IFIP Networking*, May 2011.
- [73] J. Kim, L. Iannone, and A. Feldmann. Caching locator/id mappings: An experimental scalability analysis and its implications. *Computer Networks*, 57(4):897–909, 2013.
- [74] D. Klein, M. Hartmann, M. Hoefling, and M. Menth. Integration of lisp and lisp-mn into inet. In *Proceedings of the 5th International ICST Conference on Simulation Tools and Techniques*, pages 299–306. ICST (Institute for Computer Sciences, Social-Informatics and Telecommunications Engineering), 2012.
- [75] E. Lear. Nerd: A not-so-novel endpoint id (eid) to routing locator (rloc) database. 2013.
- [76] D. Lewis, D. Meyer, D. Farinacci, and V. Fuller. Interworking between locator/id separation protocol (LISP) and non-LISP sites. RFC 6832, Internet Engineering Task Force, January 2013.
- [77] Y. Li, A. Abouseif, L. Iannone, and D. Saucez. Lisp-views: Monitoring lisp at large scale. In *Teletraffic Congress (ITC 29), 2017 29th International*, volume 1, pages 178–186. IEEE, 2017.
- [78] Y. Li and et al. Performance evaluation of locator/identifier separation protocol through ripe atlas. In *Proceedings of the 2016 conference on ACM SIGCOMM 2016 Conference*, pages 561–562. ACM, 2016.
- [79] Y. Li and L. Iannone. Using ripe atlas to evaluate the locator/id separation protocol. In *Computer Communications Workshops (INFOCOM WKSHPS), 2016 IEEE Conference on INFOCOM 2016 Conference*, pages 696–697. IEEE, 2016.
- [80] Y. Li and L. Iannone. Assessing locator/identifier separation protocol interworking performance through RIPE Atlas. *Computer Networks*, pages –, 2017.
- [81] Y. Li, D. Saucez, L. Iannone, and B. Donnet. Stability and consistency of the lisp pull routing architecture. In *Proc. Workshop on Network Measurement (WNM)*, November 2016.
- [82] M. Mahalingam, D. Dutt, K. Duda, P. Agarwal, L. Kreeger, T. Sridhar, M. Bursell, and C. Wright. Virtual extensible local area network (vxlan): A framework for overlaying virtualized layer 2 networks over layer 3 networks. Technical report, 2014.
- [83] F. Maino, V. Ermagan, A. Cabellos, and D. Saucez. Lisp-security (lisp-sec). *draft-ietf-lisp-sec-14*, Internet Engineering Task Force, 2017.
- [84] T. Marill and L. G. Roberts. Toward a cooperative network of time-shared computers. In *Proceedings of the November 7-10, 1966, fall joint computer conference*, pages 425–431. ACM, 1966.
- [85] M. Menth, D. Klein, and M. Hartmann. Improvements to lisp mobile node. In *Teletraffic Congress (ITC), 2010 22nd International*, pages 1–8. IEEE, 2010.

- [86] D. Minoli. Security considerations for mipv6. *Mobile Video with Mobile IPv6*, pages 173–189.
- [87] P. Mockapetris. Rfc 1034: Domain names-concepts and facilities, 1987. URL: <ftp://ftp.isi.edu/in-notes/rfc1034.txt>.
- [88] M. Möller, R. Bye, K. Bsufka, S. A. Camtepe, and S. Albayrak. From simulation to emulation: an integrated approach for network security evaluation. *Lecture Notes in Informatics*, 192, 2011.
- [89] D. Montero, M. Siddiqui, R. Serral-Gracia, X. Masip-Bruin, and M. Yannuzzi. Securing the lisp map registration process. In *Global Communications Conference (GLOBECOM), 2013 IEEE*, pages 2145–2151. IEEE, 2013.
- [90] R. Moskowitz, T. Heer, P. Jokela, and T. Henderson. Host identity protocol version 2 (hipv2). Technical report, 2015.
- [91] R. Moskowitz and M. Komu. Host identity protocol architecture. *draft-ietf-hip-rfc4423-bis-18*, 2017.
- [92] A. R. Natal, L. Jakab, M. Portolés, V. Ermagan, P. Natarajan, F. Maino, D. Meyer, and A. C. Aparicio. Lisp-mn: mobile networking through lisp. *Wireless personal communications*, 70(1):253–266, 2013.
- [93] P. Nikander, A. Gurtov, and T. R. Henderson. Host identity protocol (hip): Connectivity, mobility, multi-homing, security, and privacy over ipv4 and ipv6 networks. *IEEE Communications Surveys & Tutorials*, 12(2):186–204, 2010.
- [94] E. Nordmark and M. Bagnulo. Shim6: Level 3 multihoming shim protocol for ipv6. Technical report, 2009.
- [95] M. O’DELL. Gse-an alternate addressing architecture for ipv6. <http://www.watersprings.org/pub/id/draft-ietf-ipngwgseaddr-00.txt>, 2006.
- [96] C. Perkins. RFC 3344. *IP Mobility Support for IPv4*, 8, 2002.
- [97] C. Perkins, D. Johnson, and J. Arkko. Mobility support in ipv6. Technical report, 2011.
- [98] C. E. Perkins. Mobile ip. *IEEE communications Magazine*, 35(5):84–99, 1997.
- [99] D. Phoomkiattisak and S. N. Bhatti. Control plane handoff analysis for ip mobility. In *Wireless and Mobile Networking Conference (WMNC), 2016 9th IFIP*, pages 65–72. IEEE, 2016.
- [100] D. Phung, S. Secci, D. Saucez, and L. Iannone. The openlisp control plane architecture. *IEEE Network*, 28(2):34–40, 2014.
- [101] A. Raheem, A. Lasebae, M. Aiash, and J. Loo. Supporting communications in the iots using the location/id split protocol: a security analysis. In *Future Generation Communication Technology (FGCT), 2013 Second International Conference on*, pages 143–147. IEEE, 2013.

- [102] M. K. Rana, B. Sardar, S. Mandal, and D. Saha. Implementation and performance evaluation of a mobile ipv6 (mipv6) simulation model for ns-3. *Simulation Modelling Practice and Theory*, 72:1–22, 2017.
- [103] I. RFC0791. Internet protocol. *J. Postel. September*, 1981.
- [104] A. Rodriguez-Natal, S. Barkai, V. Ermagan, D. Lewis, F. Maino, and D. Farinacci. Software defined networking extensions for the locator/id separation protocol. *IETF LISP Working Group Internet-Draft*, 2014.
- [105] J. Rosenberg et al. Sip: Session initiation protocol (jun. 2008) retrieved at <http://tools.ietf.org/html/rfc3261>. *Relevant pages provided*.
- [106] D. Saucez and et al. A first measurement look at the deployment and evolution of the locator/id separation protocol. *ACM SIGCOMM Computer Communication Review*, 43(1):37–43, April 2013.
- [107] D. Saucez and L. Iannone. O. bonaventure," locator/id separation protocol (lisp) threat analysis. Technical report, RFC 7835, DOI 10.17487/RFC7835, April 2016,< <http://www.rfc-editor.org/info/rfc7835>.
- [108] D. Saucez, L. Iannone, O. Bonaventure, et al. Openlisp: An open source implementation of the locator/id separation protocol. *ACM SIGCOMM Demos Session*, pages 1–2, 2009.
- [109] D. Saucez, L. Iannone, O. Bonaventure, and D. Farinacci. Designing a deployable internet: The locator/identifier separation protocol. *IEEE Internet Computing*, 16(6):14–21, 2012.
- [110] D. Saucez, L. Iannone, A. Cabellos, and F. Coras. Locator/id separation protocol (lisp) impact. Technical report, 2016.
- [111] A. Stockmayer, M. Schmidt, and M. Menth. jlisp: An open, modular, and extensible java-based lisp implementation. In *Teletraffic Congress (ITC 28), 2016 28th International*, volume 1, pages 205–208. IEEE, 2016.
- [112] Z. Tang, Y. Zhou, W. Deng, and B. Wang. Lisp-hnm: Integrated fast host and network mobility control in lisp networks. In *Modeling and Optimization in Mobile, Ad Hoc, and Wireless Networks (WiOpt), 2017 15th International Symposium on*, pages 1–6. IEEE, 2017.
- [113] R. Thayer, N. Doraswamy, and R. Glenn. Rfc 2411: Ip security document roadmap. *Internet standard, November*, 1998.
- [114] C. Vogt. Six/one: A solution for routing and addressing in ipv6. *INTERNET DRAFT, draft-vogt-rrg-six-one-01.txt*, 2007.
- [115] Y. Wu, K. Chen, K. Xue, and D. Ni. Nemo-based mobility management in lisp network. In *Wireless Communications and Signal Processing (WCSP), 2014 Sixth International Conference on*, pages 1–6. IEEE, 2014.
- [116] H. Zhang, M. Chen, and Y. Zhu. Evaluating the performance on IDLoc mapping. In *Proc. IEEE Global Communications Conference (GLOBECOM)*, November 2008.

List of Publications

- Y. Li, D. Saucez, L. Iannone, and B. Donnet. Stability and consistency of the lisp pull routing architecture. In *Proc. Workshop on Network Measurement (WNM)*, November 2016
- Y. Li and L. Iannone. Using ripe atlas to evaluate the locator/id separation protocol. In *Computer Communications Workshops (INFOCOM WKSHPS), 2016 IEEE Conference on INFOCOM 2016 Conference*, pages 696–697. IEEE, 2016
- Y. Li and et al. Performance evaluation of locator/identifier separation protocol through ripe atlas. In *Proceedings of the 2016 conference on ACM SIGCOMM 2016 Conference*, pages 561–562. ACM, 2016
- Y. Li, A. Abouseif, L. Iannone, and D. Saucez. Lisp-views: Monitoring lisp at large scale. In *Teletraffic Congress (ITC 29), 2017 29th International*, volume 1, pages 178–186. IEEE, 2017
- Y. Li and L. Iannone. Assessing locator/identifier separation protocol interworking performance through RIPE Atlas. *Computer Networks*, pages –, 2017

Appendix A

Source Code

Some of our implementations are modified from previous open source projects. So we upload related codes to GitHub in my personal repositories as <https://github.com/SeleneLI>. It is convenient to download and explore the codes. A brief introduction of each repository is provided as follows.

A.1 LISP Stability and Consistency analyzer

<https://github.com/SeleneLI/TracesAnalyzer>

This project is totally originally written to analyze the stability and consistency of LISP mapping system.

A.2 LISP-Views

<https://github.com/SeleneLI/LISP-Views>

This project is the original codes for LISP-Views architecture.

A.3 LISP measurements on Atlas

<https://github.com/SeleneLI/Atlas>

This project is used to conduct the experiment on RIPE Atlas and analyze the ping / traceroute results for LISP interworking performance.

A.4 LISP implementation on ns-3

https://github.com/SeleneLI/LISP_ns-3

This project is an implementation of LISP simulator on ns-3 to evaluate LISP mobility performance.